Costly State Verification with Limited Commitment

Amirreza Ahmadzadeh*

September 20, 2025

This paper examines a principal-agent model. The principal mandates actions and conducts costly inspections without transfers. The principal prefers lower actions, while the agent prefers higher actions. The agent has private information about his type and is protected by ex-post participation constraints and he rejects any action below his private type. The principal faces a trade-off between the benefit from mandating lower actions and the risk that the agent rejects actions and chooses his outside option. We analyze various levels of the principal's commitment ability. First, if the principal can commit to both stochastic inspection and the action in case inspection does not take place, and if the principal's fear of ruin is greater than the agent's, then a deterministic inspection policy is optimal. Second, if the principal cannot commit to either inspections or actions, the principal's highest equilibrium payoff involves only deterministic inspection strategies. Finally, if the inspection cost is low and the principal commits to inspecting whenever requested by the agent, the principal can achieve the payoff of the optimal deterministic inspection policy.

Keywords: Costly state verification, mechanism design, cheap talk, inspection, limited commitment, regulation.

I am indebted to my advisors (in alphabetical order) Johannes Hörner, Thomas Mariotti, Anna Sanktjohanser, and Jean Tirole for their support and guidance. I am especially grateful to Marina Halac for her insights and advice. I also thank Christian Hellwig, Michelle Hyun, Bruno Jullien, Deniz Kattwinkel, David Martimort, Nour Meddahi, Teemu Pekkarinen, François Salanie, Anna Sanktjohanser, Sara Shahanaghi, Patrick Rey, Stephan Waizmann, as well as the audiences at the TSE Microeconomic Theory workshop, EARIE 2023 Rome, and Oligo workshop 2023 Padova for their comments and feedback.

*Toulouse School of Economics (TSE), amirreza.ahmadzadeh@tse-fr.eu.

1. Introduction

Addressing the climate crisis requires widespread carbon dioxide reduction to prevent or lessen severe environmental impacts. However, market incentives alone often fall short in prompting companies to cut pollution. Therefore, government intervention plays a crucial role in adopting and advancing essential technologies.

Government action can take various forms, including subsidies, regulatory standards, inspections, R&D support, and emission-related taxes or tradable permits. Regulatory standards are the most common and encompass various approaches. Performance standards, for instance, set limits on emissions per unit of product, such as restricting CO_2 emissions to a certain amount per kilowatt-hour of electricity generated.[1]

In this paper, we examine two key regulatory tools: standards and inspections. We explore the optimal regulatory behavior for enforcing standards and conducting inspections. We aim to understand how different levels of the regulator's commitment power affect the optimal regulatory policies and we identify the key regulatory instruments that require commitment.

Firms incur different costs for adjusting their technology to meet mandated standards. The regulator can set these standards to encourage firms to adopt new technologies. Additionally, the regulator has the authority to inspect firms to assess the costs associated with implementing these technologies.

Our model presumes that the benefits of adopting a technology are known to the regulator, while the costs are known only to the firms. The regulator can learn these costs through inspection. Importantly, we assume that the regulator cannot compel a firm to adopt a technology. The firm can move the business to another country or state. Therefore, the regulator's objective is to develop a cost-effective strategy that incentivizes efficient firms to adopt abatement technologies.

We analyze the impact of varying levels of the regulator's commitment ability. Lack of commitment can stem from several factors: First, *legal enforcement:* For legal enforcement to be effective, the court must remain impartial and avoid collusion with any party. Additionally, the contract must be clearly defined and ideally complete. Verifying deviations from the contract can be particularly challenging, especially when it involves non-deterministic promises. Second, *political pressure:* in the context of environmental regulation, regulators may face political pressure

[1] See: Solomon (2007), section 13.2.1.1 regulations and standards.

by not implement ex-post efficient standards that are ex-ante optimal. For instance, it may be optimal to commit to allowing a firm to pollute more after inspection takes place. Third, *reputation:* reputation can play a significant role in ensuring that the regulator adheres to pre-committed promises, especially when interacting with multiple agents, either sequentially or simultaneously. Such interactions may enable the regulator to commit to a consistent decision-making process. However, this may require additional assumptions, such as the absence of macroeconomic shocks and the presence of post-facto incentives to align observations with incentive schemes.

In Section 2, we introduce the model. A Principal (regulator) can commit to the probability of inspection and actions (standards) if inspection does not take place.[2] The Agent (firm) prefers higher actions (e.g., more pollution and emissions, lower performance standards), while the Principal prefers lower actions (e.g., less pollution and emissions, higher performance standards) and cannot use transfers. The Agent is protected by an ex-post participation constraint; therefore, the Principal faces a trade-off between the benefit from mandating lower actions and the risk that the Agent rejects actions and chooses his outside option.

In Section 3, we present the Principal's optimization problem and demonstrate that if the Principal's fear of ruin is greater than the Agent's, then deterministic inspection is the optimal choice.[3] In Section 3.1, we focus on deterministic inspection. We show that the optimal mechanism with deterministic inspection is a cutoff policy that divides types into three regions. Efficient types are never inspected and face a cap on their actions. Intermediate types are inspected and are mandated to their first-best action. Finally, inefficient types are excluded. This structure highlights the importance of inspecting intermediate types, as it limits the efficient types' rents while ensuring a low action for efficient types. Finally, we show that when restricting attention to deterministic mechanisms, commitment to actions in the case of inspection (post-inspection) does not increase the Principal's payoff.

In Section 4, we explore various levels of the Principal's commitment ability. Section 4.1 assumes that the Principal cannot commit to the probability of inspection

[2] We assume the Principal cannot commit to actions when inspection occurs. Later in the section, we discuss the properties of the optimal mechanism if the Principal can commit to actions when inspection takes place.

[3] We provide an upper bound on the payoff of stochastic inspection relative to deterministic inspection. This upper bound is calculated as the aggregate difference in fear of ruin between the Principal and the Agent, scaled by the Principal's marginal utility.

but can commit to actions when inspection does not take place. We demonstrate that there exists an equilibrium identical to the optimal deterministic policy (by restricting attention to deterministic inspection and when the Principal can fully commit to all instruments). Furthermore, we show that if the Agent's utility exhibits log-supermodularity, the Principal's preferred equilibrium will involve only deterministic inspections.

Section 4.2 contrasts with Section 4.1 by assuming that the Principal can only commit to an inspection policy. We demonstrate that if the Principal cannot commit to actions, focusing on semi-separating equilibria with only two groups of pooling types and deterministic inspections is sufficient to find the Principal's preferred equilibrium. Furthermore, we show this equilibrium involves the inspection of inefficient types.

In Section 4.3, we assume the Principal lacks the power to commit to any of her instruments. The structure of the Principal's preferred equilibrium is similar to the optimal deterministic policy, though the thresholds may differ. Finally, in Section 4.4, we explore a partial commitment setting. If the cost of inspection is not high, and the Principal commits to inspection whenever the Agent requests, then the Principal can achieve the optimal deterministic inspection policy through this partial commitment.

Relationship to the literature. The paper contributes to the literature on costly state verification (CSV). In addressing a contracting design problem, Townsend (1979) examines the optimal insurance contract between a risk-neutral principal (investor) and a risk-averse agent (entrepreneur). At the time of contracting, both parties possess the same information. Once the contract is established, the agent privately observes the project's income. The agent reports an income and incurs a cost ifnhe verifies this income. A contract specifies two components for each income report: whether the agent should verify the income and the transfer from the agent to the principal. Townsend's optimal contract maximizes the agent's ex-ante payoff, subject to the principal's ex-ante individual rationality constraint (IR). This optimal contract resembles a debt contract, in which the agent verifies incomes below a certain threshold.

Gale and Hellwig (1985) study a similar problem involving a risk-neutral borrower and a risk-neutral lender. The optimal mechanism maximizes the borrower's expected utility, subject to the lender's zero-profit constraint and the borrower's incentive compatibility constraints. Similarly, the agent incurs a cost to verify the project's income. Both Townsend (1979) and Gale and Hellwig (1985) focus on deterministic inspection.

Border and Sobel (1987) consider a more general mechanism involving stochastic inspections and bounded pre-inspection and post-inspection transfers. They assume that the principal cannot make a net payment to the agent. They demonstrate that the probability of an inspection should decrease with the agent's wealth. Through an example, they illustrate that if the principal aims to maximize expected revenue net of inspection costs, the optimal contract involves large rewards and infrequent inspections. Mookherjee and Png (1989) assume the borrower is risk-averse and demonstrate that the optimal contract is stochastic.

This paper differs from the existing literature in two ways. First, the Principal does not have full commitment power and cannot commit to the mechanism after inspection occurs. Second, our paper does not consider transfers as a tool for the designer.[4]

Another application of CSV is for optimal allocations of objects and collective choice problems. Ben-Porath et al. (2014) consider a Principal who allocates an object to one of I Agents. The Principal cannot use transfers but can check the private information of each Agent at a cost. The private information is the value of the object for each Agent. Mylovanov and Zapechelnyuk (2017) study a similar problem with a different verification technology and limited punishments. They assume the Principal can verify information after allocating the object and, contingent on this observation, can destroy a fraction of the Agent's payoff. Li (2020) studies the connection between costly verification and limited punishment. Patel and Urgun (2022) assume money burning as a new instrument for the Principal and study the optimal allocation problem with CSV, and Erlanson and Kleiner (2020) investigate the optimal allocation and collective choice problems.

Another branch of mechanism design with CSV is in monopoly regulation. Baron and Besanko (1984) extend Baron and Myerson (1982) to allow random and costly inspections. In their setting, monopolist pricing is a two-part tariff consisting of a fixed charge and a unit price. Palonen and Pekkarinen (2022) consider a CSV regulation principal-agent problem with a different approach. They assume the Agent can reduce the probability of being verified by engaging in costly avoidance actions. The paper assumes a linear and exogenous punishment function if the

[4]There are other differences: most works in the CSV literature on financial markets assume a competitive borrowing market, maximizing the borrower's utility subject to the lender's outside option and the borrower's truth-telling condition. Moreover, most works, unlike in this paper, the state of nature is not known by any party at the contract date.

Agent is caught being untruthful, and no reward if the Agent is truthful. The Principal maximizes the expected weighted sum of the Agent's payoff and transfers, net of monitoring costs, subject to the incentive compatibility and participation constraints. These papers depart from ours for two reasons. First, we have an ex-post participation constraint, so punishments are endogenous and restricted to the utility of the Agent. Second, they do not study different commitment abilities of the Principal.

The paper most closely related to ours is Halac and Yared (2020), which examines a CSV principal-agent delegation problem where the agent has a bias towards higher actions. When the agent's bias becomes extreme and when the principal has full commitment, their model resembles ours. However, unlike their study, we do not limit our analysis to deterministic mechanisms, and we assume the agent is not protected by ex-post participation constraint. Further differences are discussed following the introduction of a related result in sections 3.1 and 4.2.

Section 4 connects with the literature on cheap talk models, following the work of Crawford and Sobel (1982).[5] Khalil (1997) contributes to the literature on CSV without commitment by examining a model where the Principal cannot commit to an inspection policy. His model involves costly signaling rather than cheap talk. Melumad and Mookherjee (1989) also considers a setting without commitment on inspection. The main differences are that the Principal can commit to transfers in the event of an inspection, and the Agent cannot reject the mechanism ex-post.

Banks (1989) explores a model where an agent requests a budget. The principal lacks commitment power, utilities are linear in transfer, and inspection is imperfect. Unlike our setting (see Section 4.3), full pooling is the only equilibrium. In any equilibrium, no information is conveyed from the agent to the principal.

[5]In the no commitment case, our model is not exactly the same as cheap talk models, since the Agent (sender) can accept or reject the proposed action by the Principal (receiver). Therefore there is an addition move.

2. Model

Players, and information structure. There are two players, a Principal (she) and an Agent (he). The Agent has type $\theta \in [\underline{\theta}, \overline{\theta}]$ drawn from a commonly known cumulative distribution function $F(\cdot)$, and probability distribution function $f(\cdot) > 0$. The type is the Agent's private information.

Mechanism and action. The Principal chooses and commits to a mechanism.[6] The mechanism $\mathbb{M} = (M, i(m), a(m))$ has three components: the message space M, the probability of inspection $i(m)$ and the mandated action $a(m)$. The probability of inspection as a function of message $m \in M$ is $i(m) \in [0, 1]$, and inspection allows learning the true type of the Agent. The Agent's action in case of no inspection is $a(m) \in \mathbb{R}_+$. In case of inspection the Principal, without commitment, mandates $a^I(m, \theta) \in \mathbb{R}_+$. Action $a^I(m, \theta) \in \mathbb{R}_+$ is a function of the message m and the true type θ upon inspection.

Payoffs. Inspection costs $\phi > 0$ to the Principal. We assume the Agent is secured by **ex-post participation**: the Agent can accept or reject the final action. If the Agent rejects the mandated action, the payoffs of both players are zero. The payoff of the Agent with type θ and action a is $u(\theta, a)$.[7] The Principal's payoff is

$$v(\theta, a)\mathbb{1}_{\text{accept}} - \phi\mathbb{1}_{\text{inspection}}.$$

Timing. The Principal commits to the mechanism. Nature draws a type θ, and the Agent learns it privately. The Agent sends a message $m \in M$. The Principal inspects with probability $i(m)$. If he does not inspect, with commitment and according to the mechanism, he mandates an action $a(m)$. If he inspects, without commitment, he mandates an action $a^I(m, \theta)$. The Agent accepts or rejects the action. Figure 2 shows the timing of the model.

We consider pure strategy Perfect Bayesian equilibrium as the solution concept and maintain the following assumptions throughout the paper.

Assumption 1 *(Agent's utility)* (i) The Agent's utility $u(\theta, a)$ is twice continuously differentiable for all $(\theta, a) \in [\underline{\theta}, \overline{\theta}]^2$. (ii) A type $\theta \in [\underline{\theta}, \overline{\theta}]$ of the Agent gets **zero** payoff when $a = \theta$, and prefers **higher actions**. Formally

$$u(\theta, \theta) = 0, \text{ and } u_a(\theta, a) > 0 \text{ for all } (\theta, a) \in [\underline{\theta}, \overline{\theta}]^2.$$

[6]In section 4, we study different commitment ability of the Principal.

[7]Another interpretation of this model is that the Agent has a private outside option θ.

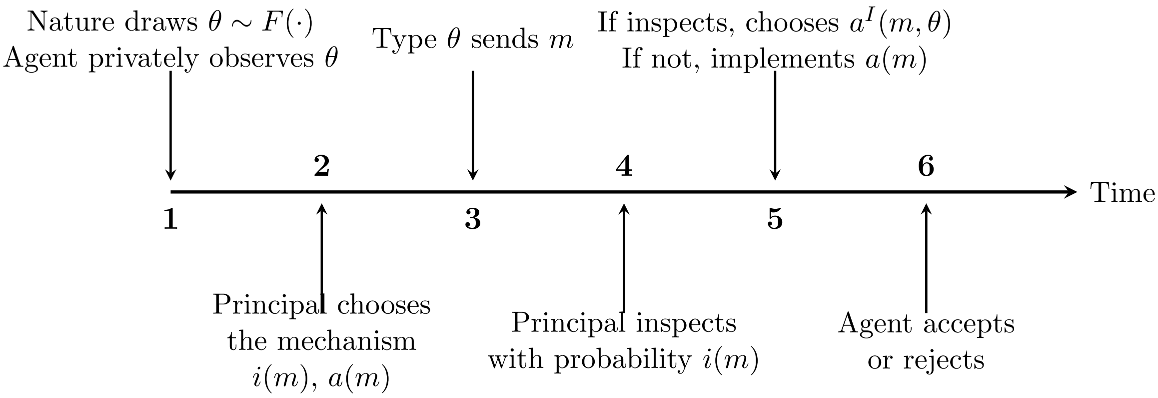

Figure 1: Timing.

Assumption 2 *(Principal's utility)* (i) Utility $v(\theta, a)$ is twice continuously differentiable for all $(\theta, a) \in [\underline{\theta}, \overline{\theta}]^2$. (ii) All types of the Agent are **valuable** for the Principal. Moreover, **lower** types are more valuable for the Principal than **higher** types. Formally $v(\theta, \theta)$ is positive and weakly decreasing in θ, for all $\theta \in [\underline{\theta}, \overline{\theta}]$. (iii) The Principal prefers lower actions for all types of the Agent. Formally

$$v_a(\theta, a) < 0 \text{ for all } (\theta, a) \in [\underline{\theta}, \overline{\theta}]^2.$$

Assumption 1 simply says the Agent's utility is zero when his action equal to his type ($a = \theta$), and it is increasing in his action. More precisely, the Agent prefers higher actions by only considering the support of CDF $F(.)$, i.e. $a \in [\underline{\theta}, \overline{\theta}]$. The assumption is silent for actions above $\overline{\theta}$.

Assumption 2 states lower types are more valuable for the Principal than higher types. This assumption ensures that the existence of all types of the Agent is valuable for the Principal. If we assume $v(\overline{\theta}, \overline{\theta}) < 0$, then the Principal can exclude this type by mandating actions to be less than $\overline{\theta}$. The principal prefers lower actions. However, if she chooses an action less than θ, the agent will reject the action (due to the ex-post participation constraint).

By Assumption 1 we can conclude that, the payoff of the Agent at final action a given type θ is $u(\theta, a)\mathbb{1}_{a \geq \theta}$, and the Principal's payoff is $v(\theta, a)\mathbb{1}_{a \geq \theta} - \phi \mathbb{1}_{\text{inspection}}$. [8]

Discussion on the commitment ability of the Principal to actions if inspection takes place, i.e. $a^I(\cdot, \cdot)$*:* Suppose the Principal is able to commit to $a^I(\hat{\theta}, \theta)$. Therefore,

[8] In Assumption 1, $u(\theta, \theta)$ can be changed with $u(\theta, a^*(\theta)) = \underline{u}(\theta)$, where $a^*(\theta)$ is the action that gives the Agent the same utility as his outside option, $\underline{u}(\theta)$.

$a^I(\hat{\theta}, \theta)$ for $\hat{\theta} \neq \theta$ is off the equilibrium path, and the Principal implements the maximum punishment: $a^I(\hat{\theta}, \theta) \leq \theta$. The main challenge is to determine the optimal action when the Principal inspects and the agent is truthful, i.e., $a^I(\theta, \theta)$. In Appendix (Full Commitment), we show that if $\lim_{a\to\infty} -v(\theta, a) = \lim_{a\to\infty} -\frac{u(\theta,a)}{v(\theta,a)} = \infty$, then the optimal mechanism does not exist.[9] Moreover, we show that the first-best outcome can be approximated through mechanisms that set the ex-post efficient action in the case of no inspection, $a(\theta) = \theta$, and inspect all types with a small probability. In the case of inspection, these mechanisms set a large action if the agent is truthful, i.e., $a^I(\theta, \theta)$ approaches infinity.

In the context of environmental regulation, a large $a^I(\theta, \theta)$ is analogous to the following policy: the regulator allows the agent to release a substantial amount of pollutants into the environment. However, we do not see these types of incentive schemes in practice. The non-existence of the optimal solution has been raised in settings with transfers.[10] To implement this mechanism, the regulator must possess strong commitment power. One possible approach, as explored by Border and Sobel (1987) in a setting with transfers, is to impose an exogenous upper bound on rewards, i.e., $a^I \leq \overline{a}$. Another approach is to assume that the principal cannot commit to actions following an inspection.[11]

3. Results

In this section, we begin by formulating the Principal's problem as an optimization problem. We then derive properties of the optimal mechanism and identify the conditions under which deterministic inspection is optimal. Next, we examine how different levels of *fear of ruin* (see section ... for the definition) of parties impact the optimal mechanism. Finally, we discuss the value of stochastic inspection when deterministic inspection is not optimal.

In any equilibrium, the Principal mandates action $a^I(m, \theta) = \theta$. This is due to the lack of commitment of the Principal after inspection occurs. By the revelation principle we can restrict the message space M to the types space $[\underline{\theta}, \overline{\theta}]$, and mechanisms to direct mechanisms. The Principal chooses a direct mechanism $\mathbb{M} = \big(i(\hat{\theta}), a(\hat{\theta})\big)$, where $\hat{\theta}$ is the reported type of the Agent. The Agent's expected

[9] An example is $v(\theta, a) = -|\theta - a|^{\frac{1}{2}}$ and $u(\theta, a) = a - \theta$.

[10] See, for example, Border and Sobel (1987) and Ahmadzadeh and Waizmann (2024).

[11] This idea has also been suggested by Halac and Yared (2020). See the discussion on stochastic inspection in this paper.

payoff given his type θ and the report $\hat{\theta}$ is

$$\begin{aligned}\pi(\hat{\theta},\theta) &\equiv (1-i(\hat{\theta}))\Big(u(\theta,a(\hat{\theta}))\Big)\mathbb{1}_{a(\hat{\theta})\geq\theta} + i(\hat{\theta})\Big(u(\theta,a^I(\hat{\theta},\theta))\Big)\mathbb{1}_{a^I(\hat{\theta},\theta)\geq\theta} \\ &= (1-i(\hat{\theta}))\Big(u(\theta,a(\hat{\theta}))\Big)\mathbb{1}_{a(\hat{\theta})\geq\theta}.\end{aligned}$$

The Principal's expected payoff if the Agent with type θ reports $\hat{\theta}$ is

$$(1-i(\hat{\theta}))\Big(v(\theta,a(\hat{\theta}))\Big)\mathbb{1}_{a(\hat{\theta})\geq\theta} + i(\hat{\theta})\Big(v(\theta,\theta)-\phi\Big).$$

The Principal's problem:

Define the problem $\mathbb{P}$ as follows:

$$\mathbb{P} \quad : \quad \max_{i(\cdot),a(\cdot)} \mathbb{E}\Big[(1-i(\theta))\Big(v(\theta,a(\theta))\Big)\mathbb{1}_{a(\theta)\geq\theta} + i(\theta)\Big(v(\theta,\theta)-\phi\Big)\Big],$$

subject to combined ex-post participation constraints and the truth telling conditions (IC) for the Agent:

$$\pi(\theta,\theta) \geq \sup_{\hat{\theta}} \pi(\hat{\theta},\theta),$$

for all $\theta \in [\underline{\theta},\overline{\theta}]$. For simplicity, let $\pi(\theta)=\pi(\theta,\theta)$. Note that the mandated action without inspection $a(\hat{\theta})$, can be less than the true type of the Agent, which implies that the Agent rejects the offered action, or in other words, the Principal can exclude some types. If the Principal wants to exclude a type (through a low $a(\cdot)$), it is without loss of generality, to assume $a(\hat{\theta})=\underline{\theta}$. We assume the solution to problem $\mathbb{P}$ exists and is piece-wise continuous.

Assumption 3 *The solution to problem $\mathbb{P}$ exists and it is piece-wise continuous.*

Types with zero payoffs:

Let us begin the analysis with types that have zero payoffs. Let $\mathbb{M}=\big(i(\cdot),a(\cdot)\big)$, satisfies IC. Define

$$\tilde{\theta}=\{\inf\theta|\pi(\theta)=0\}.$$

Note that $\tilde{\theta}\leq\overline{\theta}$. The following lemma studies the structure of $\mathbb{M}$ for $\theta\in[\tilde{\theta},\overline{\theta}]$.

Lemma 1 *For all IC mechanisms $\mathbb{M}$:*

(i) $\pi(\theta) = 0$ *for all* $\theta \geq \tilde{\theta}$.

(ii) If $a(\theta) < \theta$, *then it is without loss of generality that* $a(\theta) = \underline{\theta}$.

(iii) If $a(\theta) = \underline{\theta}$, *then optimality of* $\mathbb{M}$ *implies that for all* $\theta' > \theta$, $a(\theta') = \underline{\theta}$.

(iv) If $a(\theta) \leq \theta$, *then for all* $\theta' > \theta$, *either* $a(\theta') = \underline{\theta}$ *or* $i(\theta') = 1$.

The proof of Lemma 1 is in Appendix. For all $\theta \leq \tilde{\theta}$, either $a(\theta) \leq \theta$ or $i(\theta) = 1$. Therefore Lemma 1 (iii), and (iv) imply that a necessary condition for $\mathbb{M}$ to be IC is that there exists $\tilde{\tilde{\theta}} \in [\tilde{\theta}, \overline{\theta}]$, such $i(\theta) = 1$ for $\theta \in [\tilde{\theta}, \tilde{\tilde{\theta}}]$, and $a(\theta) = \underline{\theta}$ for $\theta \in [\tilde{\tilde{\theta}}, \overline{\theta}]$. Using the structure of Lemma 1, one can conclude that types $\theta \leq \tilde{\theta}$ do not have incentive to mimic types higher than $\tilde{\theta}$. In addition, a necessary and sufficient condition for incentives of types higher than $\tilde{\theta}$, is that $a(\theta) \leq \tilde{\theta}$ for $\theta \leq \tilde{\theta}$. This implies $a(\tilde{\theta}) = \tilde{\theta}$.

Corollary 1 *In any optimal* $\mathbb{M}$, $a(\tilde{\theta}) = \tilde{\theta}$.

An upper bound on the payoff of the Principal:

Now we establish an upper bound on the Principal's payoff. Using this upper bound, we then identify the conditions under which deterministic inspection is optimal. Finally, we present the conditions where the optimal outcome matches the upper bound.

Lemma 2 *Denote* $\mathbb{M} = \big(i_*(\cdot), a_*(\cdot)\big)$ *the solution to problem* $\mathbb{P}$ *replacing global IC constraints by local IC constraints. Suppose there exists* θ^* *such that* $i_*(\theta) = 0$ *for all* $\theta \leq \theta^*$. *Then the value of* $\mathbb{P}$ *is (weakly) less than*

$$\int_{\underline{\theta}}^{\theta^*} \Big(v(\theta, a_*(\underline{\theta})) - \frac{u(\theta^*, a_*(\underline{\theta}))}{u_a(\theta^*, a_*(\underline{\theta}))} v_a(\theta, a_*(\underline{\theta}))\Big)\, dF(\theta) + \int_{\theta^*}^{\tilde{\tilde{\theta}}} \Big(v(\theta, \theta) - \phi\Big)\, dF(\theta).$$

The proof of Lemma 2 is in the Appendix. Lemma 2 finds an upper bound on the Principal's payoff. Replacing global IC constraints by local IC constraints weaken the constraints. Therefore the payoff of the Principal should be weakly lower than a version of a problem only with local IC constraints.

To grasp the intuition behind the proof of Lemma 2, we first need to reformulate problem $\mathbb{P}$ by substituting the global IC constraints with local IC constraints.

$$\max_{i(.),a(.),\tilde{\theta},\tilde{\tilde{\theta}}} \int_{\underline{\theta}}^{\tilde{\tilde{\theta}}} \Big[(1-i(\theta))\Big(v(\theta,a(\theta))-v(\theta,\theta)+\phi\Big)+\Big(v(\theta,\theta)-\phi\Big)\Big]dF(\theta),$$

subject to

$$(1-i(\theta))u(\theta,a(\theta)) \geq \ (1-i(\hat{\theta}))u(\theta,a(\hat{\theta})),$$

for all θ, $\hat{\theta}$ and $a(\theta) \geq \theta$, $\tilde{\theta} \leq \tilde{\tilde{\theta}}$ and $a(\tilde{\theta}) = \tilde{\theta}$. The argument is analogous to methods used in the Calculus of Variations. A global variation in the inspection probability involves decreasing $(1-i(\theta))$ to $(1-\beta)(1-i(\theta))$ for types in $[\theta^*,\tilde{\theta}]$. In words, β percentage decrease in probability of not inspection. After this variation types above θ^* do not mimic types above θ^* since $(1-\beta)$ cancels from both sides of IC inequalities. However, types above θ^* would like to mimic types below θ^*. In order to keep the incentives unchanged, a variation on $a_*(\underline{\theta})$ is $a_*^\beta(\underline{\theta})$ such that $a_*^\beta(\underline{\theta})$ solves

$$u(\theta^*,a_*^\beta(\underline{\theta})) = (1-\beta)u(\theta^*,a_*(\underline{\theta})).$$

For small enough $\beta > 0$, the marginal change in the payoff by decreasing β percentage of not inspection is

$$\int_{\theta^*}^{\tilde{\theta}}(1-i_*(\theta))\Big(v(\theta,a_*(\theta))-v(\theta,\theta)+\phi\Big)dF(\theta).$$

For large enough $\beta > 0$, the marginal change in the payoff by changing $a_*(\underline{\theta})$ to $a_*^\beta(\underline{\theta})$ is

$$\int_{\underline{\theta}}^{\theta^*} \frac{u(\theta^*,a_*(\underline{\theta}))}{u_a(\theta^*,a_*(\underline{\theta}))}v_a(\theta,a_*(\underline{\theta}))\, dF(\theta).$$

Similar arguments hold for $\beta > 0$. At the optimal value, the total marginal change should be zero, therefore

$$\int_{\theta^*}^{\tilde{\theta}}(1-i_*(\theta))\Big(v(\theta,a_*(\theta))-v(\theta,\theta)+\phi\Big)dF(\theta) = \int_{\underline{\theta}}^{\theta^*} -\frac{u(\theta^*,a_*(\underline{\theta}))}{u_a(\theta^*,a_*(\underline{\theta}))}v_a(\theta,a_*(\underline{\theta}))\, dF(\theta).$$

Using the above equality, one can derive the payoff stated in Lemma 2. The left side represents the marginal increase in payoff by reducing one percent of not inspection for types in $[\theta^*,\overline{\theta}]$. The right side represents the marginal increase in the payoff by adjusting $a_*(\underline{\theta})$ so that the utility of type θ^* decreases by one percent. In the optimal mechanism, the gain by increasing inspection needs to trade off the gain by decreasing $a_*(\underline{\theta})$.

Lemma 2 indicates that the net gain from stochastic inspection, as opposed to deterministic inspection, is

$$\int_{\underline{\theta}}^{\theta^*} \left(v(\theta, a_*(\underline{\theta})) - \frac{u(\theta^*, a_*(\underline{\theta}))}{u_a(\theta^*, a_*(\underline{\theta}))} v_a(\theta, a_*(\underline{\theta})) - v(\theta, \theta) \right) dF(\theta)$$

If the above payoff is not strictly positive, then deterministic inspection is optimal. Particularly, if for all $\theta^* \leq a_*(\underline{\theta})$ ($\underline{\theta} \leq \theta^* \leq a_*(\underline{\theta}) \leq \overline{\theta}$) the integral is negative, then the optimal inspection is deterministic.

Fear of ruin:

Before proceeding let us define the concept of *fear of ruin* which was introduced by Aumann and Kurz (1977) in the context of taxation policies. We explain this concept in our setting. Suppose the agent's type is θ and the agent is considering a bet where he risks his entire utility $u(\theta, a)$ for a potential small increase in utility to $u(\theta, a + a')$, where $a' > 0$. The probability q of ruin must be very small for the Agent to be indifferent between taking the bet and keeping his current utility, i.e., $u(\theta, a) = (1 - q)u(\theta, a + a')$. Thus, the more unwilling the agent is to risk ruin, the smaller q will be. Therefore, q serves as an inverse measure of the agent's aversion to risking ruin. The probability of ruin (q) per potential extra gain (a') represents the fear of ruin, i.e., $\lim_{q \to 0} q/a' = u(\theta, a(\theta))/u_a(\theta, a(\theta))$. Therefore the *fear of ruin* of the principal for type θ and action a is $(v(\theta, a) - v(\theta, \theta))/(v_a(\theta, a))$. It is well-known that both $v(\theta, a)$ and $u(\theta, a)$ are concave in a and if Arrow-Pratt coefficient of the Principal is higher than the Agent, then the principal's fear of ruin is also higher.[12]

Lemma 2 provides an intuitive upper bound on the net gain from stochastic inspection, compared to deterministic inspection.

$$\int_{\underline{\theta}}^{\theta^*} v_a(\theta, a_*(\underline{\theta})) \left(\frac{v(\theta, a_*(\underline{\theta})) - v(\theta, \theta)}{v_a(\theta, a_*(\underline{\theta}))} - \frac{u(\theta^*, a_*(\underline{\theta}))}{u_a(\theta^*, a_*(\underline{\theta}))} \right) dF(\theta).$$

The upper bound is the aggregate difference of fear of ruin of the Principal for types below θ^* and the Agent for type θ^* evaluated at $a_*(\underline{\theta})$ scaled by the marginal utility of the Principal. The below theorem follows from this observation.

[12]See Proposition 4 of Foncel and Treich (2005).

Theorem 1 *If the Principal's fear of ruin is higher than the Agent's fear of ruin, then deterministic inspection is optimal. Formally if*

$$\frac{v(\theta,a)-v(\theta,\theta)}{v_a(\theta,a)} \geq \frac{u(\theta,a)}{u_a(\theta,a)},$$

for all $\theta, a \in [\underline{\theta},\overline{\theta}]^2$, then deterministic inspection is optimal.[13]

The proof of Theorem 1 is in Appendix.It seems intuitive that if the Principal has a greater fear of ruin than the Agent, it is less favorable for her to use stochastic inspection. However, employing stochastic inspection might reduce the Agent's information rent, potentially offsetting the risk the Principal assumes.

The debate on the sub-optimality of deterministic inspection was initiated by Townsend (1979). He provided an example demonstrating that stochastic inspection yields a higher payoff than deterministic inspection. However, many questions remain unanswered. For example, under what conditions is deterministic inspection optimal? To what extent is stochastic inspection superior to deterministic inspection?

Theorem 1 asserts that regardless of the inspection cost, if the Principal's fear of ruin exceeds that of the Agent, the optimal inspection strategy is deterministic. However, this does not imply that the optimal mechanism is unaffected by the inspection cost. In Section 3.1, we identify the optimal deterministic inspection policy.

Example 1 *(Linear utilities) Suppose $v(\theta,a)=\alpha(a-\theta)+b$ and $u(\theta,a)=a-\theta$, where $\alpha<0$, and $b>0$. Then Principal's fear of ruin is equal to the Agent's fear of ruin. Therefore the optimal inspection using Theorem 1 is deterministic.*

When utilities are linear in actions, they can be interpreted as transfers. The literature on CSV with transfers does not predict the optimality of deterministic inspection for several reasons. In Border and Sobel (1987) and Chander and Wilde (1998), the Principal can commit to an action (transfer) in the event of an inspection. Unlike our setting, the Agent's payoff is higher when inspected (truthful report) than when not inspected. Commitment to action in the case

[13]Note that a weaker condition similar to $\frac{v(\theta,a_*(\underline{\theta}))-v(\theta,\theta)}{v_a(\theta,a_*(\underline{\theta}))} \geq \frac{u(\theta^*,a_*(\underline{\theta}))}{u_a(\theta^*,a_*(\underline{\theta}))}$, is sufficient to show that deterministic inspection is optimal.

of inspection enhances the efficiency of stochastic inspection. The Principal can leverage rewards for truthful reporting.[14]

Example 2 *(Principal risk-averse) Suppose* $v(\theta, a) = \alpha(a-\theta)^{\frac{1}{2}} + b$ *and* $u(\theta, a) = a - \theta$*, where* $\alpha < 0$*, and* $b > 0$*. Then fear of ruin of the Principal is higher than the Agent. Therefore the optimal mechanism is deterministic.*

Harris and Raviv (1996) and Harris and Raviv (1998) study a principal-agent model in the context of capital budgeting and delegation. Agent's utility is linear but Principal has a concave utility. These papers deal with finite types (maximum three types) and predict optimality of stochastic inspection. Their argument depends on the discreteness of the type space and continuity of inspection probability.

Theorem 1 is silent when the Principal has less fear of ruin. In Mookherjee and Png (1989) and Melumad and Mookherjee (1989) the principal is risk neutral and the agent is risk-averse. The principal maximizes the agent's payoff subject to IR of the principal and IC of the Agent. Both papers predicts the optimality of stochastic inspection. Their argument depends on the discreteness of the type space and continuity of inspection probability.

Claims 1 and 2 in Appendix, given θ^* and $\tilde{\theta}$, provide the solution of problem $\mathbb{P}$ if local IC is sufficient condition for global IC and if the solution does not involve bunching in the interval $[\theta^*, \tilde{\theta}]$. Next Lemma provides a condition that local IC is sufficient for global IC.

Lemma 3 *Suppose* $u(\cdot,\cdot)$ *is log-supermodular: For all* $\theta \le a$*, and* (θ, a) *in* $[\underline{\theta}, \overline{\theta}]^2$

$$\frac{\partial^2 \ln\big(u(\theta, a)\big)}{\partial\theta\, \partial a} \ge 0.$$

Then, in any mechanism, (i) IC implies that $a(\theta)$ *and* $i(\theta)$ *are weakly increasing for all* $\theta \le \tilde{\theta}$*. (ii) Local IC and* $a(\cdot)$ *increasing, implies global IC.*

The proof of Lemma 3 is in Appendix. Log-Supermodularity closely resembles Supermodularity in the standard mechanism design literature. The reason for

[14] Palonen and Pekkarinen (2022) assumes no reward for the Agent following the inspection. Among other differences to our paper, their punishment is linear (exogenously fixed). They predict inspection probabilities take two values. This result is due to the linear punishment structure.

incorporating the logarithm is that, unlike transfers, the inspection probability multiplies the Agent's utility. Thus, a logarithmic transformation adjusts the Agent's (IC) to fit the standard mechanism design framework. However, this transformation is not applicable to the Principal's objective.

If the optimal mechanism, as derived using Claims 1 and 2 in the Appendix, does not involve bunching, then the Principal's optimal payoff will match the upper bound established in Lemma 2.

3.1. Deterministic Inspection

In this section, we derive the optimal mechanism, restricting the inspection policy to deterministic inspection. In order to state the optimal policy, we need to define two thresholds. Define the problem $\mathbb{P}_D$ as follows:

$$\mathbb{P}_D : \max_{\theta^* \in [\underline{\theta}, \overline{\theta}]} \left\{ \int_{\underline{\theta}}^{\theta^*} \Big(v(\theta, \theta^*) \Big) \, dF(\theta) + \int_{\theta^*}^{\overline{\theta}} \Big(v(\theta, \theta) - \phi \Big) \mathbb{1}_{v(\theta,\theta) \geq \phi} \, dF(\theta) \right\}.$$

Let Θ^* be the set of solutions of $\mathbb{P}_D$.[15] For $\theta^* \in \Theta^*$, define $\theta^{**} \geq \theta^*$ such that

$$\theta^{**} = \begin{cases} \theta^* & \text{if} \quad v(\theta^*, \theta^*) \leq \phi \\ \overline{\theta} & \text{if} \quad v(\overline{\theta}, \overline{\theta}) > \phi \end{cases}$$

Otherwise, define θ^{**} as the solution of $v(\theta^{**}, \theta^{**}) = \phi$. Using θ^* and θ^{**}, the following theorem states the optimal mechanism.

Theorem 2 *The optimal mechanism for all $\theta \in [\underline{\theta}, \overline{\theta}]$ is*

$$i_*(\theta) = \begin{cases} 0 & \theta < \theta^* \\ 1 & \theta^{**} \geq \theta \geq \theta^* \\ 0 & \theta > \theta^{**}, \end{cases}$$

$$a_*(\theta) = \begin{cases} \theta^* & \theta < \theta^* \\ \theta & \theta^{**} \geq \theta \geq \theta^* \\ \underline{\theta} & \theta > \theta^{**}. \end{cases}$$

The proof of Theorem 2 is in Appendix. First, Theorem 2 states that there are two thresholds. We call the first threshold (θ^*) the *inspection threshold*, and the second

[15] Problem $\mathbb{P}_D$ is continuous in $\theta^* \in [\underline{\theta}, \overline{\theta}]$; therefore, it admits a maximizer.

threshold (θ^{**}) the *exclusion threshold*. The optimal policy divides types into three areas: Efficient types ($\theta < \theta^*$), intermediate types ($\theta \in [\theta^*, \theta^{**}]$), and inefficient types ($\theta > \theta^{**}$). If the inspection cost is high ($v(\theta^*, \theta^*) \leq \phi$), then there is no inspection region. If the inspection cost is low ($v(\overline{\theta}, \overline{\theta}) > \phi$), there is no exclusion region.

Second, Theorem 2 expresses that the Principal mandates the efficient action after inspection, i.e., $a^I(\theta, \theta) = \theta$. The Principal does not give a reward for telling the truth in case of inspection. The reason is that, by excluding inefficient types and inspecting intermediate types, the Principal can decrease the informational rent (the set of types that can mimic) of efficient types and leave zero for intermediate types. Thus, intermediate types cannot mimic higher types, and the Principal can optimally set $a^I(\theta, \theta) = \theta$.

Third, the theorem says the Principal does not waste resources (cost of inspection) on inefficient types ($\theta > \theta^{**}$), so inspection is zero. Instead, she mandates a low action without inspection ($a(\cdot)$) and excludes the inefficient types. By inspecting intermediate types, the optimal policy hits two goals. First, it deters efficient types ($\theta < \theta^*$) from mimicking intermediate types. Second, by having full information on the intermediate types, she can mandate the efficient action ($a^I(\theta, \theta) = \theta$) for these types. Hence, commitment to $a^I(\cdot, \cdot)$ has no benefit for the Principal.

Fourth, Theorem 2 argues that the optimal mandated action without inspection ($a(\cdot)$) sets a fixed action equal to θ^* for efficient types. The mandated action for efficient types is fixed (the same as ours), but is strictly less than the efficient action for the lowest type that is inspected ($a < \theta^*$).[16] In our paper, if a for efficient types is less than θ^*, then types in the interval $[a, \theta^*]$ will reject the mandated action ex-post. This means that the Principal is excluding some efficient types. The aforementioned structure cannot be optimal, since the Principal can shift the inspection area toward more efficient types and instead exclude more inefficient types.

The trade-off at the inspection threshold is between inspecting the marginal type with the benefit of decreasing a for all efficient types. The trade-off at the exclusion threshold is between inspecting the marginal type with the benefit of mandating the efficient action, instead of excluding the marginal type.

Figure 2 illustrates the optimal mechanism. The Principal sets a cap on actions equal to θ^*. She inspects types between θ^* and θ^{**}, and mandates the efficient

[16] See Halac and Yared (2020), section IV, definition of TEC.

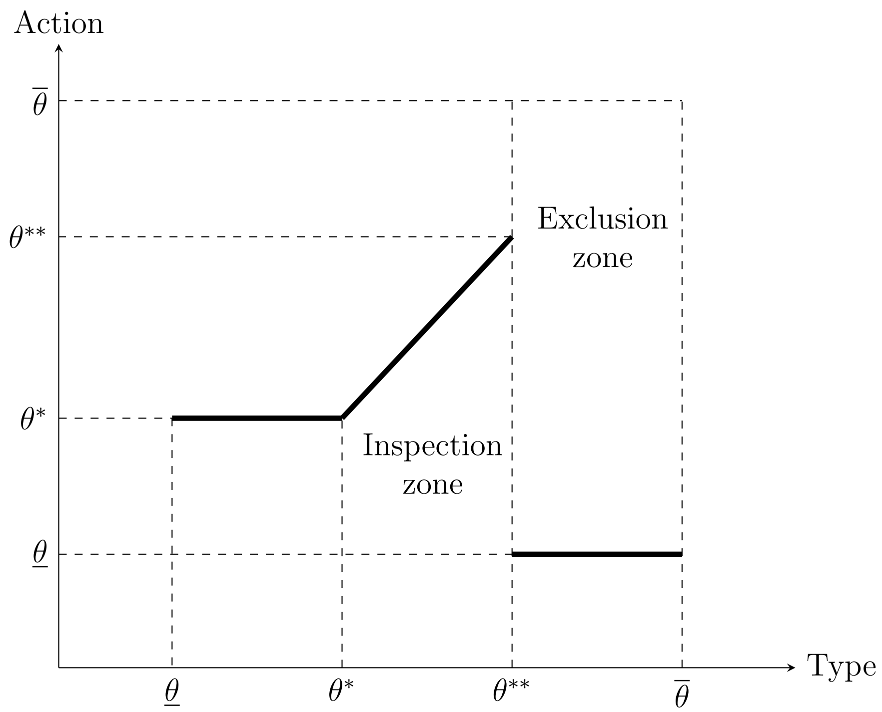

Figure 2: Optimal full commitment mechanism.

action for these types. The Principal excludes types above θ^{**} by mandating a very low action $a(\hat{\theta}) = \underline{\theta}$.

The threshold structure of the optimal policy does not come as a surprise due to deterministic inspection. A similar structure exists in Townsend (1979), Gale and Hellwig (1985), Malenko (2019) (the static case), and Halac and Yared (2020).

Besides the differences in the modeling assumptions of these papers with the current paper, the predictions of these models are different compared to ours. First, there is no exclusion region in these papers. The reason is either because the inspection cost is borne by the Agent (Townsend (1979) and Gale and Hellwig (1985)), or the first-best value of high types is increasing (Malenko (2019) and Halac and Yared (2020)), as opposed to ours, in which $v(\theta, \theta)$ is decreasing in θ.

Second, these papers assume ex-ante IR as opposed to an ex-post participation constraint in our paper.[17] Hence, inspection is used for two reasons.[18] Protecting

[17]Except Halac and Yared (2020), which assumes an ex-post participation constraint.

[18]Note that in Townsend (1979) and Gale and Hellwig (1985), the Agent pays the cost of inspection; in Malenko (2019) the Principal; and in Halac and Yared (2020) both.

inefficient types from excessive losses and reducing the rent of efficient types. Therefore, the inspection area is affected by the outside option (ex-ante) of the Agent, whereas in our model inspection does not act as insurance against excessive losses for the Agent, since the Agent can always reject the mandated action ex-post.

4. Different Commitment Ability

In this section, we examine the varying commitment capabilities of the Principal. There are two motivations for this analysis.

Firstly, it helps to identify which instruments are essential for the Principal to commit to. Understanding this is crucial for determining the effectiveness of the Principal's commitment ability across different tools.

Secondly, in practice, the Principal's commitment can vary depending on the application. For instance, in some cases, a government department specifies only audit policies without setting standards (actions in our context). In other cases, a regulator has an independent audit agency that does not specify when and under what conditions inspections are conducted. We view these different environments as representing different commitment abilities for the Principal.

By studying these variations, we can gain deeper insights into the strategic importance of each instrument and how the Principal can best utilize them in practice.

4.1. Commitment only to the action

From an ex post perspective, the Principal may lack the incentive to conduct inspections. When the Agent sends a message to the Principal, it might not be optimal for the Principal to inspect the Agent, making it challenging to ensure adherence to the contract. Additionally, if the Principal's inspection efforts are unobservable, she might reduce her effort to save on inspection costs. Especially when the optimal inspection policy is stochastic, it becomes difficult to monitor whether the Principal is adhering to the contract, and it is hard to enforce and verify deviations in court.

In this section, we examine a scenario where the Principal can commit to an action if no inspection occurs but does not commit to inspection policies. For example, consider an environmental regulator that sets standards (action in our setting) when there is no inspection. However, the regulator does not specify when inspections will occur, under what conditions they will take place, or what the mandated standards will be after an inspection.

Suppose the Principal can commit only to $a(\cdot) : \mathbb{M} \to \mathbb{R}_+$. Timing is as follows: The Principal commits to $a(\cdot)$. The Agent privately observes type $\theta \sim F(\cdot)$, and sends a message m. The Principal observes m, and decides whether to inspect or not: with inspection, she observes the true type and mandates a (without commitment); without inspection, she mandates $a(m)$ (with commitment). The Agent accepts or rejects the mandated standard.

The Agent's strategy is: $m(\theta) \in \mathbb{R}_+$. The Principal's strategy is $\big(i(m) \in [0,1], a^I(m,\theta) \in \mathbb{R}_+\big)$. The Principal's belief is: $\beta(\theta|m)$.

A related observation is that the Principal is opportunistic in the case of inspection: $a^I(m,\theta) = \theta$.

Theorem 3 *(i) The solution of the deterministic mechanism (Theorem 2) is a PBE of the game when the Principal can commit to the action. (ii) If $u(\cdot,\cdot)$ is log-supermodular, then the maximum ex-ante payoff of the Principal in the game when the Principal can commit to the action is equal to the optimal deterministic mechanism.*

The proof of Theorem 3 is in Appendix. Theorem 3 asserts that the Principal can guarantee a payoff equivalent to deterministic inspection by committing solely to $a(\cdot)$. The equilibrium strategies are as follows: types $[\underline{\theta}, \theta^*] \cup [\theta^{**}, \overline{\theta}]$ send message m_N, and types in $[\theta^*, \tilde{\theta}]$ send m_I. The Principal sets the cap θ^* on actions without inspection, i.e., $a(m) = \theta^*$, and inspects only when she receives message m_I, i.e., $i(m_I) = 1$. The Principal's off-the-equilibrium-path belief $\beta(\theta|m)$ is $\underline{\theta}$ with probability one.

We explain intuitively why the above strategies form an equilibrium. Types below θ^* do not deviate to m_I; otherwise, they receive a zero payoff. Moreover, types above θ^* receive a zero payoff anyhow.

The Principal adheres to her inspection strategy. If she does not inspect intermediate types (types that send m_I), then she will exclude types through the pre-committed action $a(m_I) = \theta^*$. By committing to this action, she effectively penalizes herself. She has committed to a low action when she does not inspect and receives m_I. The Principal does not inspect when she receives m_N. If inspecting m_N would yield a higher payoff, then inspecting both messages could generate an even higher payoff than optimal deterministic inspection (Theorem 2). However, this implies that in the optimal deterministic inspection, full inspection (inspecting all types) should be optimal. This cannot be the case unless the inspection cost, ϕ, is zero.

Note that Theorem 3 does not claim that there is no equilibrium yielding a higher payoff for the Principal than the optimal deterministic inspection. This might seem counterintuitive, as one might expect that without the ability to commit to an inspection policy, the Principal would be indifferent between inspecting and not inspecting, thus simplifying the policy to a deterministic one. However, this is not entirely accurate. While it is true that the Principal should be indifferent to inspection when the probability of inspection is between zero and one, stochastic inspection can influence the Agent's incentives and reduce the information rent. Theorem 3 (*ii*) states that if the Agent has log-supermodular utility, the Principal cannot use stochastic inspection to lower the Agent's information rent. In fact, the proof of Theorem 3 (*ii*) relies on the following monotonicity property of an equilibrium: if $m \in m(\theta)$ and $m' \in m(\theta')$, where $\tilde{\theta} > \theta' > \theta$, then $i(m') \geq i(m)$ and $a(m') \geq a(m)$. The monotonicity condition can be shown using the IC condition of the Agent for log-supermodular utilities.

Khalil (1997) examines a model of CSV where the Principal cannot commit to an inspection policy. Inspections occur ex post (after production), and the Principal cannot adjust allocations post-inspection, with penalties being exogenous. Therefore, his model is one of costly signaling rather than cheap talk. An ex post audit implies that if auditing is optimal, then inspections are stochastic and compliance is random (with the Agent playing a mixed strategy due to the presence of two types). In an audit contract, the Agent does not receive rent in either state of nature because the inefficient type is indifferent. This paper finds that there is more auditing under no commitment than under commitment. Conversely, our paper shows that the level of auditing remains the same whether there is no commitment (Section 4.1) or commitment to inspection (Section 3.1).

Melumad and Mookherjee (1989) also considers a setting with no commitment to inspection. The main differences are that the Principal can commit to transfers in case of inspection and the Agent cannot reject the mechanism ex post. In our model, the Principal can punish herself by mandating a very low action if she receives m_I and does not inspect. Then the Agent will reject. This helps the Principal to decrease the value of not inspecting without generating a big loss for the Agent. Therefore, in our model the Principal can achieve the full (deterministic) inspection payoff.

4.2. Commitment Only to Inspection

In this section, we examine a scenario where a Principal commits to an inspection policy but does not commit to actions. For instance, consider an inspection agency that outlines the timing and conditions under which inspections will occur and provides this information to a regulatory body. However, the regulatory body does not set specific standards for these inspections.

Suppose the Principal can only commit to $i(\cdot) : \mathbb{M} \to [0, 1]$. Timing is as follows: The Principal commits to $i(\cdot)$. The Agent privately observes type $\theta \sim F(\cdot)$, and sends a message m. The Principal inspects or not based on $i(m)$. With inspection, she observes the true type of the Agent and chooses an action a. Without inspection, she chooses a (possibly different) action a. The Agent accepts or rejects the mandated standard. The Agent's strategy is $m(\theta) \in \mathbb{R}_+$. The Principal's strategy is $\big(a(m) \in \mathbb{R}_+, a(m, \theta) \in \mathbb{R}_+\big)$. The Principal's belief is $\beta(\theta|m)$. Strategies and beliefs are measurable. A similar observation is that the Principal is opportunistic in the case of inspection: $a^I(m, \theta) = \theta$.

Lemma 4 *Fix $i(\cdot)$. In a PBE of the game when the Principal can only commit to inspection, (i) $a(m)$ is constant for all on-path messages m such that $i(m) < 1$. (ii) For on-path messages m_1 and m_2, if $i(m_1) < 1$ and $i(m_2) < 1$, then $i(m_1) = i(m_2)$.*

The proof of Lemma 4 is in Appendix. Lemma 4 asserts that the mandated action, in the absence of inspection, should be uniform across all messages. The reasoning is as follows: if there are two distinct mandated actions, types that are very close to the lower action will attempt to imitate those with the higher action. This phenomenon occurs regardless of the probability of inspection (as long as it is less than one). The underlying reason is that there will always be a type very close to an action; otherwise, the Principal can reduce that action.

As a result of Lemma 4, we can consolidate all messages with an inspection value less than 1 into a single message, m_N, and all messages with an inspection value of exactly 1 into another single message, m_I. Consequently, we can focus solely on equilibria with at most two messages on the equilibrium path. The following corollary formally states this observation.

Corollary 2 *For any equilibrium, there exists an outcome-equivalent equilibrium (same ex-ante payoff of the Principal and ex-ante payoffs of all types of the Agent) with at most two different messages,* m_I *and* m_N*.* $i(m_I) = 1,\ i(m_N) < 1$.

An interpretation of these two messages in Corollary 2 is that when the Agent sends m_I, he is requesting an inspection, whereas when he sends m_N, he is not requesting an inspection.

Theorem 4 *(i) The highest ex-ante payoff of the Principal is achieved by the following strategies. Agent: types in* $[\underline{\theta}, s^*] \cup [s^{**}, \overline{\theta}]$ *send* m_N *and types in* $[s^*, s^{**}]$ *send* m_I*. Principal:* $i(m_I) = 1,\ i(m_N) = 0$*, and* $a(m_N) = s^*$*. (ii)* $\theta^{**} \leq s^{**}$*. (iii) If* $\phi < v(\overline{\theta}, \overline{\theta})$ *(no exclusion region), then it is equivalent to the optimal deterministic inspection policy.*

The proof of Theorem 4 is in Appendix. Theorem 4 indicates that the equilibrium yielding the highest payoff shares a similar threshold structure with the optimal deterministic inspection policy (Theorem 2). However, the thresholds differ. The equilibrium structure does not involve stochastic inspection, suggesting that commitment to stochastic inspection is beneficial only if the Principal can also commit to an action ($a(\theta)$).

It is intuitive that the exclusion area is not in the middle of the type space. In other words, the Principal inspects any type θ' that is higher than a type θ that she excludes. Otherwise, replacing the messages sent by these types ($m(\theta) = m_I$ and $m(\theta') = m_N$) both increases the Principal's payoff and does not change the best reply of the Principal when she receives message m_N. The payoff increases since $v(\theta, \theta)$ is decreasing in θ. The best reply, $a(m_N)$, does not change since for all $a \leq a(m_N)$

$$\int_{\underline{\theta}}^{a} \left[v\left(\theta, a\right) \mathbb{1}_{a \geq \theta} \right] \beta(\theta | m_N)\, d\theta \leq \int_{\underline{\theta}}^{a(m_N)} \left[v\left(\theta, a(m_N)\right) \mathbb{1}_{a(m_N) \geq \theta} \right] \beta(\theta | m_N)\, d\theta.$$

The inequality remains valid because both sides are unchanged before and after the modification. This inequality holds for all $a \geq a(m_N)$ as well since the left side (weakly) decreases.

The theorem states that to incentivize the Principal to mandate sufficiently low actions when the Agent sends m_N, the Principal should commit to inspect inefficient types, i.e., $s^{**} \geq \theta^{**}$. The benefit from the exclusion area should be small enough that the Principal does not deviate in $a(m_N)$ and choose a higher action. Moreover, if the inspection cost is sufficiently low, meaning the exclusion area does not exist as per Theorem 2, then the highest payoff and equilibrium structure align with Theorem 2.

In the limited commitment section, Halac and Yared (2020) examines a delegation model where the Principal can commit to inspection but may alter the allowable action. This model shares some similarities with the one discussed in this section. However, Halac and Yared (2020) focuses on deterministic inspections, and $v(\theta, \theta)$ increases with θ. Consequently, there is no exclusion area.

4.3. No Commitment

Now we analyze the case in which the Principal does not have any commitment power. Timing is as follows: The Agent privately observes type $\theta \sim F(\cdot)$, and sends a message m. The Principal observes m and decides whether to inspect or not. With inspection, she observes the true type and chooses action a. Without inspection, she chooses a (possibly different) action a. The Agent accepts or rejects the mandated standard.

The Agent's strategy is: $m(\theta) \in \mathbb{R}_+$. The Principal's strategy is: $\big(i(m) \in [0,1], a(m) \in \mathbb{R}_+, a(m,\theta) \in \mathbb{R}_+\big)$. The Principal's belief is: $\beta(\theta|m)$.

By Lemma 4, since the Principal cannot commit to action without inspection, one can construct all equilibrium outcomes with at most two messages. We focus on interval strategies: the Agent's strategy $m(\theta)$ alternates between m_I and m_N in finite intervals.

Assumption 4 *For all $(\theta, \theta', s) \in [\underline{\theta}, \overline{\theta}]^3$, such that $\theta' \leq \theta \leq s$,*

$$v(\theta, s) - v(\theta, \theta) \geq v(\theta', s) - v(\theta', \theta').$$

It is easy to see that Assumption 4 is a weaker condition than supermodularity.

Theorem 5 *Suppose the Principal's utility satisfies Assumption 4. Then the equilibrium with the highest ex-ante payoff has a structure similar to the optimal deterministic inspection policy (with different thresholds).*

The proof of Theorem 5 is in Appendix. Banks (1989) examines a model in which an agent requests a budget. In the open procedure section, the principal may choose to inspect or not and then decides on a budget for the agent accordingly. Utilities are linear, and inspection is imperfect; there is a probability that the inspection does not yield any information. Imperfect inspection incentivizes types just below but very close to s^* to send m_I; therefore, (full) pooling is the only equilibrium. In any equilibrium, no information is transmitted from the agent to the principal.

In contrast, our paper focuses on finding the highest payoff for the principal, and with perfect inspection there are equilibria like Theorem 5 that reveal more information to the Principal. No new information will be transmitted from the agent to the principal. In contrast, our paper focuses on finding the highest payoff for the principal, and with perfect inspection there are equilibria like Theorem 5 that reveal more information to the Principal. Then the nonexistence of semi-separating equilibria in Banks (1989) is due to imperfect inspection. However, this is a knife-edge case. Even with imperfect inspection, the Principal can incentivize types below s^* not to send m_I. If a type below s^* sends m_I and if the inspection is successful, then the agent should pay a small part of the inspection cost.

4.4. Partial Commitment

Suppose the Principal is obligated to conduct an inspection only if the Agent requests it. In numerous applications, there are institutions in place to ensure that principals cannot refuse an audit once it has been requested. Thus, the agent always has the option to ask for an audit. We demonstrate that with this minimal level of commitment, and if the inspection cost is not high, then the highest equilibrium payoff matches the payoff described in Theorem 2.

Theorem 6 *(Partial commitment) Suppose the inspection cost is not high, i.e., $\phi < v(\overline{\theta}, \overline{\theta})$, and the Principal can commit to inspect message m_I. There exists an equilibrium such that the ex-ante payoff of the Principal is the same as in Theorem 2.*

The proof of Theorem 6 is in Appendix. The intuition behind the above theorem is as follows. Since the Principal cannot commit to actions, by employing Lemma 4 and Theorem 4, we can focus on equilibria with deterministic inspection and two messages, $\{m_I, m_N\}$. Suppose types in $[\theta^*, \overline{\theta}]$ send message m_I, while other types send m_N. The Principal has no incentive to inspect m_N. If inspecting m_N would yield a higher payoff, then inspecting both messages could generate an even higher payoff than the optimal deterministic inspection. The Principal optimally chooses action $a(m_N) = \theta^*$ after observing m_N. The reasoning is as follows: first, $a(m_N) \leq \theta^*$, since the Principal assigns zero probability to types above θ^* after observing m_N. Second, if $a(m_N) < \theta^*$, then a policy that caps actions at $a(m_N)$ and inspects types in the interval $[\theta^*, \overline{\theta}]$ yields a strictly higher payoff than the optimal policy in Problem $\mathbb{P}_D$, which is not possible.

5. Conclusion

In conclusion, this paper provides a comprehensive analysis of the regulatory tools of mandating standards and conducting inspections within a principal–agent framework. We have examined how different levels of commitment by the regulator affect the effectiveness of these tools. Our findings indicate that when the regulator's fear of ruin exceeds the firm's fear of ruin, a deterministic inspection policy proves to be optimal. This policy involves a structured approach where types are segmented into efficient, intermediate, and inefficient categories. This segmentation helps balance the trade-off between encouraging technology adoption and managing the risk of a firm rejecting standards.

Moreover, our analysis of varying commitment levels reveals that while full commitment to both actions and inspections yields higher payoffs, partial-commitment scenarios also show promising outcomes. When the cost of inspections is not high and the regulator commits to inspecting upon request, the optimal deterministic policy can still be achieved. These insights highlight the importance of strategic commitment and the potential for tailored regulatory approaches to effectively incentivize technology adoption. Overall, our study underscores the need for regulators to carefully design and commit to their policies to maximize their effectiveness in achieving desired environmental and performance standards.

References

Amirreza Ahmadzadeh and Stephan Waizmann. Mechanism design with costly inspection. *Available at SSRN 4823817*, 2024.

Robert J Aumann and Mordecai Kurz. Power and taxes. *Econometrica: Journal of the Econometric Society*, pages 1137–1161, 1977.

Jeffrey S Banks. Agency budgets, cost information, and auditing. *American journal of political science*, pages 670–699, 1989.

David P Baron and David Besanko. Regulation, asymmetric information, and auditing. *The RAND Journal of Economics*, pages 447–470, 1984.

David P Baron and Roger B Myerson. Regulating a monopolist with unknown costs. *Econometrica: Journal of the Econometric Society*, pages 911–930, 1982.

Elchanan Ben-Porath, Eddie Dekel, and Barton L Lipman. Optimal allocation with costly verification. *American Economic Review*, 104(12):3779–3813, 2014.

Kim C Border and Joel Sobel. Samurai accountant: A theory of auditing and plunder. *The Review of economic studies*, 54(4):525–540, 1987.

Parkash Chander and Louis L Wilde. A general characterization of optimal income tax enforcement. *The Review of Economic Studies*, 65(1):165–183, 1998.

Vincent P Crawford and Joel Sobel. Strategic information transmission. *Econometrica: Journal of the Econometric Society*, pages 1431–1451, 1982.

Albin Erlanson and Andreas Kleiner. Costly verification in collective decisions. *Theoretical Economics*, 15(3):923–954, 2020.

Jerome Foncel and Nicolas Treich. Fear of ruin. *Journal of Risk and Uncertainty*, 31:289–300, 2005.

Douglas Gale and Martin Hellwig. Incentive-compatible debt contracts: The one-period problem. *The Review of Economic Studies*, 52(4):647–663, 1985.

Marina Halac and Pierre Yared. Commitment versus flexibility with costly verification. *Journal of Political Economy*, 128(12):4523–4573, 2020.

Milton Harris and Artur Raviv. The capital budgeting process: Incentives and information. *The Journal of Finance*, 51(4):1139–1174, 1996.

Milton Harris and Artur Raviv. Capital budgeting and delegation. *Journal of Financial Economics*, 50(3):259–289, 1998.

Morton I Kamien and Nancy Lou Schwartz. *Dynamic optimization: the calculus of variations and optimal control in economics and management*. courier corporation, 2012.

Fahad Khalil. Auditing without commitment. *The RAND Journal of Economics*, pages 629–640, 1997.

Yunan Li. Mechanism design with costly verification and limited punishments. *Journal of Economic Theory*, 186:105000, 2020.

Andrey Malenko. Optimal dynamic capital budgeting. *The Review of Economic Studies*, 86(4):1747–1778, 2019.

Nahum D Melumad and Dilip Mookherjee. Delegation as commitment: the case of income tax audits. *The RAND Journal of Economics*, pages 139–163, 1989.

Dilip Mookherjee and Ivan Png. Optimal auditing, insurance, and redistribution. *The Quarterly Journal of Economics*, 104(2):399–415, 1989.

Tymofiy Mylovanov and Andriy Zapechelnyuk. Optimal allocation with ex post verification and limited penalties. *American Economic Review*, 107(9):2666–94, 2017.

Petteri Palonen and Teemu Pekkarinen. Regulating a monopolist with enforcement. *Available at SSRN 3729347*, 2022.

Rohit Patel and Can Urgun. Costly verification and money burning. 2022.

Atle Seierstad and Knut Sydsaeter. *Optimal control theory with economic applications*. Elsevier North-Holland, Inc., 1986.

SD Solomon. Contribution of working group i to the fourth assessment report of the intergovernmental panel on climate change 2007. *(No Title)*, 2007.

Robert M Townsend. Optimal contracts and competitive markets with costly state verification. *Journal of Economic theory*, 21(2):265–293, 1979.

A. Appendix

A.1. Full Commitment

Fix $\alpha > 0$. For all $n \in \mathbb{N}$, define mechanism $\mathbb{M}_n(\alpha)$ as follows: $a^I(\theta,\theta) = n$, $i(\theta) = \frac{-\alpha}{v(\theta,n)}$, $a(\theta) = \theta$, and $a^I(\hat{\theta},\theta) = \underline{\theta}$ if $\hat{\theta} \neq \theta$ for all $(\theta,\hat{\theta}) \in [\underline{\theta},\overline{\theta}]^2$.

Lemma 5 *Let* $\overline{u} = \sup_{(\hat{\theta},\theta)\in[\underline{\theta},\overline{\theta}]^2} u(\theta,\hat{\theta})$, *and* $\alpha > 0$. *Suppose* $\lim_{a\to\infty} v(\theta,a) = -\infty$, *and* $\lim_{a\to\infty} -\frac{u(\theta,a)}{v(\theta,a)} > \frac{\overline{u}}{\alpha}$, *for all* $\theta \in [\underline{\theta},\overline{\theta}]$. *Then the payoff of mechanisms* $\mathbb{M}_n(\alpha)$ *converges to* $\mathbb{E}\big[v(\theta,\theta)\big] - \alpha$, *as* n *goes large.*

Proof. For n large enough the mechanism $\mathbb{M}_n$ is incentive compatible, i.e,

$$
\begin{aligned}
&(1-i(\theta))\Big(u(\theta,a(\theta))\Big)\mathbb{1}_{a(\theta)\geq\theta} + i(\theta)\Big(u(\theta,a^I(\theta,\theta))\Big)\mathbb{1}_{a^I(\theta,\theta)\geq\theta} \\
&= \frac{-\alpha}{v(\theta,n)}u(\theta,n) \geq \sup_{\hat{\theta}}(1-i(\hat{\theta}))\Big(u(\theta,a(\hat{\theta}))\Big)\mathbb{1}_{a(\hat{\theta})\geq\theta} \\
&= \sup_{\hat{\theta}}(1-i(\hat{\theta}))u(\theta,\hat{\theta}).
\end{aligned}
$$

The reason is $\lim_{n\to\infty} \frac{-\alpha}{v(\theta,n)}u(\theta,n) > \overline{u}$, and $\sup_{\hat{\theta}}(1-i(\hat{\theta}))u(\theta,\hat{\theta})$ is weakly less than $\overline{u}$. Note that all types accept $\mathbb{M}_n$ ex-post. Now we compute the payoff of $\mathbb{M}_n$ for the Principal.

$$
\begin{aligned}
&\mathbb{E}\Big[(1-i(\theta))\Big(v(\theta,a(\theta))\Big)\mathbb{1}_{a(\theta)\geq\theta} + i(\theta)\Big(-\phi + v(\theta,a^I(\theta,\theta))\mathbb{1}_{a^I(\theta,\theta)\geq\theta}\Big)\Big] \\
&= \mathbb{E}\Big[(1-\frac{-\alpha}{v(\theta,n)})\Big(v(\theta,\theta)\Big) + \frac{-\alpha}{v(\theta,n)}\Big(-\phi + v(\theta,n)\Big)\Big]
\end{aligned}
$$

Fixing α, the limit of the above payoff as n goes large is $\mathbb{E}\big[v(\theta,\theta) - \alpha\big]$. ■

As a result if $\lim_{a\to\infty} -\frac{u(\theta,a)}{v(\theta,a)} = \infty$, then $\lim_{a\to\infty} -\frac{u(\theta,a)}{v(\theta,a)} > \frac{\overline{u}}{\alpha}$ for all $\alpha > 0$. Therefore mechanism $\mathbb{M}_n(\alpha)$ approximates first-best as n goes large and α goes small.

Lemma 6 *Suppose* $\lim_{a\to\infty} v(\theta,a) = -\infty$, *and* $\lim_{a\to\infty} -\frac{u(\theta,a)}{v(\theta,a)} = \infty$, *for all* $\theta \in [\underline{\theta},\overline{\theta}]$. *Then the solution to problem* $\mathbb{P}$ *does not exist.*

Proof. Assume it exits, therefore the payoff of the Principal should be $\mathbb{E}\big(v(\theta,\theta)\big)$. Based on the objective of the principal in problem $\mathbb{P}$, if the payoff is $\mathbb{E}\big(v(\theta,\theta)\big)$, then $i(\theta) \overset{\text{a.e}}{=} 0$, $a(\theta) \overset{\text{a.e}}{=} 0$. This mechanism clearly cannot satisfy the truth-telling condition. ■

A.2. Proof of Lemma 1

Proof.

(i) By contradiction let $\theta_2 > \theta_1$ and $\pi(\theta_2) > \pi(\theta_1) = 0$. Then $i(\theta_2) < 1$ and $a(\theta_2) > \theta_2$. But θ_1 can mimic θ_2 and receive a strictly positive payoff. A contradiction.

(ii) When $a(\theta) < \theta$, it means that the Principal excludes type θ, if she does not inspect. In this case, type θ's expected payoff does not change if the principal decreases $a(\theta)$ to $\underline{\theta}$.

(iii) Let $a(\theta) = \underline{\theta}$. Then for all types $\theta^{'} > \theta$, we have $a(\theta^{'}) \leq \theta^{'}$. If $a(\theta^{'}) < \theta^{'}$, then we can assume $a(\theta^{'}) = \underline{\theta}$. If $a(\theta^{'}) = \theta^{'}$, then $i(\theta^{'}) = 1$, otherwise type θ, can mimic $\theta^{'}$. If $i(\theta^{'}) = 1$, the Principal can strictly gain by switching the policies for type θ, and $\theta^{'}$. Excluding type θ, while keeping $\theta^{'}$ is not efficient; i.e $v(\theta, \theta) - \phi > v(\theta^{'}, \theta^{'}) - \phi$.

(iv) If $a(\theta) \leq \theta$, then either $a(\theta^{'}) \leq \theta^{'}$ or $i(\theta^{'}) = 1$ for all $\theta^{'} > \theta$. If $a(\theta^{'}) < \theta^{'}$, we can assume $a(\theta^{'}) = \underline{\theta}$. If $a(\theta^{'}) = \theta^{'}$, then $i(\theta^{'}) = 1$, otherwise θ can mimic $\theta^{'}$.

■

A.3. Proof of Lemma 2

Proof. For ease of notation we drop $*$. For $\theta \leq \tilde{\tilde{\theta}}$, $a(\theta) \geq \theta$. Therefore the optimization becomes

$$\max_{i(.), a(.), \tilde{\tilde{\theta}}} \int_{\underline{\theta}}^{\tilde{\tilde{\theta}}} \Big[(1 - i(\theta))\Big(v(\theta, a(\theta)) - v(\theta, \theta) + \phi\Big) + \Big(v(\theta, \theta) - \phi\Big)\Big] f(\theta) d\theta,$$

subject to the IC condition. By Lemma 1, there exists $\tilde{\theta}$ such that for all $\theta < \tilde{\theta}$, $a(\theta) > \theta$, $a(\tilde{\theta}) = a(\tilde{\theta})$. For types $\theta > \tilde{\theta}$, $a(\cdot)$ is irrelevant. Therefore the optimization becomes

$$\max_{i(.), a(.), \tilde{\theta}, \tilde{\tilde{\theta}}} \int_{\underline{\theta}}^{\tilde{\theta}} \Big[(1 - i(\theta))\Big(v(\theta, a(\theta)) - v(\theta, \theta) + \phi\Big)\Big] f(\theta) d\theta + \int_{\underline{\theta}}^{\tilde{\tilde{\theta}}} \Big[\Big(v(\theta, \theta) - \phi\Big)\Big] f(\theta) d\theta$$

subject to

$$(1-i(\theta))u(\theta,a(\theta)) \geq\ (1-i(\hat{\theta}))u(\theta,a(\hat{\theta})),$$

$a(\theta) \geq \theta$ and $\tilde{\theta} \leq \tilde{\tilde{\theta}}$. Fixing $\tilde{\theta}$ and $\tilde{\tilde{\theta}}$, we relax the global IC constraint by local IC as follows. Given a type θ and $\hat{\theta} \in (\theta, a(\theta))$. Therefore $a(\hat{\theta}) > \theta$. The local IC condition is

$$(1-i(\theta))u(\theta,a(\theta)) \geq\ (1-i(\hat{\theta}))u(\theta,a(\hat{\theta})).$$

By employing the Envelope Theorem, the local truth telling (except those points that $a(\cdot)$ has a discontinuity) condition gives

$$\frac{\partial \ln\big[\pi(\theta)\big]}{\partial\theta} = \frac{\partial ln\big[u(\theta,a(\theta))\big]}{\partial\theta},$$

or equivalently for $\pi(\theta) > 0$ ($\theta < \tilde{\theta}$),

$$\dot{\pi}(\theta) = \pi(\theta)\frac{u_\theta(\theta,a(\theta))}{u(\theta,a(\theta))}.$$

The relaxed problem of $\mathbb{P}$ by replacing global IC with local IC is

$$\max_{\pi(\cdot),a(\cdot)} \int_{\underline{\theta}}^{\tilde{\theta}} \Big[\frac{\pi(\theta)}{u(\theta,a(\theta))}\Big(v(\theta,a(\theta)) - v(\theta,\theta) + \phi\Big)\Big] f(\theta)d\theta + \int_{\underline{\theta}}^{\tilde{\tilde{\theta}}} \Big(v(\theta,\theta) - \phi\Big) f(\theta)d\theta,$$

subject to (for $\theta < \tilde{\theta}$)

$$\dot{\pi}(\theta) = \pi(\theta)\frac{u_\theta(\theta,a(\theta))}{u(\theta,a(\theta))},$$
$$\pi(\theta) \leq u(\theta,a(\theta)).$$

Note that the state equation does not need to be valid at discontinuity points of $a(\cdot)$. We analysis the relaxed problem by using the Pontryagin's maximum principle (π is the state and a is the control variable).[19] The Hamiltonian for $\theta < \tilde{\theta}$ is

$$H(a,\pi,\mu,w,\theta) = \frac{\pi}{u(\theta,a)}\Big(v(\theta,a) - v(\theta,\theta) + \phi\Big) f(\theta) + \mu\,\pi\frac{u_\theta(\theta,a)}{u(\theta,a)} + w(u(\theta,a) - \pi),$$

where the Lagrangian multiplier for $\pi(\theta) \leq u(\theta,a(\theta))$ is $w(\theta)$. From the Pontryagin principle for the co-state variable $\mu(\theta)$ we have (except at discontinuity points of

[19]Note that we can write the Pontryagin's maximum principle for piece-wise continuous control, and piece-wise differentiable state. For more information see Seierstad and Sydsaeter (1986), Chapter 2, Theorem 2, or Kamien and Schwartz (2012), section 14.

$a(\cdot))$

$$\dot{\mu}(\theta) = -\frac{\partial H}{\partial \pi} = -\frac{1}{u(\theta, a(\theta))}\Big(v(\theta, a(\theta)) - v(\theta, \theta) + \phi\Big) f(\theta) - \mu(\theta)\ \frac{u_\theta(\theta, a(\theta))}{u(\theta, a(\theta))} + w(\theta).$$

Since $u(\theta, a(\theta)) > 0$, then

$$\dot{\mu}(\theta) u(\theta, a(\theta)) + \mu(\theta)\ u_\theta(\theta, a(\theta)) - w(\theta) u(\theta, a(\theta)) = -\Big(v(\theta, a(\theta)) - v(\theta, \theta) + \phi\Big) f(\theta). \tag{1}$$

For types $\theta \in [\theta^*, \tilde{\theta}]$, $w(\theta) = 0$. Hence

$$\frac{d\ \mu(\theta)\pi(\theta)}{d\ \theta} = -\frac{\pi(\theta)}{u(\theta, a(\theta))}\Big(v(\theta, a(\theta)) - v(\theta, \theta) + \phi\Big) f(\theta).$$

Employing Newton–Leibniz theorem, and the fact that $\mu(\cdot)$, and $\pi(\cdot)$ are continuous functions:

$$\int_{\theta^*}^{\tilde{\theta}} \frac{d\ \mu(\theta)\pi(\theta)}{d\ \theta}\ d\theta = \int_{\theta^*}^{\tilde{\theta}} -\frac{\pi(\theta)}{u(\theta, a(\theta))}\Big(v(\theta, a(\theta)) - v(\theta, \theta) + \phi\Big) f(\theta)\ d\theta.$$

Since $\pi(\tilde{\theta}) = 0$,

$$\mu(\theta^*)\pi(\theta^*) = \int_{\theta^*}^{\tilde{\theta}} \frac{\pi(\theta)}{u(\theta, a(\theta))}\Big(v(\theta, a(\theta)) - v(\theta, \theta) + \phi\Big) f(\theta)\ d\theta. \tag{2}$$

In the next step we compute $\mu(\theta^*)$. From the Pontryagin's maximum principle, we know for all $\theta < \tilde{\theta}$, a maximizes the following

$$\frac{\pi(\theta)}{u(\theta, a)}\Big(v(\theta, a) - v(\theta, \theta) + \phi\Big) f(\theta) + \mu(\theta)\ \pi(\theta)\frac{u_\theta(\theta, a)}{u(\theta, a)} + w(\theta)(u(\theta, a) - \pi(\theta)).$$

Sine the above is a C^1 function of a, as a necessary condition we can write

$$\begin{aligned} H_a = \pi(\theta)\Bigg(&\frac{u(\theta, a(\theta))\big(v_a(\theta, a(\theta)) f(\theta) + \mu(\theta) u_{\theta a}(\theta, a(\theta))\big)}{u^2(\theta, a(\theta))} \\ &- \frac{u_a(\theta, a(\theta))\Big((v(\theta, a(\theta)) - v(\theta, \theta) + \phi) f(\theta) + \mu(\theta)\ u_\theta(\theta, a(\theta))\Big)}{u^2(\theta, a(\theta))}\Bigg) \\ &+ w(\theta) u_a(\theta, a(\theta)) = 0. \end{aligned} \tag{3}$$

Pluging 1, we get

$$H_a = \pi(\theta)u(\theta,a(\theta))\Big(\frac{v_a(\theta,a(\theta))f(\theta)+\mu(\theta)u_{\theta a}(\theta,a(\theta))+u_a(\theta,a(\theta))\big(\dot{\mu}(\theta)-w(\theta)\big)}{u^2(\theta,a(\theta))}\Big) + w(\theta)u_a(\theta,a(\theta)) = 0.$$

Simple algebra gives us

$$\pi(\theta)u(\theta,a(\theta))\Big(\frac{v_a(\theta,a(\theta))f(\theta)+\mu(\theta)u_{\theta a}(\theta,a(\theta))+u_a(\theta,a(\theta))\dot{\mu}(\theta)}{u^2(\theta,a(\theta))}\Big) = u_a(\theta,a(\theta))w(\theta)\big(\frac{\pi(\theta)}{u(\theta,a(\theta))}-1\big)=0.$$

The last equality is due to the fact that $w(\theta)(u(\theta,a(\theta))-\pi(\theta))=0$. Finally since $\pi(\theta)$, and $u(\theta,a(\theta))$ are strictly positive we have

$$\mu(\theta)u_{\theta a}(\theta,a(\theta))+u_a(\theta,a(\theta))\dot{\mu}(\theta) = -v_a(\theta,a(\theta))f(\theta).$$

The above equation for $\theta \leq \theta^*$ implies

$$\int_{\underline{\theta}}^{\theta} \frac{d\;\mu(t)u_a(t,a(t))}{d\;t}\;dt = \int_{\underline{\theta}}^{\theta} -v_a(t,a_*(t))f(t)\;dt.$$

The transversality condition at $\underline{\theta}$ is $\mu(\underline{\theta})=0$. Hence

$$\mu(\theta)u_a(\theta,a(\underline{\theta})) = \int_{\underline{\theta}}^{\theta} -v_a(t,a(\underline{\theta}))f(t)\;dt. \tag{4}$$

Employing equations 2 and 4, and the fact that $\pi(\theta^*)=u(\theta^*,a(\underline{\theta}))$, the Principal's payoff becomes

$$\int_{\underline{\theta}}^{\theta^*}\Big(v(\theta,a(\underline{\theta}))-v(\theta,\theta)+\phi\Big)f(\theta)\;d\theta + \frac{u(\theta^*,a(\underline{\theta}))}{u_a(\theta^*,a(\underline{\theta}))}\int_{\underline{\theta}}^{\theta^*} -v_a(\theta,a(\underline{\theta}))f(\theta)\;d\theta + \int_{\underline{\theta}}^{\tilde{\tilde{\theta}}}\Big(v(\theta,\theta)-\phi\Big)f(\theta)d\theta.$$

Equivalently

$$\int_{\underline{\theta}}^{\theta^*} \Big(v(\theta, a(\underline{\theta})) - \frac{u(\theta^*, a(\underline{\theta}))}{u_a(\theta^*, a(\underline{\theta}))} v_a(\theta, a(\underline{\theta}))\Big) f(\theta)\, d\theta + \int_{\theta^*}^{\tilde{\tilde{\theta}}} \Big(v(\theta,\theta) - \phi\Big) f(\theta) d\theta.$$

■

Claim 1 *Suppose local IC is sufficient for global IC. Assume $a_*(\theta)$ is strictly increasing in $[\theta^*, \tilde{\theta}]$, then there exists a C^1 function $\mu(\theta)$ such that $a_*(\theta)$ satisfy the below equations:*

$$-\dot{\mu}(\theta) = \frac{v(\theta, a_*(\theta)) - v(\theta,\theta) + \phi}{u(\theta, a_*(\theta))} f(\theta) + \mu(\theta) \frac{u_\theta(\theta, a_*(\theta))}{u(\theta, a_*(\theta))},$$

$$\frac{\partial\Big(\frac{v(\theta,a_*(\theta)) - v(\theta,\theta)+\phi}{u(\theta,a_*(\theta))}\Big)}{\partial a} f(\theta) + \mu(\theta) \frac{\partial^2 \ln\big(u(\theta, a_*(\theta))\big)}{\partial\theta\, \partial a} = 0,$$

for all $\theta \in [\theta^*, \tilde{\theta}]$. Moreover, if $a_*(\cdot)$ is piecewise C^1 then $a_*(\cdot)$ satisfies (at all differentiable and continuous points of a^*) the below first order differential equation

$$\dot{a}(\theta) = \Big[\psi(\theta, a_*(\theta)) - \frac{\partial\Psi(\theta, a_*(\theta))}{\partial\theta}\Big]\Big(\frac{\partial\Psi(\theta, a_*(\theta))}{\partial a}\Big)^{-1},$$

with end point condition $a_*(\tilde{\theta}) = \tilde{\theta}$.

Proof. For all $\theta \in [\theta^*, \tilde{\theta}]$, from equation 3:

$$\mu(\theta) = \frac{-\frac{\partial\Big(\frac{v(\theta,a_*(\theta)) - v(\theta,\theta)+\phi}{u(\theta,a_*(\theta))}\Big)}{\partial a} f(\theta)}{\frac{\partial^2 \ln\big(u(\theta,a_*(\theta))\big)}{\partial\theta\,\partial a}} = \Psi(\theta, a_*(\theta)),$$

and from equation 1:

$$-\dot{\mu}(\theta) = \frac{v(\theta, a_*(\theta)) - v(\theta,\theta) + \phi}{u(\theta, a_*(\theta))} f(\theta) + \mu(\theta) \frac{u_\theta(\theta, a_*(\theta))}{u(\theta, a_*(\theta))}.$$

Therefore

$$-\dot{\mu}(\theta) = \frac{\big(u_\theta(\theta, a_*(\theta)) v_a(\theta, a_*(\theta)) - u_{a\theta}(\theta, a_*(\theta))(v(\theta, a_*(\theta)) - v(\theta,\theta) + \phi)\big) f(\theta)}{-u_a(\theta, a_*(\theta)) u_\theta(\theta, a_*(\theta)) + u(\theta, a_*(\theta)) u_{\theta a}(\theta, a_*(\theta))}$$
$$\equiv \psi(\theta, a_*(\theta)).$$

Finally a simple calculation gives us the differential equation. ■

Claim 2 *(The optimal inspection) Let $a_*(.)$ be the optimal mandated action in case of no inspection. Then the expected payoff for the Agent, and the optimal inspection policy for $\theta \leq \tilde{\theta}$, are*

$$\pi(\theta) = \pi(\underline{\theta})\ exp\Big(\int_{\underline{\theta}}^{\theta} \frac{u_\theta(t, a_*(t))}{u(t, a_*(t))}\ dt\Big),$$

and

$$i_*(\theta) = 1 - \frac{\pi(\underline{\theta})\ exp\Big(\int_{\underline{\theta}}^{\theta} \frac{u_\theta(t,a_*(t))}{u(t,a_*(t))}\ dt\Big)}{u(\theta, a_*(\theta))}.$$

Proof.

Using Newton–Leibniz theorem, and the fact that $\pi(.)$ is a continuous function:

$$ln(\pi(\theta)) - ln(\pi(\underline{\theta})) = \int_{\underline{\theta}}^{\theta} \frac{d\ ln(\pi(t))}{d\ t}\ dt = \int_{\underline{\theta}}^{\theta} u_\theta(t, a(t))dt.$$

Note that the above equality is valid even if $i_*(.)$ or $a_*(.)$ has discontinuities. Finally

$$\pi(\theta) = \pi(\underline{\theta})\ exp\Big(\int_{\underline{\theta}}^{\theta} \frac{u_\theta(t, a_*(t))}{u(t, a_*(t))}\ dt\Big).$$

We know $\pi(\theta) = (1 - i_*(\theta))u(\theta, a_*(\theta))$, so $i_*(\theta) = 1 - \frac{\pi(\theta)}{u(\theta, a_*(\theta))}$. ■

A.4. Proof of Lemma 3

Proof. (i) For all $\theta < \tilde{\theta}$, $i(\theta) < 1$ and $a(\theta) > \theta$. First we show $a(\cdot)$ is weakly increasing. For all θ, and $\hat{\theta} \in (\theta, a(\theta))$, IC conditions imply

$$\ln(1 - i(\theta)) + \ln\big(u(\theta, a(\theta))\big) \geq\ \ln(1 - i(\hat{\theta})) + \ln\big(u(\theta, a(\hat{\theta}))\big),$$

and

$$\ln(1 - i(\hat{\theta})) + \ln\big(u(\hat{\theta}, a(\hat{\theta}))\big) \geq\ \ln(1 - i(\theta)) + \ln\big(u(\hat{\theta}, a(\theta))\big).$$

Assume $\theta > \hat{\theta}$. Adding up the above inequalities gives us

$$\frac{\ln\big(u(\theta, a(\theta))\big) - \ln\big(u(\hat{\theta}, a(\theta))\big)}{\theta - \hat{\theta}} \geq \frac{\ln\big(u(\theta, a(\hat{\theta}))\big) - \ln\big(u(\hat{\theta}, a(\hat{\theta}))\big)}{\theta - \hat{\theta}}.$$

Using log-Supermodular property of $u(\cdot,\cdot)$, we conclude $a(\theta) \geq a(\hat{\theta})$. Now we show $i(\cdot)$ is weakly increasing. By contradiction assume there exist types $\theta < \theta^{'}$ such that $i(\theta) > i(\theta^{'})$. Since $a(\cdot)$ is a weakly increasing function, then type θ prefers to mimic type $\theta^{'}$, i.e.

$$(1 - i(\theta))u(\theta, a(\theta)) \leq (1 - i(\theta^{'}))u(\theta, a(\theta^{'})).$$

A contradiction.

(ii) Define $y(\theta) = ln(\pi(\theta))$, and $u(\theta, a(\hat{\theta})) = ln(u(\theta, a(\hat{\theta})))$. A logarithm transformation of the local truth-telling condition for all t gives us

$$\dot{y}(t) = u_{\theta}(t, a(t)),$$

for all t such that $a(\cdot)$ is continuous at t. At discontinuity points of $i(.)$, and $a(.)$ we should replace the derivatives with left or right derivatives, so

$$\dot{y}(t^{-}) = u_{\theta}(t, a(t^{-})), \text{ and } \dot{y}(t^{+}) = u_{\theta}(t, a(t^{+})).$$

Let $t > t^{'}$. Using Newton–Leibniz theorem, and the fact that $y(.)$ is a continuous function:

$$y(t) - y(t^{'}) = \int_{t^{'}}^{t} u_{\theta}(s, a(s))ds, \tag{5}$$

We need to show for all t, and $t^{'}$.

$$\pi(t) \geq \pi(t^{'})\frac{u(t, a(t^{'}))}{u(t^{'}, a(t^{'}))},$$

or

$$y(t) - y(t^{'}) \geq u(t, a(t^{'})) - u(t^{'}, a(t^{'})) = \int_{t^{'}}^{t} u_{\theta}(s, a(t^{'}))ds$$

Employing 5, we should show

$$\int_{t^{'}}^{t} \left[u_{\theta}(s, a(s)) - u_{\theta}(s, a(t^{'}))\right]ds \geq 0.$$

By log-Supermodular property of $u(\cdot,\cdot)$, the integral is positive. Now let $t^{'} > t$. We need to show

$$y(t) - y(t^{'}) \geq u(t, a(t^{'})) - u(t^{'}, a(t^{'})) = -\int_{t}^{t^{'}} u_{\theta}(s, a(t^{'}))ds.$$

Rewrite equation 5, $y(t) - y(t^{'}) = -\int_t^{t^{'}} u_\theta(s, a(s))ds$, and plugging in the above inequality, we can conclude

$$\int_t^{t^{'}} \big[u_\theta(s, a(t^{'})) - u_\theta(s, a(s))\big] ds \geq 0.$$

Again using log-Supermodular property of $u(\cdot, \cdot)$ the integral is positive. ■

Now given $a_*(.)$, we can find the payoff of the Agent $\pi(.)$, and the optimal inspection policy $i_*(.)$.

A.5. Proof of Theorem 2

Proof. Due to the maximum punishment rule, and Assumption 1, $a^I(\hat{\theta}, \theta) \leq \theta$ if $\hat{\theta} \neq \theta$. Observe that if $a(\theta) < \theta$, then the mechanism can choose $a(\theta) = \underline{\theta}$. Now given a mechanism $(i(\hat{\theta}), a^I(\hat{\theta}, \theta), a(\hat{\theta}))$, define $\theta^* = \sup\{a(\theta) | i(\theta) = 0\}$. Assumption 1, and the global IC imply that

$$\begin{aligned} i(\theta)\Big(u(\theta, a^I(\theta,\theta))\Big)\mathbb{1}_{a^I(\theta,\theta)\geq\theta} + (1 - i(\theta))\Big(u(\theta, a(\theta))\Big)\mathbb{1}_{a(\theta)\geq\theta} \\ \geq i(\hat{\theta})\Big(u(\theta, a^I(\hat{\theta},\theta))\Big)\mathbb{1}_{a^I(\hat{\theta},\theta)\geq\theta} + (1 - i(\hat{\theta}))\Big(u(\theta, a(\hat{\theta}))\Big)\mathbb{1}_{a(\hat{\theta})\geq\theta} \\ = (1 - i(\hat{\theta}))\Big(u(\theta, a(\hat{\theta}))\Big)\mathbb{1}_{a(\hat{\theta})\geq\theta}. \end{aligned}$$

for all $\hat{\theta} \neq \theta$, and $(\theta, \hat{\theta}) \in [\underline{\theta}, \overline{\theta}]^2$. The last inequality comes from the fact that either $a^I(\hat{\theta}, \theta) = \theta$, then $u(\theta, \theta) = 0$ (by Assumption 1), or $a^I(\hat{\theta}, \theta) < \theta$, then $u(\theta, a^I(\hat{\theta}, \theta)\mathbb{1}_{a^I(\hat{\theta},\theta)\geq\theta} = 0$. Thus we have

$$i(\theta)\Big(u(\theta, a^I(\theta,\theta))\Big)\mathbb{1}_{a^I(\theta,\theta)\geq\theta} + (1 - i(\theta))\Big(u(\theta, a(\theta))\Big)\mathbb{1}_{a(\theta)\geq\theta} \geq \Big(u(\theta, \theta^*)\Big)\mathbb{1}_{\theta^*\geq\theta}.$$

First it means that if $i(\theta) = 0$, and $a(\theta) = \underline{\theta}$, then $\theta^* \leq \theta$. Second if $i(\theta) = 0$, and $a(\theta) \neq \underline{\theta}$ means that $a(\theta) \geq \theta$, and then $a(\theta) \geq \theta^*$. However, by the definition of θ^*, we conclude $a(\theta) = \theta^*$. Therefore if $i(\theta) = 0$, then

$$a(\theta) = \begin{cases} \theta^* & \theta > \theta^* \\ \{\theta^*, \underline{\theta}\} & \theta = \theta^* \\ \underline{\theta} & \theta < \theta^*. \end{cases} \tag{6}$$

Third if $i(\theta) = 1$, and $\theta \leq \theta^*$, then $a^I(\theta, \theta) \geq \theta^*$. On the other hand, we know if $i(\theta) = 1$, then $a^I(\theta, \theta) \geq \theta$. To see this by contradiction assume $a^I(\theta, \theta) < \theta$, then if the Principal chooses $i(\theta) = 0$, and $a(\theta) = \underline{\theta}$, will have higher payoff with no effect on the global IC. The higher payoff comes from the fact that inspection has a positive cost $\phi > 0$. Thus from IC, two necessary conditions are $a^I(\theta, \theta) \geq \max\{\theta^*, \theta\}$, and condition 6.

Rewrite the Principal's problem

$$\max_{i(.), a^I(.,.), a^N(.)} \int_{\underline{\theta}}^{\overline{\theta}} \Big[i(\theta) v(\theta, a^I(\theta, \theta)) + (1 - i(\theta)) v(\theta, a(\theta)) \mathbb{1}_{a(\theta) \geq \theta} - \phi i(\theta) \Big] \, dF(\theta),$$

subject to the IC constraints. For the moment we consider weaker conditions, that are $a^I(\theta, \theta) \geq \max\{\theta^*, \theta\}$, and condition 6. Later we check the global IC condition. By Assumption 2, the objective function is decreasing in $a^I(\theta, \theta)$. Therefore $a^I(\theta, \theta) = \max\{\theta^*, \theta\}$. Rewriting the objective function

$$\max_{i(.), \theta^*} \left\{ \int_{\underline{\theta}}^{\overline{\theta}} \Big[i(\theta) \Big(v(\theta, \max\{\theta^*, \theta\}) - \phi \Big) + (1 - i(\theta)) \Big(v(\theta, \theta^*) \Big) \mathbb{1}_{\theta^* \geq \theta} \Big] \, df(\theta) \right\}.$$

Now if $\theta \leq \theta^*$, then

$$v(\theta, \max\{\theta^*, \theta\}) - \phi < v(\theta, \theta^*),$$

therefore the optimal policy chooses $i(\theta) = 0$ for all $\theta \leq \theta^*$. For $\theta > \theta^*$, the optimal policy chooses $i(\theta) = 1$, iff $v(\theta, \theta) - \phi \geq 0$. Therefore the objective becomes to solve problem $\mathbb{P}_D$

$$\max_{\theta^* \in [\underline{\theta}, \overline{\theta}]} \left\{ \int_{\theta^*}^{\overline{\theta}} \Big(v(\theta, \theta) - \phi \Big) \mathbb{1}_{v(\theta, \theta) \geq \phi} \, dF(\theta) + \int_{\underline{\theta}}^{\theta^*} \Big(v(\theta, \theta^*) \Big) \, dF(\theta) \right\}.$$

The above explanation, and the optimal θ^* for problem $\mathbb{P}_D$, together suggests that the following policy is optimal for the weaker IC conditions.

$$a^I(\hat{\theta}, \theta) = \theta,$$

$$i(\hat{\theta}) = \begin{cases} 0 & \hat{\theta} \leq \theta^* \\ 1 & \theta^{**} \geq \hat{\theta} > \theta^* \\ 0 & \hat{\theta} > \theta^{**}, \end{cases}$$

$$a(\theta) = \begin{cases} \theta^* & \hat{\theta} \le \theta^* \\ \theta & \theta^{**} \ge \hat{\theta} > \theta^* \\ \underline{\theta} & \hat{\theta} > \theta^{**}, \end{cases}$$

Now we have to show the above policy is globally IC. The argument is as follows. Types $\theta \le \theta^*$, cannot mimic types $\hat{\theta} \ge \theta^*$, since they are either inspected or excluded. Types $\theta \le \theta^*$, are indifferent to mimic types $\hat{\theta} < \theta^*$, since $a(\hat{\theta}) = \theta^*$. Types $\theta > \theta^*$, cannot mimic types $\hat{\theta} \ge \theta^*$, since they are either inspected or excluded. Types $\theta > \theta^*$, cannot mimic types $\hat{\theta} < \theta^*$, since $a(\hat{\theta}) = \theta^* < \theta$. ■

A.6. Proof of Theorem 3

Proof. (i) Let m_I and m_N be two different messages. Consider the following strategy for the agent. $m(\theta) = m_N$ for $\theta \in [\underline{\theta}, \theta^*] \cup [\theta^{**}, \overline{\theta}]$ and $m(\theta) = m_I$ for $\theta \in (\theta^*, \theta^{**})$. The Principal commits to $a(\cdot)$ as follow: $a(m_N) = \theta^*$ and $a(m) = \underline{\theta}$ for all $m \ne m_I$. The inspection strategy of the Principal is as follows: $i(m_I) = 1$ and $i(m) = 0$ for all $m \ne m_I$. Let the principal's belief for off the equilibrium path messages be (with probability one) equal to type $\underline{\theta}$.

The Agent does not have any deviation strategy. Types in $[\underline{\theta}, \theta^*]$ strictly prefer to send message m_N. Types in $(\theta^*, \overline{\theta}]$ are indifferent to send any message since their payoff will be zero anyhow.

The Principal does not have any deviation strategy. If the Principal does not inspect message m_I, then types (θ^*, θ^{**}) will reject the action $a(m_I) = \underline{\theta}$. The payoff of inspection is strictly higher since $v(\theta, \theta) - \phi > 0$ for all $\theta \in (\theta^*, \theta^{**})$. The Principal does not inspect message m_N. By contradiction, suppose the ex-ante payoff of inspecting types in $[\underline{\theta}, \theta^*] \cup [\theta^{**}, \overline{\theta}]$ is higher than mandating action $a(\theta) = \theta^*$. We know, the ex-ante payoff of inspection types in (θ^*, θ^{**}) is strictly higher than mandating action $a(\theta) = \theta^*$ for these types. This implies that inspecting all types yields a higher payoff than mandating action $a(\theta) = \theta^*$ for types $[\underline{\theta}, \theta^*] \cup [\theta^{**}, \overline{\theta}]$ and inspecting types in (θ^*, θ^{**}). Hence inspecting all types has a higher payoff than the maximizer to problem $\mathbb{P}_D$. A contradiction.

(ii) Suppose $u(\cdot, \cdot)$ is log-Supermodular. We find all equilibria (ex-ante) payoffs for the Principal and show payoffs are (weakly) less than the ex-ante payoff of the equilibrium in (i). Fix an equilibrium E. Define $\tilde{\theta} = \sup\{\theta | \pi(\theta) > 0\}$. Let $\mathcal{M}$ be the set of on equilibrium path messages. By a similar argument to Lemma 3, one can show if $m \in \mathcal{M}(\theta)$ and $m' \in \mathcal{M}(\theta')$, where $\tilde{\theta} > \theta' > \theta$, then $i(m') \ge i(m)$ and

$a(m') \geq a(m)$. Denote $\check{i}(\theta) = i(m(\theta))$ and $\check{a}(\theta) = a(m(\theta))$. Therefore both $\check{i}(\cdot)$ and $\check{a}(\cdot)$ are weakly increasing. For ease of notation we use $a(\cdot)$ instead of $\check{a}(\cdot)$ and $i(\cdot)$ instead of $\check{i}(\cdot)$. Hence without loss of generality we can assume $i(\cdot)$ and $a(\cdot)$ are weakly increasing.

Define $\Theta_I = \{\theta | i(m(\theta)) > 0\}$, $\Theta_N = \{\theta | i(m(\theta)) = 0\}$. It is easy to see $a^* \equiv a(m(\theta)) = a(m(\theta'))$ for all θ, θ' in Θ_N. The ex-ante payoff of The Principal under E is

$$\mathbb{E}\big[v(\theta, a^*)\mathbb{1}_{a^* \geq \theta} | \theta \in \Theta_N\big]\mathfrak{M}(\Theta_N) + \mathbb{E}\big[v(\theta,\theta) - \phi | \theta \in \Theta_I\big]\mathfrak{M}(\Theta_I),$$

where $\mathfrak{M}(\Theta_N)$ and $\mathfrak{M}(\Theta_I)$ are the probability measure of Θ_N and Θ_I respectively. $a(\cdot)$ and $i(\cdot)$ are increasing therefore for all $\theta_I \in \Theta_I$ and $\theta_N \in \Theta_N$, $\theta_N \leq \theta_I$ if $\pi(\theta_N) > 0$. Moreover, $\theta_N \leq a^* \leq \theta_I$ if $\pi(\theta_N) > 0$. Therefore all equilibria payoffs for the Principal can be represented by the value of problem $\mathbb{P}_a$.

$$\mathbb{P}_a(a^*, \Theta_I, \Theta_N): \quad \mathbb{E}\big[v(\theta, a^*)\mathbb{1}_{a^* \geq \theta} | \theta \in \Theta_N\big]\mathfrak{M}(\Theta_N) + \mathbb{E}\big[v(\theta,\theta) - \phi | \theta \in \Theta_I\big]\mathfrak{M}(\Theta_I),$$

subject to the incentive of the Agent: $\Theta_I \subset [a^*, \overline{\theta}]$ and the incentive of the Principal:

$$\mathbb{E}\big[v(\theta,\theta) - \phi | \theta \in \Theta_N\big] \leq \mathbb{E}\big[v(\theta, a^*)\mathbb{1}_{a^* \geq \theta} | \theta \in \Theta_N\big].$$

Note that we do not need to check the incentive of the Principal to inspect m_I since the Principal can commit to a very inefficient action $a(m_I)$ and excludes all types.

In order to find the maximum of $\mathbb{P}_a$, fix a^*. The optimal Θ_I, by Assumption 2 is an interval starting from a^*. Thus the above optimization yields a similar payoff to problem $\mathbb{P}_D$.

The constraint $\mathbb{E}\big[v(\theta,\theta) - \phi | \theta \in \Theta_N\big] \leq \mathbb{E}\big[v(\theta, a^*)\mathbb{1}_{a^* \geq \theta} | \theta \in \Theta_N\big]$, implies that the minimum of $\mathbb{P}_a$ is $\max\{\mathbb{E}[v(\theta,\theta) - \phi], 0\}$. Now suppose $\Theta_I = [a^*, a^{**}]$. The $\mathbb{P}_a$ is continuous in a^* and a^{**}. Therefore $\mathbb{P}_a$ contains all values between the min and max. ∎

A.7. Proof of Lemma 4

Proof. (i) Let m_1 and m_2 be two messages on the equilibrium path that the Principal inspects with probability less than one, i.e. $i(m_1) < 1$ and $i(m_2) < 1$.

Let $\Theta_i = \{\theta \in [\underline{\theta}, \overline{\theta}] : m(\theta) = m_i\}$ for $i \in \{1, 2\}$. Let

$$\theta_{m_i} = \sup \left\{ \theta | m = m_i, \text{ and } \int_{t<\theta} \beta(\theta|m = m_i)\, d\theta = 0 \right\}.$$

By definition, the probability measure of types less than θ_{m_i} that send message m_i is zero. By a simple computation, for a small enough $\delta > 0$

$$\begin{aligned}
&\int_{\underline{\theta}}^{\theta_{m_i}+\delta} \left[v(\theta, \theta_{m_i} + \delta) \right] \beta(\theta|m = m_i)\, d\theta \\
&= \int_{\underline{\theta}}^{\theta_{m_i}} \left[v(\theta, \theta_{m_i} + \delta) \right] \beta(\theta|m = m_i)\, d\theta + \int_{\theta_{m_i}}^{\theta_{m_i}+\delta} \left[v(\theta, \theta_{m_i} + \delta) \right] \beta(\theta|m = m_i)\, d\theta \\
&= \int_{\theta_{m_i}}^{\theta_{m_i}+\delta} \left[v(\theta, \theta_{m_i} + \delta) \right] \beta(\theta|m = m_i)\, d\theta > 0.
\end{aligned}$$

From second line to third line is by the fact that probability measure of types less than θ_{m_i} is zero and $v(\cdot, \cdot)$ is bounded. The above value is strictly positive since by Assumption 2, $v(\theta, \theta) > 0$ and the probability measure of types between θ_{m_i} and $\theta_{m_i} + \delta$ is strictly positive. As a consequence, for $i \in \{1, 2\}$ we have

$$\max_{\tilde{a}} \int_{\underline{\theta}}^{\tilde{a}} \left[v(\theta, \tilde{a}) \right] \beta(\theta|m = m_i)\, d\theta > 0.$$

Best reply of the Principal $a(m)$ is

$$a(m) = \operatorname*{argmax}_{\tilde{a}} \int_{\underline{\theta}}^{\overline{\theta}} \left[V(\theta, \tilde{a})\, \mathbb{1}_{\tilde{a} \geq \theta} \right] \beta(\theta|m)\, d\theta = \int_{\underline{\theta}}^{\tilde{a}} \left[V(\theta, \tilde{a}) \right] \beta(\theta|m)\, d\theta.$$

The maximum is strictly positive, therefore $a(m_i) > \inf\{\Theta_i\}$. Moreover for all $\epsilon > 0$ small enough we can find a type $\theta_i \in \Theta_i$ such that $a(m_i) - \theta_i < \epsilon$. Otherwise $a(m_i)$ is not optimal and the Principal can decrease it.

Now by contradiction assume $a(m_1) > a(m_2)$, then types in Θ_2 that are very close but less that $a(m_2)$ want to deviate and send message m_1. Formally

$$(1 - i(m_1))u(\theta, a(m_1))\mathbb{1}_{a(m_1) \geq \theta} > (1 - i(m_2))u(\theta, a(m_2))\mathbb{1}_{a(m_2) \geq \theta},$$

for $\theta \in \Theta_2$ and small ϵ such that $\theta \in (a(m_2) - \epsilon, a(m_2))$. The reason is for small enough ϵ, the right side goes to zero, and the left side is always higher than a positive amount. A contradiction. Hence $a(m_1) = a(m_2)$.

(ii) By contradiction assume $i(m_1) > i(m_2)$. By (i) $a(m_1) = a(m_2)$. Let $\Theta_i = \{\theta \in [\underline{\theta}, \overline{\theta}] : m(\theta) = m_i\}$ for $i \in \{1, 2\}$. Then types in Θ_1 that are less than $a(m_1)$ want to deviate to m_2. A contradiction. ■

A.8. Proof of Theorem 4

Proof. By Lemma 4 we can focus on equilibria with two messages. Fix an equilibrium E. Define $\Theta_I = \{\theta | m(\theta) = m_I\}$ and $\Theta_N = \{\theta | m(\theta) = m_N\}$ under E. Denote $\mathfrak{M}(\Theta_N)$ and $\mathfrak{M}(\Theta_I)$ are the probability measure of Θ_N and Θ_I respectively.

Claim 3 *Suppose the Principal cannot commit to $a(\cdot)$. Then in any equilibrium $\Theta_I \subset [a(m_N), \overline{\theta}]$. Moreover, the payoff of any equilibrium is equivalent to the value of problem $\mathbb{P}_i$*

$$\begin{aligned}\mathbb{P}_i : &\mathbb{E}\big[(1 - i(m_N))v(\theta, a(m_N))\mathbb{1}_{a(m_N)\geq\theta} | \theta \in \Theta_N\big]\mathfrak{M}(\Theta_N) \\ &+ \mathbb{E}\big[i(m_N)(v(\theta, \theta) - \phi) | \theta \in \Theta_N\big]\mathfrak{M}(\Theta_N) + \mathbb{E}\big[v(\theta, \theta) - \phi | \theta \in \Theta_I\big]\mathfrak{M}(\Theta_I),\end{aligned}$$

subject to $\Theta_I \subset [a(m_N), \overline{\theta}]$ and

$$a(m_N) \in \operatorname*{argmax}_{\tilde{a}} \int_{\underline{\theta}}^{\overline{\theta}} \left[v\left(\theta, \tilde{a}\right) \mathbb{1}_{\tilde{a}\geq\theta}\right] \beta(\theta|m_N)\, d\theta.$$

Proof. $a(m_N)$ should be less than all types in Θ_I, otherwise types $\theta \in (\inf \Theta_I, a(m_N))$ deviate and send message m_N. $\Theta_I \subset [a(m_N), \overline{\theta}]$ is no deviation constraint for the Agent and the last equation of the Claim is no deviation constraint for the Principal. ■

(i) Let $a(m_N) = s^*$. The value of $i(m_N)$ in $\mathbb{P}_i$ does not affect any constraint. Therefore, in order to find the maximum, $i(m_N)$ should be 0 if $\mathbb{E}\big[v(\theta, s^*)\mathbb{1}_{s^*\geq\theta} | \theta \in \Theta_N\big] \geq \mathbb{E}\big[(v(\theta, \theta) - \phi) | \theta \in \Theta_N\big]$, otherwise $i(m_N)$ should be 1. If $i(m_N) = 1$, then full inspection will be the maximum payoff. Therefore let $i(m_N) = 0$. Second we show $\Theta_I = [s^*, s^{**}]$ for some $s^{**} < \overline{\theta}$. Fix an optimal Θ_I. By chaining Θ_I to $[s^*, s^{**}]$ (keeping the same probability measure), s^* does not change but the ex-ante value of the principal will (weakly) increase since $v(\theta, \theta)$ is an increasing function. After change, s^* does not change since for all $s \leq s^*$

$$\int_{\underline{\theta}}^{s} \left[v\left(\theta, s\right) \mathbb{1}_{s\geq\theta}\right] \beta(\theta|m_N)\, d\theta \leq \int_{\underline{\theta}}^{s^*} \left[v\left(\theta, s\right) \mathbb{1}_{s^*\geq\theta}\right] \beta(\theta|m_N)\, d\theta.$$

The above inequality holds since both sides remain the same before and after the change. The above inequality is correct for all $s \geq s^*$. The reason is the left side (weakly) decreases. Therefore we can rewrite the optimization problem as follows

$$\mathbb{P}_2 : \max_{s^*,s^{**}} \left\{ \int_{\underline{\theta}}^{s^*} \Big(v(\theta, s^*) \Big) \, dF(\theta) + \int_{s^*}^{s^{**}} \Big(v(\theta,\theta) - \phi \Big) \, dF(\theta) \right\},$$

subject to

$$s^* \in \operatorname*{argmax}_{\tilde{s}} \int_{\theta \in [\underline{\theta}, s^*] \cup [s^{**}, \overline{\theta}]} \Big[v\left(\theta, \tilde{s}\right) \mathbb{1}_{\tilde{s} \geq \theta} \Big] \, dF(\theta).$$

(ii) Now we show if $\phi < v(\overline{\theta}, \overline{\theta})$, (no exclusion region), then the maximum payoff is equivalent to optimal deterministic inspection with full commitment. In this case, the payoff of $\mathbb{P}_2$ is maximized when $s^* = \theta^*$ and $s^{**} = \overline{\theta}$. We just need to show $a(m_N) = \theta^*$. By contradiction, first assume $a(m_N) < \theta^*$, then

$$\int_{\underline{\theta}}^{a(m_N)} \Big[v(\theta, a(m_N)) \Big] \, f(\theta) d\theta > \int_{\underline{\theta}}^{\theta^*} \Big[v(\theta, \theta^*) \Big] \, f(\theta) d\theta.$$

Therefore

$$\begin{aligned}
&\int_{\underline{\theta}}^{a(m_N)} \Big[v(\theta, a(m_N)) \Big] \, f(\theta) d\theta + \int_{a(m_N)}^{\overline{\theta}} \Big[v(\theta,\theta) - \phi \Big] \, f(\theta) d\theta \\
&\qquad > \int_{\underline{\theta}}^{a(m_N)} \Big[v(\theta, a(m_N)) \Big] \, f(\theta) d\theta + \int_{\theta^*}^{\overline{\theta}} \Big[v(\theta,\theta) - \phi \Big] \, f(\theta) d\theta \\
&\qquad\qquad > \int_{\underline{\theta}}^{\theta^*} \Big[v(\theta, \theta^*) \Big] \, f(\theta) d\theta + \int_{\theta^*}^{\overline{\theta}} \Big[v(\theta,\theta) - \phi \Big] \, f(\theta) d\theta.
\end{aligned}$$

The value of the first line cannot be higher than the third line, since θ^* is the solution to $\mathbb{P}_D$. A contradiction.

Second, assume $a(m_N) > \theta^*$. We can conclude $a(m_N) > \theta^{**}$ since after observing m_I by Principal she puts zero probability in interval (θ^*, θ^{**}). However, $a(m_N) > \theta^{**}$ is impossible since $\theta^{**} = \overline{\theta}$.

(iii) Now we show $s^{**} \geq \theta^{**}$. By contradiction assume $s^* \leq s^{**} < \theta^{**}$. By changing s^{**} to θ^{**}, the value of $\mathbb{P}_2$ increases. We show this change cannot affect the constraint of $\mathbb{P}_2$. After the change, s^* does not change since for all $s \leq s^*$

$$\int_{\underline{\theta}}^{s} \left[v\left(\theta, s\right) \mathbb{1}_{s \geq \theta} \right] \beta(\theta | m_N) \, d\theta \leq \int_{\underline{\theta}}^{s^*} \left[v\left(\theta, s\right) \mathbb{1}_{s^* \geq \theta} \right] \beta(\theta | m_N) \, d\theta.$$

The above inequality holds since by multiplying both sides by $\frac{F(s^*)+1-F(\theta^{**})}{F(s^*)+1-F(s^{**})}$, the inequality transforms to the inequality before the change. The above inequality for the same reason is correct for all $s \geq \theta^{**} > s^{**}$. ■

A.9. Proof of Theorem 5

Proof. Note that since the Principal cannot commit to $a(\cdot)$, we can use both Lemma 4, and Claim 3. Let $a(m_N) = s^*$. By Assumption 4, the Agent's strategy $m(\theta)$ alternates between m_I, and m_N in finite intervals for types in $[s^*, \overline{\theta}]$.

We show $\Theta_I = [s^*, s^{**}]$ for some $s^{**} \leq \overline{\theta}$. If there exists $s^{**} \in [s^*, \overline{\theta}]$ such that $m(\theta) = m_I$ for $\theta \in [s^*, s^{**})$ and $m(\theta) = m_N$ for $\theta \in [s^{**}, \overline{\theta}]$, the claim has been proved. Otherwise, by contradiction, there exist two intervals $[t_1, t_2]$ and $[t_2, t_3]$ such that $m(t) = m_N$ for $t \in [t_1, t_2]$ and $m(t) = m_I$ for $t \in [t_2, t_3]$. The idea is to switch two small sub intervals. Let $[t_1', t_2] \subset [t_1, t_2]$ and $[t_2, t_2'] \subset [t_2, t_3]$ such that the probability measure of $[t_1', t_2]$ and $[t_2, t_2']$ are the same in $F(\cdot)$. Therefore after the switch $m(t) = m_I$ for $t \in [t_1', t_2]$ and $m(t) = m_N$ for $t \in [t_2, t_2']$. Denote β_s and β the posterior belief of the Principal after and before the switch respectively. This switch, does not change $a(m_N)$ since for all $\theta \leq s^*$, $\beta_s = \beta$ and for all $a \leq s^*$

$$\int_{\underline{\theta}}^{a} \left[v\left(\theta, s\right) \mathbb{1}_{a \geq \theta} \right] \beta_s(\theta | m_N) \, d\theta \leq \int_{\underline{\theta}}^{s^*} \left[v\left(\theta, s^*\right) \mathbb{1}_{s^* \geq \theta} \right] \beta_s(\theta | m_N) \, d\theta.$$

The above inequality holds since both sides remain the same before and after the switch. The above inequality is also correct for all $a \geq s^*$. The reason is the value of left side (weakly) decreases after the switch and the value of the right side remains the same. This implies that the value of not inspecting m_N does not change. However, the value of inspecting m_N strictly decreases after the switch.

Now we need to show after the switch, the Principal still wants to inspect message m_I. Denote $a_s(m_I)$, the action of the Principal after the switch and when observes m_I. First $a_s(m_I) \notin (t_1', t_2')$. The reason is interval (t_1', t_2') is very small, and by decreasing $a_s(m_I)$ to t_1 the Principal's value strictly increases. The value strictly increases since the Principal puts zero probability (posterior belief) on types in interval $[t_1, t_1')$ after observing m_I. For the same reason $a_s(m_I) \notin (t_1, t_1')$. Second if $a_s(m_I)$ is less than t_1, then $a(m_I)$ is also less than t_1 (before the switch). The reason is after the switch, the value of not inspecting m_I (weakly) increases, but action $a = a_s(m_I)$ is available for the Principal before the switch. Action $a = a_s(m_I)$ yields the same payoff as the payoff of not inspecting after the switch. Therefore the switch does not change the value of not inspecting messages m_I. However, the value of inspecting m_I strictly increases. Third, let $a_s(m_I)$ be higher than t_2'. The switch increases the value of inspection by

$$\Big(\int_{t_1'}^{t_2} v(\theta,\theta)\, dF(\theta) - \int_{t_2}^{t_2'} v(\theta,\theta)\, dF(\theta)\Big)\frac{1}{\mathfrak{M}(\Theta_I)}.$$

We argue the value of not inspecting m_I increases (the value after minus before) at most

$$\Big(\int_{t_1'}^{t_2} v(\theta, a_s(m_I))\, dF(\theta) - \int_{t_2}^{t_2'} v(\theta, a_s(m_I))\, dF(\theta)\Big)\frac{1}{\mathfrak{M}(\Theta_I)}.$$

The reason is

$$\int_{\underline{\theta}}^{\overline{\theta}} \Big[V\left(\theta, a(m_I)\right) \mathbb{1}_{a(m_I)\geq\theta}\Big]\, \beta(\theta|m_I)\, d\theta \geq \int_{\underline{\theta}}^{\overline{\theta}} \Big[V\left(\theta, a_s(m_I)\right) \mathbb{1}_{a_s(m_I)\geq\theta}\Big]\, \beta(\theta|m_I)\, d\theta,$$

and the value of not inspecting m_I is at most

$$\begin{aligned}
&\int_{\underline{\theta}}^{\overline{\theta}} \Big[V\left(\theta, a_s(m_I)\right) \mathbb{1}_{a_s(m_I)\geq\theta}\Big]\, \big(\beta_s(\theta|m_I) - \beta(\theta|m_N)\big)\, d\theta \\
&= \Big(\int_{t_1'}^{t_2} v(\theta, a_s(m_I))\, dF(\theta) - \int_{t_2}^{t_2'} v(\theta, a_s(m_I))\, dF(\theta)\Big)\frac{1}{\mathfrak{M}(\Theta_I)}.
\end{aligned}$$

Finally by Assumption 4

$$\begin{aligned}
&\Big(\int_{t_1'}^{t_2} v(\theta,\theta)\, dF(\theta) - \int_{t_2}^{t_2'} v(\theta,\theta)\, dF(\theta)\Big)\frac{1}{\mathfrak{M}(\Theta_I)} \\
&\geq \Big(\int_{t_1'}^{t_2} v(\theta, a_s(m_I))\, dF(\theta) - \int_{t_2}^{t_2'} v(\theta, a_s(m_I))\, dF(\theta)\Big)\frac{1}{\mathfrak{M}(\Theta_I)},
\end{aligned}$$

which implies the value of inspecting message m_I after the switch is still higher than not inspecting. We conclude $\Theta_I = [s^*, s^{**}]$. ■

Therefore we can rewrite the optimization problem as follows

$$\mathbb{P}_3: \max_{s^*, s^{**}} \left\{ \int_{\underline{\theta}}^{s^*} \Big(v(\theta, s^*)\Big)\, dF(\theta) + \int_{s^*}^{s^{**}} \Big(v(\theta,\theta) - \phi\Big)\, dF(\theta) \right\},$$

subject to

$$s^* \in \operatorname*{argmax}_{\tilde{s}} \int_{\theta \in [\underline{\theta}, s^*] \cup [s^{**}, \overline{\theta}]} \Big[v\,(\theta, \tilde{s})\, \mathbb{1}_{\tilde{s} \geq \theta}\Big]\, dF(\theta),$$

and

$$\mathbb{E}\big[v(\theta,\theta) - \phi | \theta \in [s^*, s^{**}]\big]\mathfrak{M}(\Theta_I) \geq \max_{\tilde{s}} \int_{[s^*, s^{**}]} \Big[v\,(\theta, \tilde{s})\, \mathbb{1}_{\tilde{s} \geq \theta}\Big]\, dF(\theta).$$

A.10. Proof of Theorem 6

Proof. An ideal equilibrium suggestion for the highest equilibrium payoff for the Principal is as follows

- Strategies:

$$m(\theta) = \begin{cases} m_{NI} & \theta \in [\underline{\theta}, \theta^*], \\ m_I & \theta \in (\theta^*, \overline{\theta}] \end{cases}$$

$$i(m) = \begin{cases} 1 & m = m_I, \\ 0 & \text{otherwise} \end{cases}$$

$$a(m) = \begin{cases} \theta^* & m = m_{NI}, \\ \underline{\theta} & \text{otherwise} \end{cases}$$

$$a^I(m, \theta) = \theta.$$

- Beliefs: On the equilibrium path, beliefs are consistent with strategies and the off-path belief puts probability 1 on $\underline{\theta}$, i.e.

$$\beta(\underline{\theta} | m \neq m_I, m_{NI}) = 1.$$

The above strategy and beliefs generates the commitment payoff according to the Proposition 2. We need to check incentives of both players. The Agent does not have any incentive to deviate. Types above θ^* do not want to deviate to m_{NI} since the mandated action becomes less than their types. Types above θ^* will not deviate to m_I since they will be inspected and their payoff become zero. The principal can commit to inspect m_I, so we do not check the incentive of the Principal after observing m_I. Now we show $a(m_N) = \theta^*$, where

$$a(m_N) = \underset{\tilde{a}}{\text{argmax}} \int_{\underline{\theta}}^{\tilde{a}} \Big[V(\theta, \tilde{a}) \Big] \beta(\theta|m_{NI})\, d\theta.$$

By contradiction, first assume $a(m_N) < \theta^*$, then

$$\int_{\underline{\theta}}^{a(m_N)} \Big[v(\theta, a(m_N)) \Big] f(\theta) d\theta > \int_{\underline{\theta}}^{\theta^*} \Big[v(\theta, \theta^*) \Big] f(\theta) d\theta.$$

Therefore

$$\begin{aligned}\int_{\underline{\theta}}^{a(m_N)} \Big[v(\theta, a(m_N)) \Big] f(\theta) d\theta + \int_{\theta^*}^{\overline{\theta}} \Big[v(\theta, \theta) - \phi \Big] f(\theta) d\theta \\ > \int_{\underline{\theta}}^{\theta^*} \Big[v(\theta, \theta^*) \Big] f(\theta) d\theta + \int_{\theta^*}^{\overline{\theta}} \Big[v(\theta, \theta) - \phi \Big] f(\theta) d\theta.\end{aligned}$$

The above strict inequality is a contradiction since by the definition, the threshold θ^* should generate higher value than the threshold $a(m_N)$. Thus the left side should not be higher than the right side. A contradiction. Second, assume $a(m_N) > \theta^*$. We can conclude $a(m_N) > \theta^{**}$ since after observing m_{NI} by Principal she puts zero probability in interval (θ^*, θ^{**}). However, $a(m_N) > \theta^{**}$ is impossible since $\theta^{**} = \overline{\theta}$. Now we need to show by observing m_{NI} the principal does not inspect, formally

$$\int_{\underline{\theta}}^{a(m_N)} \Big[v(\theta, a(m_N)) \Big] \beta(\theta|m_{NI}) d\theta \geq \int_{\underline{\theta}}^{\theta^*} \Big[v(\theta, \theta) - \phi \Big] \beta(\theta|m_{NI}) d\theta,$$

we can replace $a(m_N)$ by θ^*

$$\int_{\underline{\theta}}^{\theta^*} \Big[v(\theta, \theta^*) \Big] \beta(\theta|m_{NI}) d\theta \geq \int_{\underline{\theta}}^{\theta^*} \Big[v(\theta, \theta) - \phi \Big] \beta(\theta|m_{NI}) d\theta.$$

By contradiction assume the reverse, thus we have

$$\int_{\underline{\theta}}^{\theta^*} \Big[v(\theta,\theta) - \phi\Big] f(\theta)d\theta + \int_{\theta^*}^{\overline{\theta}} \Big[v(\theta,\theta) - \phi\Big] f(\theta)d\theta$$
$$> \int_{\underline{\theta}}^{\theta^*} \Big[v(\theta,\theta^*)\Big] f(\theta)d\theta + \int_{\theta^*}^{\overline{\theta}} \Big[v(\theta,\theta) - \phi\Big] f(\theta)d\theta.$$

This is a contradiction by the definition θ^*. So the Principal does not any incentive to inspect message m_{NI}.

■