Trust with Evidence

Amirreza Ahmadzadeh[*] [†]

September 15, 2025

We study a dynamic principal–agent relationship in which an agent must exert costly effort to learn a privately observed binary state before taking an action. The principal wants to match the action with the state, while the agent is biased toward one action, generating both a moral hazard (effort choice) and an adverse selection (action choice) problem. The principal disciplines the agent through verification (at a cost), reduced workload and termination. We show reduced workload is always a valuable instrument, even when the cost of verification is small and the loss from shirking is large. By promising a reduced workload in the future, the principal can lower verification costs across multiple periods. For high biases, verification and reduced workload are insufficient instruments, and the principal must rely on firing along the equilibrium path. The threat of future firing complements verification and saves verification costs over time.

1 Introduction

Organizations often rely on experts or subordinates to take actions that must align with evolving states of the world, privately observed by the agent. Examples include a compliance officer deciding whether to escalate a case, a credit officer aligning a loan decision with borrower fundamentals, or a product manager selecting features to match latent user demand. In such settings, the principal cannot directly observe the state and faces two incentive problems: the agent must exert costly effort to learn the state (a moral hazard problem), and he may privately favor one action over another (a bias, akin to adverse selection). To address these challenges, the principal can demand evidence at a cost (verification), adjust the agent's workload over time, or, if necessary, terminate the relationship. Our question is how these instruments should be combined and how their use evolves over time.[1]

We study an infinite-horizon principal–agent relationship. In each period, the agent can exert costly effort to learn a binary state (L, R) and then chooses a binary action (L, R).[2] The principal decides whether to request evidence (verification at a cost) and whether to retain or fire the agent; if retained, the agent receives a fixed wage. The agent is biased toward R, deriving a private benefit from choosing it, independently of the state. Payoffs are discounted and unobservable to both players. Our analysis has two objectives. First, we characterize the the behavior underlying equilibria giving payoffs on the upper boundary of equilibrium payoffs. Second, we characterize the long-run dynamics of the relationship.

Our first theorem characterizes *behavior* along the upper boundary of the equilibrium payoff set as a function of the agent's continuation utility ($u \in [0, \overline{u}]$). Three regions emerge, separated by thresholds u_f and u_e ($0 \leq u_f \leq u_e \leq \overline{u}$). (i) *Firing region:* when the continuation utility is low, the principal uses firing as a disciplinary tool; the firing probability declines linearly from one at $u = 0$ to zero at $u = u_f$ and is never used thereafter. (ii) *Verification region:* for intermediate utilities $[u_f, u_e]$, the principal relies on verification, using it more often after the agent's preferred action (R); verification intensities are constant outside the interior interval and, inside it, fall with utility in a way that differentially penalizes the biased action. (iii) *Reduced-workload region:* for high utilities $[u_e, \overline{u}]$, the principal tolerates shirking occasionally; the probability of working declines strictly with utility, reflecting implicit delegation earned by past performance.

An implication of the characterization in Theorem 1 is that, counterintuitively,

---

the principal may benefit from the agent's bias. Although bias makes it harder to enforce action–state matching when the agent works, it simultaneously increases the agent's motivation to work and to demonstrate that the state is R. The benefit comes from two sources: first, the agent's reward from shirking and choosing R becomes stronger, which allows the principal to sustain working over a larger interval $[0, u_e]$; and second, the principal's concern about shirking diminishes, since she can lower the intensity of verification.

The second theorem characterizes the *structure of continuation utilities* that implement the upper boundary. Three properties are central. First, the agent's incentives to work and to match the action with the state are *decoupled*. Continuation utilities following verification sustain effort, while continuation utilities without verification ensure action–state matching.

Second, continuation utilities following verification are independent of the agent's action.[3] Intuitively, hard information obtained through verification dominates soft information inferred from actions when no verification occurs.

Third, when no verification occurs, continuation utility weakly increases if the agent takes his less-preferred action (L) and weakly decreases if he takes his preferred action (R). In other words, the agent is rewarded for selecting L and punished for selecting R.

Theorem 2 identifies two regions with important implications for the dynamics: the *evidence-based region* and the *hybrid region*. In the evidence-based region, the principal verifies after both actions with positive probability. Verification following the less-preferred action (L) ensures that the agent does not shirk while still choosing L. When no verification occurs, continuation utilities remain unchanged, so utilities evolve only through hard information (through verification). Since both verification probabilities are strictly positive, the principal relies on differential verification rather than distorting continuation utilities.

In the *hybrid region*, the principal never verifies after the agent takes L. However, the continuation utility in case of no verification strictly increases when the agent takes L and strictly decreases when he takes R. In this stage of the relationship, continuation utilities are shaped by both hard information (evidence) and soft information (unverified actions).

The third theorem analyzes the *dynamics* of the principal's preferred equilibrium. Two findings stand out. First, allowing the agent to shirk is always a valuable instrument, even when the cost of verification is small and the loss from not working (and thus potentially mismatching the action with the state) is large. By promising a lighter workload in the future, the principal can lower verification costs across multiple periods leading up to the present. We refer to this phenomenon as

---

[3]We focus only on on-path behavior. If the agent deviates and fails to provide evidence when the principal verifies, he is fired.

snowballing.

Second, if the evidence-based region is empty (which occurs under high biases), verification and reduced workload are insufficient instruments, and the principal must rely on firing along the equilibrium path. The threat of future firing complements verification and reduces verification costs across multiple periods. Hence, snowballing—through the promise of future punishment—renders the use of firing unavoidable. By contrast, when the evidence-based region is nonempty, increased verification substitutes for firing and prevents its snowballing effect.

Furthermore, we show that once the relationship leaves the evidence-based region, it never returns. If the hybrid region is empty, then once the relationship enters the reduced-workload region, it remains there indefinitely.

On the technical side, the history of the game can be fully summarized by the agent's continuation utility. The upper boundary of the equilibrium payoff set is self-generating in the sense of Abreu, Pearce, and Stacchetti (1990), which allows us to formulate the principal's problem as a recursive program subject to incentive constraints. By introducing endogenous variables for verification, firing, and working, we obtain a tractable formulation that highlights the distinct role of each instrument in sustaining equilibrium.

Roadmap.   Section 2 situates the paper within the existing literature. Section 3 introduces the model and the equilibrium concept. Section 4 presents the main results and builds intuition through a simple three-state automaton. Section 5 develops the recursive formulation, characterizes the upper boundary of equilibrium payoffs, and states the three main theorems. Section 5.3 provides interpretations and extensions, including the roles of commitment and transfers.

2 Literature Review

Our paper contributes to two strands of the literature. The first strand is dynamic monitoring. The work most closely related to ours is Bhaskar (2024), who study long-term principal–agent relationships with both adverse selection and moral hazard. The agent can work or shirk and has private types (high- or low-cost of effort). The principal monitors effort, using transfers and firing to provide incentives.[4] Monitoring serves to screen out low types by inducing shirking. The optimal self-enforcing contract has two phases: an initial screening phase, where the first monitored action reveals the agent's type, and a continuation phase, in which only high types remain. This second phase resembles our evidence-based structure when the agent is unbiased: effort is exerted until seniority, and

---

[4]An earlier version of their paper also considers the case without transfers.

4

monitoring intensity declines with the agent's track record. Unlike Bhaskar (2024), our agent has no persistent type but instead acquires private information about the state through costly effort (covert information acquisition).[5]

Recent work on dynamic mechanism design incorporates various forms of endogenous monitoring. Halac and Prat (2016) and Piskorski and Westerfield (2016), as well as Orlov (2022), study endogenous arrival rates of Poisson signals (good news in the former, bad news in the latter). Fahim et al. (2021) and Zeng (2022) allow the principal to choose the precision of Brownian monitoring. Liu (2011) and Marinovic et al. (2018) endogenize the number of past observations, while Varas et al. (2020) focus on the timing of observations. Dai et al. (2024) study the joint design of monitoring and compensation in a continuous-time moral hazard model with transfers. The principal allocates limited monitoring capacity between confirmatory and contradictory evidence of effort. Wong (2023), studies general monitoring schemes under a fixed exogenous wage. He shows that the optimal scheme is non-stationary, focusing on negative evidence of effort with increasing precision but decreasing frequency over time. All of these papers focus on broader aspects of monitoring design, but they do not address the role of bias or covert information acquisition in a dynamic relationship.[6]

The second strand is dynamic delegation. The closest paper in this literature is Lipnowski and Ramos (2020), who study an infinite-horizon game where, in each period, the principal decides whether to delegate a project adoption choice to the agent. The state is binary and i.i.d. across time, while the agent's preferences are state-dependent but biased toward one action. They characterize the set of equilibrium payoffs under fixed discounting and show that, unlike in dynamic agency models with commitment, the agent's autonomy diminishes over time rather than being rewarded through backloading. Li et al. (2017) study a repeated project-selection game in which, each period, a biased agent and an uninformed principal jointly select a project and then simultaneously choose implementation effort. The effort decision effectively grants the principal commitment power, as it provides each player with a tool to unilaterally punish the other. Unlike our paper, in the both paper the principal does not have endoegouse verification tool.

---

[5]There is a large literature on dynamic moral hazard, but these papers abstract from endogenous monitoring; see, for instance, Ray (2002), Levin (2003), Board (2011), Halac (2012), Nikolowa (2017), Fong (2008), and Guo and Hörner (2021).

[6]Solan and Zhao (2021) study a repeated inspection game between a principal and two agents. In each period, the principal—who has commitment power—can inspect at most one agent, while each agent decides whether to adhere or violate. They show that postponing rewards frees up resources for verification.

3 The Setup

A principal (she) aims to match an action with an unknown state. She employs an agent (he) to learn the state and take an action. They interact repeatedly. Time is discrete, and the horizon is infinite. In each period, $t = 1, 2, ...$, the interaction unfolds as follows.

Timing:

1. The principal chooses whether to fire (F) or retain $(\neg F)$ the agent. Firing ends the game.

2. The agent privately chooses to work (W) or shirk $(\neg W)$. If the agent works, he privately learns the state θ and has access to a set of evidence $S = \{\theta, \emptyset\}$. If the agent does not work, then $S = \{\emptyset\}$. The state $\theta \in \{L, R\}$ is i.i.d. across periods and the probability of $\theta = L$ is $\rho \geq 1/2$. Finally, the agent sends a message $m \in \{\hat{L}, \hat{R}\}$ to the principal.[7]

3. The agent publicly takes an action $a \in \{L, R\}$.

4. The principal publicly verifies the evidence (V) or not $(\neg V)$. If the principal verifies, the agent discloses an element $s \in S$.[8]

Payoffs: If the agent's action matches the state, the principal gets a flow payoff of $\pi > 0$. Otherwise, she gets 0. If the agent is retained, the principal pays him a per-period fixed wage $w > 0$. Verifying the evidence costs $c_P > 0$ to the principal and working costs $c_A \geq 0$ to the agent. In addition, action R yields a flow benefit $b \geq 0$ to the agent. That is, the agent does not care about the state, and he is biased toward action R.[9] Payoffs are not observed by either player.

Once the agent is fired, if ever, the game ends and both players get 0. Let $\pi(\theta, a) = \pi \mathbb{1}_{\{\theta = a\}}$. Both players share the same discount factor $\delta \in [0, 1)$. Therefore, if the agent is fired in period $T \in \mathbb{N} \cup \{0, +\infty\}$, the principal's payoff is

$$\sum_{t=0}^{T-1} \delta^t \big(\pi(\theta^t, a^t) - w - c_P \mathbb{1}_V\big).$$

---

[7]The message may be interpreted as a report about the state when the agent works and learns the state. Allowing a richer message space does not affect our results.

[8]Note that information acquisition is modeled differently from the literature on costly state verification. Here the principal cannot learn the state directly.

[9]In Section 5.3, we discuss how the results change when the agent is biased toward $a = L$.

Let $u(L) = w$ and $u(R) = w + b$. The agent's payoff is

$$\sum_{t=0}^{T-1} \delta^t (u(a^t) - c_A \mathbb{1}_W).$$

We assume the interaction is mutually beneficial: $\pi \geq w \geq c_A$, i.e., the agent's value to the principal exceeds the wage, and the wage covers the cost of working. We also assume the wage exceeds the principal's payoff when the agent always chooses the more likely state, i.e., $w > \rho\pi$. Of course interesting dynamics arise only if the agent is sufficiently patient: $\frac{\delta}{1-\delta}(w + (1-\rho)b) \geq c_A$, meaning the present value of future benefits exceeds today's working cost.

Throughout, players condition their play on the outcome of a public randomization device (prd). Denote the outcome of the prd regarding firing by $f_\xi \in \{F, \neg F\}$, working by $w_\xi \in \{W, \neg W\}$, action recommendation by $a_\xi \in \{L, R\}$ and verification by $u_\xi \in \{V, \neg V\}$.

Histories and Strategies: The solution concept we use is perfect public Bayesian equilibrium (PPBE).[10] We focus on pure strategy equilibrium. Let $f \in \{F, \neg F\}$, $w \in \{W, \neg W\}$, and $v \in \{V, \neg V\}$ denote generic choices of firing, working, and verification. Let ν be the outcome of verification: if $v = V$, $\nu = s$ and if $v = \neg V$, $\nu = \emptyset$. The set $\mathcal{H}^t$ of public histories h^t of the form

$$h^t = \Big((f_\xi^1, w_\xi^1, m^1, a_\xi^1, a^1, v_\xi^1, v^1, \nu^1), ..., (f_\xi^t, w_\xi^t, m^t, a_\xi^t, a^t, v_\xi^t, v^t, \nu^t) \Big),$$

corresponds to the public history of the game at the end of period t.[11]

A firing strategy of the principal is a map from the set of histories of the form (h^{t-1}, f_ξ^t) to $\{F, \neg F\}$, a verification strategy is a map from the set of histories of the form $(h^{t-1}, f_\xi^t, w_\xi^t, m^t, a_\xi^t, a^t, v_\xi^t)$ to $\{V, \neg V\}$.

A working strategy of the agent is a map from the set of histories of the form $(h^{t-1}, f_\xi^t, w_\xi^t)$ to $\{W, \neg W\}$. A messaging strategy of the agent is a map from the set of histories of the form $(h^{t-1}, f_\xi^t, w_\xi^t, S^t)$ to $\{\hat{L}, \hat{R}\}$. A strategy for the action of the agent is a map from the set of histories of the form $(h^{t-1}, f_\xi^t, w_\xi^t, S^t, m^t, a_\xi^t)$ to $\{L, R\}$. A strategy for evidence disclosure of the agent is a map from the set of histories of the form $(h^{t-1}, f_\xi^t, w_\xi^t, S^t, m^t, a_\xi^t, v_\xi^t, v^t)$ to S^t.

---

[10]We follow Athey and Bagwell (2008): A PPBE is a Perfect Bayesian Equilibrium (PBE) in which strategies depend on public histories and payoff-relevant private information, that is, information that pertains only to the current round. Allowing for more general strategies does not affect the set of equilibrium payoffs.

[11]For ease of notation we do not include the firing decision of the principal in the history of the game since if $f = F$ the game ends. Therefore in all histories we assume $f = \neg F$.

To fix ideas, we focus on the principal-preferred equilibrium[12] that displays the following feature: whenever the agent works, he matches the action with the state.[13] Solving for this equilibrium calls for solving for the entire upper boundary of the equilibrium payoff set. Hence, the characterization of other equilibria (for instance, the agent-preferred equilibrium) follows as a by-product.

Figure 1: The upper boundary of the feasible payoff set.

Before analyzing the game, it is useful to characterize the set of feasible payoffs without incentive constraints. Figure 5 illustrates this set's upper boundary, which exhibits several interesting features. On the upper boundary, the principal does not use verification and there are four extreme points: $(0,0)$: firing, $(\frac{w+(1-\rho)b-c_A}{1-\delta}, \frac{\pi-w}{1-\delta})$: the agent works and matches the action with the state, yielding the highest feasible payoff for the principal, $(\frac{w}{1-\delta}, \frac{\rho\pi-w}{1-\delta})$: the agent does not work and always takes action L, $(\frac{w+b}{1-\delta}, \frac{\pi(1-\rho)-w}{1-\delta})$: the agent does not work and always takes action R, yielding the highest feasible utility for the agent.[14] Note that the principal's highest

feasible payoff cannot be achieved in equilibrium, as she does not verify, and the agent has no incentive to work and learn the state.

The minmax payoff vector, $(0, 0)$, is an equilibrium payoff vector: if the principal expects that the agent does not work, she fires him. Conversely, if the agent expects that the principal fires him at every opportunity, not working is optimal. Without loss of generality, in the principal's preferred equilibrium, we assume any observable deviation triggers the minmax equilibrium.[15]

There is a range of parameters for which immediate firing is the unique equilibrium. To avoid trivialities, we assume throughout that equilibria involving strictly positive payoffs exist. It will be clear from our explicit formulas what restriction this entails on the parameters.

---

on the upper boundary if and only if $\frac{1-\rho}{\rho} \frac{c_A}{2-\rho} > b$. Intuitively, this condition holds when the bias is sufficiently small.

[15]The following observations are also made without loss of generality: 1) If the agent works and truthfully reveals the state through message m, and the principal verifies, then the agent chooses $s = \theta$; otherwise, the principal fires the agent. 2) If the agent deviates—either by not working or by working and misreporting the state—and the principal verifies, then the agent chooses $s = \emptyset$. Even if the agent were to reveal the true state ($s = \theta$) after sending an incorrect message ($m \neq \theta$), the principal would still fire the agent. Note that, in this case, learning the state provides no benefit to the principal, since the agent's action has already been taken.

4 Three-State Automaton: Illustrating the Main Results

To build intuition about the trade-offs involved, we begin by examining three-state automata under the assumption that the principal has commitment power. We then show how these automata can be adjusted to incorporate the principal's incentives, so that the three-state automata constitute equilibria of our game. While the principal-preferred equilibrium does not belong to this class, it shares some important properties with the best equilibrium within that class.

The three states are the following. **Firing** (state $\mathcal{F}$): The principal fires the agent, both players receive zero payoff, and the game ends. **Retiring** (state $\mathcal{R}$): The agent never works, and the principal does not verify. Moreover, the agent always chooses his favorite action, $a = R$. **Matching** (state $\mathcal{M}$): The agent works, learns the state, and matches the action with the state (complies).[16]

Both $\mathcal{F}$ and $\mathcal{R}$ are absorbing states in which the players' actions are fixed. Hence, we focus on state $\mathcal{M}$. We determine (i) the probability of verification and (ii) the transition probabilities from $\mathcal{M}$ to $\mathcal{R}$ and $\mathcal{F}$, as functions of verification and the agent's action, in order to maximize the principal's payoff while ensuring that compliance is optimal for the agent.[17]

Let v^a denote the probability of verification after the agent takes action a, and let u be the agent's utility in state $\mathcal{M}$. Rather than working with transition probabilities, it is convenient to use continuation utilities. Denote the continuation utility when action a is taken and verification occurs by u_V^a,[18] and the continuation utility when action a is taken and verification does not occur by $u_{\neg V}^a$.[19]

In state $\mathcal{M}$, the principal must account for two incentive constraints: 1) Complying must be optimal for the agent. Formally,

$$\mathbb{E}\left[u(\theta) + v^\theta \delta u_V^\theta + (1 - v^\theta)\delta u_{\neg V}^\theta\right] - c_A \geq \max_a \{u(a) + (1 - v^a)\delta u_{\neg V}^a\}. \qquad \text{(WM)}$$

The left-hand side is the ex-ante utility if the agent complies. The right-hand side is the utility of the agent from shirking and choosing action a. In this case,

---

[16]From now on, when we say the agent "complies," we mean that he works, learns the state, and matches the action with the state.

[17]Transitions across states are implemented with the public randomization device. Hence, the prevailing state is common knowledge.

[18]As discussed in the model, it is without loss of generality to assume that, after verification, if $s \neq \theta$, the agent is fired. Therefore, for simplicity, u_V^a refers to the continuation utility in case the agent provides evidence, i.e., $s = \theta$.

[19]Continuation utilities below (above) u imply a transition to $\mathcal{F}$ ($\mathcal{R}$) with positive probability. Utilities are fixed at states $\mathcal{F}$ and $\mathcal{R}$ (0 and $\frac{w+b}{1-\delta}$). Hence, the transition probabilities are given by the ratio of the absolute difference between the current utility and the continuation utility to the absolute difference between the current utility and the continuation utility ($\mathcal{F}$ or $\mathcal{R}$).

the agent is fired after the verification and the continuation utility is zero. 2) Conditional on working, the agent must prefer to match the action with the state. Formally, for $a, a' \in \{L, R\}$,

$$u(a) + v^a \delta u_V^a + (1 - v^a) \delta u_{\neg V}^a \geq u(a') + (1 - v^{a'}) \delta u_{\neg V}^{a'}. \qquad \text{(M)}$$

We denote by (M–L) when $a = L$, $a' = R$, and by (M–R) when $a = R$, $a' = L$ in the constraint (M).

To understand how two distinct frictions—moral hazard (shirking) and adverse selection (mismatching the action with the state)—shape the equilibrium structure, we analyze we analyze two special cases: an unbiased agent ($b = 0$) and an agent with zero cost of effort ($c_A = 0$).

1) Unbiased agent. First, consider an unbiased agent. It is without loss of generality to assume 1) $v := v^L = v^R$ and 2) if the agent works, then he matches the action with the state. Intuitively, the agent does not have any incentive to mismatch the action after learning the state, and so the principal has no reason to condition verification on the chosen action. Therefore (M) is slack and the only binding incentive constraint is (WM).

Four observations are in order. First, the agent's utility from shirking and choosing L must be the same as shirking and choosing R. If shirking and choosing a is worse than shirking and choosing a', a mean-preserving contraction in u_V^a (verification) and $u_{\neg V}^a$ (no verification) benefits the principal.[20] Hence (WM) simplifies to

$$\mathbb{E}_\theta \left[v^\theta \delta u_V^\theta \right] \geq c_A.$$

Second, the continuation utilities after verification must be equal, i.e., $u_V := u_V^L = u_V^R$. A mean-preserving contraction of u_V^R and u_V^L does not affect (WM). Therefore, by concavity of the value function, the continuation utilities in the case of verification must be identical. Intuitively, when the agent provides evidence after verification, rewards or punishments are independent of the agent's actions.

Third, when (M) is slack, distorting $u_{\neg V}^L$ or $u_{\neg V}^R$ is not beneficial, i.e., $u_{\neg V}^L = u_{\neg V}^R = u$. The continuation utilities in the case of no verification, $u_{\neg V}^a$, do not appear in the incentive constraints. Thus, spreading them away from the initial utility u in state $\mathcal{M}$ is unnecessary for the principal. Hence, in the absence of verification, the state remains in $\mathcal{M}$.

---

[20]Note that if $u_V^a = u_{\neg V}^a$, then a mean-preserving contraction between continuation utilities cannot be implemented. However, she can reduce the cost of verification by decreasing the probability of verification v^a, provided that $v^a \neq 0$. If $v^a = 0$, the agent would (weakly) prefer to deviate.

Fourth, when (M) is slack, (WM) must bind, i.e., $v\delta u_V = c_A$. Otherwise, the principal could reduce the probability of verification without affecting the incentive constraint. Therefore, the probability of verification, v, and the continuation utility, u_V, serve as substitute tools for providing incentives to the agent: reducing one necessitates increasing the other.

All of these observations extend to the principal's preferred equilibrium, without relying on the assumption of the principal's commitment power or restricting attention to three-state automata.

These observations, along with the agent's promise-keeping, pin down v and u_V.[21] We are left with solving for the optimal utility u in state $\mathcal{M}$. It is easy to show that $u_V \geq u$. Intuitively, if verification occurs and the principal finds that the agent has complied, then she does not expose the agent to firing. When the only binding constraint is (WM), the agent is not fired on the equilibrium path.

Figure 2: Three-state automaton. When verification does not occur, the state remains unchanged. If verification occurs and the agent provides evidence, the state transitions to $\mathcal{R}$ with probability $p = (u_V - u)/(\frac{w+b}{1-\delta} - u)$. If, after verification, the agent does not provide evidence ($s \neq \theta$), the state transitions to $\mathcal{F}$ (dashed arrow).

To sum up, there are only two candidates for the optimal three-state automaton. If the verification cost is small ($c_P \leq \bar{c}_P$), the state always remains at $\mathcal{M}$, i.e., $u_V = u$.[22] In this case, the agent is neither fired nor retired (since, as noted earlier,

---

[21]Formally, $u = \mathbb{E}\left[u(\theta) + v^\theta \delta u_V^\theta + (1-v^\theta)\delta u_{\neg V}^\theta\right] - c_A$. Therefore $v = \frac{w + \delta u - u}{\delta u}$ and, $u_V = \frac{c_A}{\delta v} = \frac{c_A u}{w + \delta u - u}$.

[22]Explicitly, $\bar{c}_P = \frac{\delta \frac{w}{1-\delta}\left(\frac{\rho\pi}{c_A} - 1\right)}{2 + \frac{c_A}{w - c_A}}$.

$u_{\neg V}^L = u_{\neg V}^R = u_V = u$). If instead, $c_P > \bar{c}_P$, the state transitions from $\mathcal{M}$ to $\mathcal{R}$ with positive probability, i.e., $u_V > u$. This result is intuitive: when c_P is small, the principal prefers to verify more often and never retire the agent, rather than to reduce the probability of verification and retire the agent with positive probability.

Interestingly, by considering more complete strategy profiles (beyond three state automata), building on this automaton, retiring the agent becomes a valuable tool too, independently of how small the cost of verification is. The principal can promise a lighter workload in the future and reduce the verification costs over multiple periods leading up to that time.

To see this, let us consider using a sequence of three-state automata. First, consider an auxiliary three-state automaton with states $\mathcal{F}$, $\mathcal{R}$, and $\mathcal{M}_1$. In state $\mathcal{M}_1$, the agent always complies; the principal retires the agent after the first verification. If verification does not occur, the state remains $\mathcal{M}_1$. The principal verifies after both actions with the same probability v, chosen so that the agent is indifferent between working and shirking.[23]

Next, we replace $\mathcal{R}$ with $\mathcal{M}_1$ in the initial three-state automaton.[24] Simple algebra shows that $u_V > u$ if and only if $c_P > \bar{c}_P^1$, where $\bar{c}_P^1 < \bar{c}_P$. Intuitively, retirement now occurs after two rounds of verification: once from $\mathcal{M}$ to $\mathcal{M}_1$ and once within $\mathcal{M}_1$, from $\mathcal{M}_1$ to $\mathcal{R}$. Retirement after two rounds saves on verification costs compared to a scenario in which retirement does not occur at all. Therefore, for a larger range of verification costs, retirement with positive probability becomes optimal.


Figure 3: Replacing $\mathcal{R}$ with $\mathcal{M}_1$.

---

[23]Formally, $v\delta \frac{w}{1-\delta} = c_A$.

[24]Formally, replace $\mathcal{R}$ with the payoffs of $\mathcal{M}_1$ and treat it as an absorbing state.

13

This intuition can be generalized. Introduce another auxiliary three-state automaton with states $\mathcal{F}$, $\mathcal{M}_1$, and $\mathcal{M}_2$. In state M_2, the agent always complies and after the first verification, the state transitions to $\mathcal{M}_1$, and if verification does not occur, the state remains in $\mathcal{M}_2$. The principal verifies after both actions with the same probability v, chosen so that the agent is indifferent between working and shirking. We then replace $\mathcal{R}$ with $\mathcal{M}_2$ in the original three-state automaton. Simple algebra shows that $u_V > u$ if and only if $c_P > \bar{c}_P^2$, where $\bar{c}_P^2 < \bar{c}_P^1$. Similarly, one can introduce $\mathcal{M}_n$ and show that the sequence $\{\bar{c}_P^n\}$ is decreasing and converging to zero when n goes to infinity.[25]

By retiring the agent after n verifications, the principal reduces the verification probability at each of the n steps. The promise of future retirement creates incentives in all earlier periods. The savings from reduced verification costs accumulate across n periods, whereas the disutility of retirement is incurred only once. As a result, the ratio of the benefit from reduced verification costs to the disutility of retirement grows at rate n. For sufficiently large n, the benefit always outweighs the disutility. We refer to this phenomenon as "snowballing."[26]

Note that *snowballing* differs from the standard *backloading* result. While delaying rewards—i.e., *backloading*—is effective for sustaining incentives and is indeed operative in our model, the accumulating impact of such deferrals—*snowballing*—renders the deferred rewards effectively unavoidable, regardless of how small the verification cost.

2) The agent with zero cost of working. Now suppose $c_A = 0$, so that without loss of generality we may assume that the agent knows the state. In contrast to the unbiased case, (WM) is slack. Intuitively, since the agent knows the state, working is not a concern for the principal; the only concern is that he matches the action with the state.

Four results are worth highlighting. First, the probability of verification after R is strictly higher than after L, i.e., $v^R > v^L$. Since the agent is biased toward

---

[25]Explicitly $\bar{c}_P^n = \dfrac{\delta \frac{w}{1-\delta}\left(\frac{\rho\pi}{c_A}-1\right)}{n+2+\frac{c_A}{w-c_A}}$

[26]Note that while the *ratio* of benefits to disutility grows at rate n, the *net gain*—that is, the benefit minus the disutility—converges to zero as n becomes large. Let π_n denote the principal's payoff in state $\mathcal{M}_n$, and let π^* denote the payoff when retirement never occurs. Therefore,

$$\pi_{n-1} - \pi^* = \frac{\left(\frac{c_A}{w}\right)^n \left(n\frac{c_P}{\delta} - \frac{\rho\pi(w-c_A)}{(1-\delta)c_A}\right)}{1 - \left(\frac{c_A}{w}\right)^n}.$$

Hence, the benefit from reduced verification cost is: $\dfrac{n\frac{c_P}{\delta}\left(\frac{c_A}{w}\right)^n}{1-\left(\frac{c_A}{w}\right)^n}$, and the disutility is: $\dfrac{\frac{\rho\pi(w-c_A)}{(1-\delta)c_A}\left(\frac{c_A}{w}\right)^n}{1-\left(\frac{c_A}{w}\right)^n}$.

14

action R, to incentivize him to match the action with the state, verification occurs more frequently after R than after L.

Second, there is no need to verify after L, i.e., $v^L = 0$. Verification after L only matters for the incentive to work, which is absent here. Shifting verification from L to R always strengthens incentives.

Third, since (WM) is not binding, the continuation utility after verification remains undistorted, i.e., $u_V^R = u$. After verification, if the agent provides evidence, the state does not transition to $\mathcal{R}$ or $\mathcal{F}$. Since (M–R) never binds and the only binding constraint is (M–L), u_V^R does not appear in any binding constraint. Thus, u_V^R is not used to provide incentives; it must therefore remain undistorted.

Fourth, in the absence of verification, the continuation utility is lower than the initial utility when the agent takes R, and higher when the agent takes L; that is, $u_{\neg V}^L \geq u \geq u_{\neg V}^R$. Intuitively, when verification does not occur, the agent is exposed to firing (punished) after taking his preferred action (R) and to retiring (rewarded) after his less preferred action (L). If $u_{\neg V}^L < u_{\neg V}^R$, then a mean-preserving contraction of $u_{\neg V}^L$ and $u_{\neg V}^R$ relaxes (M–L) and benefits the principal. Hence, it must be that $u_{\neg V}^L \geq u_{\neg V}^R$. Moreover, the envelope theorem implies that the continuation utilities, $u_{\neg V}^L$ and $u_{\neg V}^R$, must bracket u.

All of these four observations extend to the principal's preferred equilibrium (without the restriction to three-state automata). Figure 4 displays the three-state automaton when $c_A = 0$.

Figure 4: Three-state automaton with $c_A = 0$. The transition probability from $\mathcal{M}$ to $\mathcal{R}$ under verification is $(u_V - u)/(\frac{w+b}{1-\delta} - u)$. If verification does not occur and the action is L, the transition probability from $\mathcal{M}$ to $\mathcal{R}$ is $(u_{\neg V}^L - u)/(\frac{w+b}{1-\delta} - u)$. If verification does not occur and the action is R, the transition probability from $\mathcal{M}$ to $\mathcal{F}$ is $(u - u_{\neg V}^R)/(u - 0)$.

We can characterize the use of retirement and firing in terms of two thresholds,

$c_P^f \geq c_P^r > 0$. Retirement occurs with strictly positive probability if and only if $c_P > c_P^r$, whereas firing occurs with strictly positive probability if and only if $c_P > c_P^f$. The intuition is simple: when verification cost is low, verification is the most efficient instrument, and the principal exclusively relies on it. As c_P increases, it becomes too costly to rely solely on verification, and the principal introduces retirement as a device: promising eventual retirement allows her to reduce the frequency of verification while still giving the agent incentives to comply. Firing, by contrast, is the harshest instrument because it terminates the relationship. Hence, firing appears only when verification costs are so high that neither verification nor retirement alone can sustain incentives.

Fixing c_P, an analogous set of thresholds $b^f \geq b^r > 0$ can be defined in terms of the agent's bias b. For small biases, verification alone suffices to discipline the agent. As the bias increases b^r, the temptation to mismatch the action with the state grows, and the principal must rely on retirement as an additional incentive device. When the bias becomes large ($b > b^f$), retirement is not enough, and firing is used with positive probability.

In summary, for small biases or verification costs, the state always remain at $\mathcal{M}$ and the principal does not use firing or retiring.

Interestingly, in the principal's preferred equilibrium (without the restriction to three-state automata), allowing the agent to shirk is always valuable: the principal benefits from snowballing even when c_P is small and π is large. To satisfy (M–L), the principal can either increase verification (v^R) while keeping $u_{\neg V}^R$ undistorted at u, or reward the agent by raising $u_{\neg V}^L$ above u. However, the latter is better. The cumulative savings in verification costs always dominate.

Moreover, firing the agent—even without deviation—serves as an effective instrument if (M–L) binds on the equilibrium path (as in the case $c_A = 0$). To satisfy (M–L), she must either increase verification (v^R) while keeping $u_{\neg V}^R$ undistorted at u, or punish the agent by reducing $u_{\neg V}^R$ below u. The threat of future firing complements verification and reduces verification costs across multiple periods. Snowballing, by promising future punishment, renders the use of firing unavoidable, regardless of how small b is.

We now combine the two cases, allowing for both a positive bias and a cost of working. In the principal's preferred equilibrium, (1) allowing shirking is always valuable; (2) if (M–L) binds along the entire equilibrium path, then firing occurs with probability one, which happens when the expected benefit from bias exceeds the cost of working, i.e., $(1 - \rho)b > c_A$; and (3) if (WM) binds only on part of the equilibrium path, the principal verifies both actions with positive probability to detect shirking, and firing does not occur. In this case, greater verification substitutes for firing and prevents its snowballing effect.

The principal's incentives. So far, we have assumed that $\mathcal{F}$ and $\mathcal{R}$ are

absorbing states with predetermined behaviors. We have also assumed that the principal never deviates from verifying whenever she is "supposed to."[27] Since $\mathcal{F}$ corresponds to the minmax equilibrium, it is sequentially rational. However, the agent cannot be permanently retired in $\mathcal{R}$; otherwise, the principal would deviate and fire him. This problem can be addressed with a slight modification of the automaton introduced above. To prevent this deviation, the agent must work at least occasionally. By replacing $\mathcal{R}$ with $\mathcal{RW}$ (reduced workload), we allow transitions from $\mathcal{RW}$ to $\mathcal{M}$. Finally, the transitions must be chosen such that the principal's payoff in $\mathcal{RW}$ exceeds $\frac{c_P}{\delta}$. This ensures that the principal has sufficient incentive to verify in $\mathcal{M}$.

Figure 5: Three-state automaton when $b = 0$, where $\mathcal{R}$ is replaced by $\mathcal{RW}$. The state transitions from $\mathcal{RW}$ to $\mathcal{M}$ with probability q.

5 Results

5.1 Principal's Program

Our goal is to characterize the preferred equilibrium of the principal. The history of the game can be summarized by the agent's continuation utility. The upper boundary of the equilibrium payoff set is self-generating, in the sense of Abreu, Pearce, and Stacchetti (1990). Accordingly, the problem can be formulated as a recursive program, subject to incentive constraints.

---

[27]When we say the principal is "supposed to" verify or fire, we mean that the prd recommends working, verifying or firing.

Lemma 3 in the Appendix shows that, without loss of generality, we may assume that continuation utilities depend only on the agent's action and are independent of messages when he is supposed to work. Moreover, when the agent is not supposed to work, he does not learn the state, and the continuation utility is independent of both his action and the messages.

If the agent is supposed to shirk, the principal does not verify. If the agent shirks when he is supposed to, verification yields no additional information and only incurs a cost.

Therefore, along the equilibrium path, the agent's continuation utility at the end of each period falls into three cases:

1. u_V^a, when the agent is supposed to work, he sends message $a \in \{L, R\}$, and the principal verifies.

2. $u_{\neg V}^a$, when the agent is supposed to work, he sends message $a \in \{L, R\}$ and the principal does not verify.

3. $u_{\neg W}$, when the agent is recommended not to work.[28]

Let $\Pi(u)$ be the principal's highest equilibrium payoff consistent with the agent's equilibrium payoff being u. Let $\Pi(u) = -\infty$ if no equilibrium exists that gives the agent u. Let $\overline{u}$ denote the maximum promise such that $\Pi(u) > -\infty$. The principal's program is:

$$\mathcal{P}: \quad \Pi(u) = \sup \mathbb{E}\Big[\mathbb{1}_{\neg F}\big(\pi(\theta, a) - c_P\,\mathbb{1}_V - w + \delta\,\Pi(u')\big)\Big],$$

subject to the agent's incentive constraints:
1) Working and matching the action with the state

$$\mathbb{E}\left[u(\theta) + v^\theta \delta u_V^\theta + (1 - v^\theta)\delta u_{\neg V}^\theta\right] - c_A \geq \max_a\{u(a) + (1 - v^a)\delta u_{\neg V}^a\}. \quad \text{(WM)}$$

2) Conditional on working, matching the action with the state,

$$u(a) + v^a \delta u_V^a + (1 - v^a)\delta u_{\neg V}^a \geq u(a') + (1 - v^{a'})\delta u_{\neg V}^{a'}. \quad \text{(M)}$$

for all $a, a' \in \{L, R\}$. 3) The promise-keeping constraint:

$$u = \mathbb{E}\big[\mathbb{1}_{\neg F}\left(u(\theta) - c_A\,\mathbb{1}_W + \delta\,u'\right)\big]. \quad \text{(PK)}$$

---

[28]Subscript "V" stands for *working and verifying*, subscript "$\neg V$" stands for *working and not verifying* and subscript "$\neg W$" stands for *not working*.

Principal's incentive constraints:
1) Firing

$$\mathbb{E}\Big[\pi(\theta, a) - c_P \mathbb{1}_V - w + \delta\Pi(u')\Big] \geq 0, \qquad\qquad \text{(Fi)}$$

2) Verifying

$$\delta\Pi(u_V^a) - c_P \geq 0, \qquad\qquad \text{(Ve)}$$

for all $a \in \{L, R\}$. The supremum is taken over the continuation utility u', the probability of firing, the probability of working, the probability of verifying and the probability that the agent takes action R when he does not work. The continuation utility u' can take the three aforementioned forms, $u_V^a, u_{\neg V}^a, u_{\neg W} \in [0, \overline{u}]$.[29]

The working and matching constraint (WM) ensures that the agent works and matches the action with the state. The matching constraint (M) guarantees that if the agent works, he matches the action with the state. The promise-keeping constraint (PK) requires that the agent's promised utility is equal to his actual expected utility.

The firing incentives of the principal are as follows: (i) The principal fires the agent when she is supposed to. Not firing is an observable deviation and triggers the minmax equilibrium. Therefore, this incentive is trivially satisfied, and there is no need to impose a constraint. (ii) The principal retains the agent when she is supposed to. In this case, an incentive constraint is required. The principal's payoff must be positive for all on-path utility levels of the agent. This incentive constraint is given by equation (Fi).

The verification incentives of the principal are as follows: (i) The principal does not verify when she is not supposed to. This constraint is trivially satisfied, as the principal knows that the agent has worked and has matched the action to the state, making verification unnecessary. (ii) The principal verifies whenever she is supposed to. In this case, an incentive constraint is needed. The principal's incentive to verify when she is supposed to is that the discounted expected payoff from verification, $\delta\Pi(u_V^a)$ (for $a \in \{L, R\}$), must be at least as large as the cost c_P, as represented by equation (Ve).

---

[29]Expectations are taken with respect to the random outcomes, namely the state, firing status, working status, verification status and the agent's action (when the agent is supposed to shirk).

5.2 Analysis

This section presents three theorems. The first characterizes behavior as a function of the agent's utility along the upper boundary of the equilibrium payoff set. The second examines the evolution of continuation utilities on this boundary. The third combines these results to describe the dynamic evolution of the relationship in the principal's preferred equilibrium.

Let $\tilde{u}$ denote the agent's utility when he consistently shirks and chooses L, and the principal neither verifies nor fires him, i.e., $\tilde{u} := \frac{w}{1-\delta}$.

Theorem 1 (Behavior). *There are unique thresholds u_f and u_e, with $0 < u_f \leq u_e < \overline{u}$, such that the behavior along the upper boundary of the equilibrium payoff set is characterized as follows:*

1. ***Firing Region:*** *The firing probability declines linearly from 1 at $u = 0$ to 0 at $u = u_f$. For all $u \geq u_f$, the agent is never fired.*

2. ***Verification Region:*** *The probability of verification after R is strictly higher than after L for all $u \in [0, \overline{u}]$, i.e., $v^R > v^L$. Both v^R and v^L are continuous in u and remain constant on the intervals $[0, u_f]$ and $[u_e, \overline{u}]$. On interval $[u_f, u_e]$, v^R is strictly decreasing, while v^L is zero for $\tilde{u} \leq u_f$; otherwise it is strictly decreasing up to $\min\{\tilde{u}, u_e\}$ and, if $\tilde{u} \leq u_e$, equal to zero in $[\tilde{u}, u_e]$.*

3. ***Reduced Workload Region:*** *The agent works with probability one whenever $u \leq u_e$. The probability of working declines strictly over $[u_e, \overline{u}]$. Moreover, the agent chooses R if the principal's marginal payoff at $\overline{u}$ is less than $\frac{\pi(1-2\rho)}{b}$, and chooses L otherwise.*

Proof in Appendix.

Figures 6 and 7 illustrate the firing, verification, and working probabilities as functions of continuation utilities. In what follows, we discuss the intuition behind Theorem 1 for each region.

Firing. When the agent's target utility is low ($u \leq u_f$), the principal lacks the flexibility to use reduced workload (lower effort) or verification alone to sustain incentive compatibility.

In this regime, she relies on firing with positive probability as the last-resort disciplinary tool. As u increases, the probability of firing decreases linearly, eventually reaching zero at the threshold u_f. Beyond this point, the incentives maintain without resorting to termination.

Verification. Because the agent is biased toward R, when the state is L, verifying more frequently after R than after L (i.e., $v^R > v^L$), allows the principal to raise the cost of choosing R when the state is L.

20

Figure 6: Working and firing probabilities as functions of continuation utilities.

Figure 7: Verification probabilities as functions of continuation utilities.

Verification remains constant in regions where the principal primarily relies on other instruments to deliver the promised utility—namely, firing at low utility levels ($u \in [0, u_f]$) and reduced workload at high utility levels ($u \in [u_e, \overline{u}]$). In the low-utility regime, verification is heavy; in the high-utility regime, it is light.

In the intermediate region $[u_f, u_e]$, the principal reduces verification, as u increases. In particular, the verification probability following L drops to zero once the agent's utility exceeds the utility from permanently shirking and choosing L.

Reduced Workload. For all $u \leq u_e$, the agent works with probability one. Once $u > u_e$, the principal begins to tolerate some shirking: the agent is no longer expected to exert effort in every period, and the probability of work declines as u increases. This shift marks a move toward implicit delegation, where the agent's past performance earns him greater discretion. Nevertheless, the agent never shirks with probability 1, as it yields a negative payoff to the principal.

In the reduced-workload region, for small biases, the agent's private gain from choosing R while shirking, is negligible relative to the principal's loss. As the bias grows large, the agent's gain becomes comparable to the principal's loss, making

shirking with R part of the equilibrium. Formally, if the agent's net benefit from shirking with R rather than L exceeds the principal's marginal payoff, the principal tolerates R while shirking.

A consequence of the characterization in Theorem 1 is that the principal may, in fact, benefit from the agent's bias. Although bias makes it harder to enforce action–state matching when the agent works, it simultaneously increases the agent's motivation to work and to demonstrate that the state is R. The benefit comes from two sources: first, the agent's reward from shirking and choosing R becomes stronger, which allows the principal to sustain working over a larger interval $[u_f, u_e]$; and second, the principal's concern about shirking diminishes, since she can lower v^L.

Before characterizing continuation utilities, we first discuss their role in shaping incentives. Continuation payoffs serve as the key incentive lever in the absence of transfers. When combined with verification, they allow the principal to reward compliance and punish shirking or mismatching the action and the state.

Decoupling of Incentives. Using continuation utilities, the principal has two distinct instruments for providing incentives: continuation utilities when verification occurs and continuation utilities when verification does not occur. The former induces the agent to work, while the latter encourages matching the action with the state. We explain this result in the following.

If the agent deviates by shirking whenever he is supposed to work, he obtains the same utility from taking L or R (See Lemma 3 in Appendix). That is,

$$u(L) + (1 - v^L)\delta u_{\neg V}^L = u(R) + (1 - v^R)\delta u_{\neg V}^R. \tag{I}$$

The intuition behind this lemma is as follows: if shirking and taking a yields less utility than taking a', then a mean-preserving contraction with respect to u_V^a (in the case of verification) and $u_{\neg V}^a$ (in the case of no verification) would benefit the principal. Narrowing the gap between continuation utilities increases the principal's payoff without affecting the agent's incentives.[30]

The above observation allows us to replace (M), without loss of generality, with the stronger condition (I). Furthermore, (WM) can be replaced by

$$\rho v^R \delta u_V^R + (1 - \rho)v^L \delta u_V^L \geq c_A. \tag{W}$$

Thus, the cost of working (c_A) must be compensated by the probability of verification and the continuation utility in the case of verification. Conditions (I) and (W) together imply that the agent's incentives are decoupled. The incentive

---

[30]Note that if $u_V^a = u_{\neg V}^a$, then the principal cannot implement a mean-preserving contraction between continuation utilities. However, she can reduce the cost of verification by decreasing the probability of verification v^a.

to work is guaranteed through the probability of verification and the continuation utility that follows verification. In contrast, the incentive to match the action with the state is supported by the probability of not verification and the continuation utility in the absence of verification.[31]

The upper boundary is linearly increasing in u for $u \in [0, u_f]$, concave for $u \in [u_f, u_e]$, and linearly decreasing in $u \in [u_e, \overline{u}]$ (see Claims 1 and 5 in Appendix).[32]

Due to linearity of the upper boundary for low and high utility, a straightforward approach for implementing the equilibrium along these segments is to randomize over the segment's extreme points, assigning weights proportional to the delivered utility.

Theorem 2 characterizes continuation utilities for $u \in \{[u_f, u_e] \cup \overline{u}\}$.[33] The result highlights a crucial threshold, $\tilde{u}$, that separates regimes in which the principal verifies after both actions from those where verification after R suffices.

Theorem 2 (Continuation Utilities). *For all $u \in [u_f, u_e]$,*

- $u_V := u_V^R = u_V^L > u_{\neg V}^R$,

- *if $\tilde{u} < u_f$, then, $u_{\neg V}^R < u < u_{\neg V}^L$,*

- *if $\tilde{u} \in [u_f, u_e]$, then for $u \leq \tilde{u}$, $u_{\neg V}^R = u = u_{\neg V}^L$ and for $u > \tilde{u}$, $u_{\neg V}^R < u < u_{\neg V}^L$,*

- *if $\tilde{u} > u_e$, then, $u_{\neg V}^R = u_{\neg V}^L = u$.*

Moreover, at $u = \overline{u}$, the continuation utilities under working are equal to those at u_e, and $u_{\neg W} = u$.

Proof in Appendix.

Theorem 2 describes the structure of continuation utilities. The first takeaway is that when the principal verifies and the agent provides evidence, the continuation utility is independent of his action, i.e., $u_V := u_V^R = u_V^L$. The decoupling of

---

[31]Although incentives are decoupled through u_V and $u_{\neg V}$, they remain connected through the probability of verification v.

[32]Intuitively, firing with probability f scales down the utility by $1 - f$. Due to the multiplicative structure of the principal's payoff with respect to the firing probability, this also proportionally reduces the principal's value by $1 - f$. The concavity of the upper boundary in the intermediate utility region arises from the prd. The principal's preferred equilibrium is achieved at an interior utility $u^* \in [u_f, u_e]$. Decreasing the probability of working (e) in $[u_e, \overline{u}]$ increases the utility linearly and decreases the principal's payoff linearly. Linearity arises for two reasons. First, in the agent's utility, the expected cost of exerting effort is given by ec_A, so the probability of shirking affects the agent's stage-game payoff linearly. Second, since the principal's payoff is scaled by the probability that the agent exerts effort, variations in shirking also induce a linear effect on the principal's continuation value.

[33]The unique equilibrium at $u = 0$ involves firing.

23

incentives implies that continuation utilities after verification matter only for the working constraint (W). A mean-preserving contraction of u_V^L and u_V^R does not affect the incentive constraints. Therefore, by the concavity of the value function, the continuation utilities in the case of verification must be equal.[34] Intuitively, hard information obtained through verification dominates soft information inferred from actions. Moreover, the agent is rewarded whenever the principal verifies and the agent provides evidence; that is, if $v^a > 0$, then $u_V > u_{\neg V}^a$.

The second takeaway is that, in the absence of verification, for continuation utilities lower than $\tilde{u}$, the agent is neither rewarded nor punished for taking L or R; that is, $u_{\neg V}^R = u = u_{\neg V}^L$. We refer to the region $[u_f, \tilde{u}]$ (if it is nonempty) as the evidence-based region, since continuation utilities evolve only based on hard information.

In the evidence-based region, the agent's continuation utility lies below what he could secure by shirking, choosing L, and repeatedly obtaining the flow payoff w. Since the principal cannot detect or punish such deviations without verification, the principal must verify with positive probability, even after the agent selects the less-preferred action. Verification probabilities are the more effective instrument, because of linearity, it gains a constant return. Hence, as long as both verification probabilities are strictly positive, the principal prefers to use differential verification rather than dispersing continuation utilities.

In the evidence-based region, the only binding constraint is (WM), while (M) is slack. From the decoupling of incentives, continuation utilities in the case of no verification are used only to incentivize the agent to match the action with the state. Therefore, when (M) is not binding, distorting them away from the initial utility u is unnecessary.

The third takeaway is that for utilities higher than $\tilde{u}$, the agent is punished for taking his preferred state and rewarded for taking the less-preferred one; that is, $u_{\neg V}^R < u < u_{\neg V}^L$. The relationship is now shaped by both hard (evidence) and soft (unverified) information. We refer to the region $[\tilde{u}, u_e]$ (if it is nonempty) as the hybrid region.

In the hybrid region, once $v^L = 0$ (i.e., for all $u \geq \tilde{u}$), the principal loses access to one of her instruments. At this point, she must rely on spreading out continuation utilities (as a last resort) to induce the agent to match the action to the state when the state is L. The agent must be rewarded for choosing his less-preferred action, and punished otherwise. If $u_{\neg V}^L < u_{\neg V}^R$, then a mean-preserving contraction of $u_{\neg V}^L$ and $u_{\neg V}^R$ relaxes the matching constraint for state L and benefits the principal. Therefore, it must hold that $u_{\neg V}^L \geq u_{\neg V}^R$. Moreover, the envelope theorem implies that the continuation utilities $u_{\neg V}^L$ and $u_{\neg V}^R$ must bracket u. Figure 8 depicts the

---

[34]Note that a mean-preserving contraction of u_V^R and u_V^L weakly increases $\min\{\Pi(u_V^L), \Pi(u_V^R)\}$, and thus does not affect (Ve).

structure of continuation utilities in the verification region.

Figure 8: Continuation utilities in the verification region.

In the reduced-workload region ($[u_e, \overline{u}]$), the most efficient adjustment of utilities is achieved by allowing the agent to shirk with different probabilities. Accordingly, the continuation utilities when the agent is supposed to work ($u_V, u_{\neg V}^L, u_{\neg V}^R$) remain constant. Moreover, the continuation utility when the agent is not supposed to work ($u_{\neg W}$) remains undistorted at u.

Agent's preferred equilibrium. In this equilibrium, the agent's utility is $\overline{u}$, the principal's payoff is zero ($\Pi(\overline{u}) = 0$), and the workload is minimized. If the agent is required to work, the game transitions to u_e; when u_e lies in the evidence-based region, the agent may remain indefinitely in the reduced-workload region without returning to verification, whereas if u_e lies in the hybrid region, the agent moves to the verification region with probability one. If the agent is not required to work, then depending on the extent of biases (as discussed in Theorem 1), the agent chooses either L or R. In the following period, after not working, the continuation utility stays at $u_{\neg W}$. Figures 9 displays the continuation utilities at $\overline{u}$.

Figure 9: Continuation utilities at $\overline{u}$.

Let u^* denote the utility level of the agent that maximizes the principal's equilibrium payoff, i.e., $u^* = \arg\max_u \Pi(u)$. Note that u^* always lies in the verification region, $u^* \in [u_f, u_e]$, because the value function is strictly increasing in the firing region and strictly decreasing in the verification region. Intuitively, the relationship begins when the agent complies—works, learns the state, and matches the action with the state. This raises the question of whether the principal ever uses firing or a reduced workload. If verification is more efficient than firing or a reduced workload—because the verification cost c_p is small and the loss from firing or reduced workload π is large—should the principal still employ these less efficient instruments? Equivalently, can the equilibrium path transition from the verification region to the firing or reduced-workload regions? Theorem 3 addresses these questions.

Theorem 3 (Dynamics). *In the principal's preferred equilibrium,*

1. *The principal uses reduced-workload with positive probability.*

2. *The principal fires the agent with probability one if and only if the evidence-based region is empty.*

The first part of Theorem 3 states that allowing the agent to shirk is always a valuable instrument, even when the cost of verification is small and the loss from not working (and thus potentially mismatching the action with the state) is large.

By promising a lighter workload in the future, the principal can lower verification costs across multiple periods leading up to the present. As explained in Section 4, we refer to this phenomenon as *snowballing*.

In the evidence-based region, the agent's reward for complying takes the form of an increase in continuation utility ($u_V > u$). Such rewards move the relationship either into the hybrid region or directly into the reduced-workload region. In the hybrid region, compliance is rewarded through continuation utility when the agent chooses his less-preferred action, L, i.e., $u^L_{\neg V} > u$, which leads with positive probability to the reduced-workload region.[35]

The second part of Theorem 3 asserts that if the evidence-based region is empty, verification and reduced workload are insufficient tools, and the principal must rely on firing along the equilibrium path. The threat of future firing complements verification and lowers verification costs across multiple periods. Hence, snowballing—through the promise of future punishment—renders the use of firing unavoidable.

In the hybrid region, the principal punishes the agent—despite compliance—through continuation utilities when verification does not occur and the agent selects his preferred action, i.e., $u^R_{\neg V} < u$. Such punishments shift the relationship toward the firing region if the evidence-based region is empty, in which case the agent is eventually fired with probability one. By contrast, if the evidence-based region is nonempty, the principal verifies both actions with positive probability to detect shirking, and firing does not occur. Greater verification substitutes for firing and prevents its snowballing effect.

When the bias is large ($(1 - \rho)b > c_A$), the evidence-based region is empty. If the benefit from the agent's bias ($(1 - \rho)b$) exceeds the cost of working (c_A), the agent prefers to comply rather than shirk and choose L, i.e., $w + (1 - \rho)b - c_A > w$. Otherwise, (M–L) does not bind, and the evidence-based region is nonempty. In this case, the evidence-based region serves as a barrier to firing.

---

[35]Except when (Ve) binds, $\Pi'(u_V) \leq \Pi'(u)$, which implies $u_V > u$ if $\Pi(\cdot)$ is strictly concave at u. Therefore, even in the hybrid region, the agent might also be rewarded through continuation utility if verification takes place.

Principal's preferred equilibrium. Employing Theorems 1–3, we describe the dynamics in the principal's preferred equilibrium as follows. As discussed earlier, the relationship begins in the verification region, i.e., $u^* \in [u_e, u_f]$. We consider three cases.

First, suppose the evidence-based region is empty ($\tilde{u} < u_f$). In this case, the entire verification region falls within the hybrid region. If the agent chooses L, the principal does not verify and the continuation utility increases, i.e., $u^L_{\neg V} > u$. The relationship may remain in the verification region, but with a lower probability of verification, or it may transition to the reduced-workload region. If it transitions to the reduced-workload region, the game may move either to u_e or to the agent's preferred equilibrium, $\overline{u}$.[36] Once the relationship enters the reduced-workload region, however, it ultimately returns to the hybrid region with probability one.

If the agent instead chooses R, the principal verifies with positive probability. If verification does not occur, the agent is punished through a reduction in continuation utility, i.e., $u^R_{\neg V} < u$. The relationship then either transitions to the firing region or remains in the verification region but with a higher probability of verification. Ultimately, the agent is fired with probability one. Figure 10 illustrates these transitions when the evidence-based region is empty.

Figure 10: Equilibrium-path transitions between regions when the evidence-based region is empty.

Second, suppose the hybrid region is empty ($\tilde{u} \geq u_e$). The relationship begins in the evidence-based region, where the principal verifies after both actions, though more frequently following R. In this region, continuation utility increases only when the principal verifies and the agent provides evidence. This, in turn, implies

---

[36]When we say the relationship moves to u, we mean it moves to the game associated with continuation utility u.

that verification in the next period occurs with lower probability for both actions. If verification does not occur, continuation utilities remain unchanged ($u^L \neg V = u^R \neg V = u$), and the probability of verification in the next period equals that of the current period. Following verification, the relationship either remains in the evidence-based region or transitions to the reduced-workload region. Once it moves into the reduced-workload region, it remains there indefinitely. Figure 11 illustrates the equilibrium-path transitions when the hybrid region is empty.

Figure 11: Equilibrium-path transitions between regions when the hybrid region is empty.

Third, suppose both the evidence-based and hybrid regions are nonempty. If $u^* \in [u_f, \tilde{u}]$, the relationship begins in the evidence-based region. After verification, it either remains in this region, transitions to the hybrid region, or moves to the reduced-workload region. Once the relationship leaves the evidence-based region, it cycles between the hybrid and reduced-workload regions but never returns to the evidence-based region. Figure 12 illustrates this case.

If $u^* \in [\tilde{u}, u_e]$, the relationship begins in the hybrid region, and the dynamics follow the same pattern as in the previous case after the transition to the hybrid region. In this third case, firing never occurs.
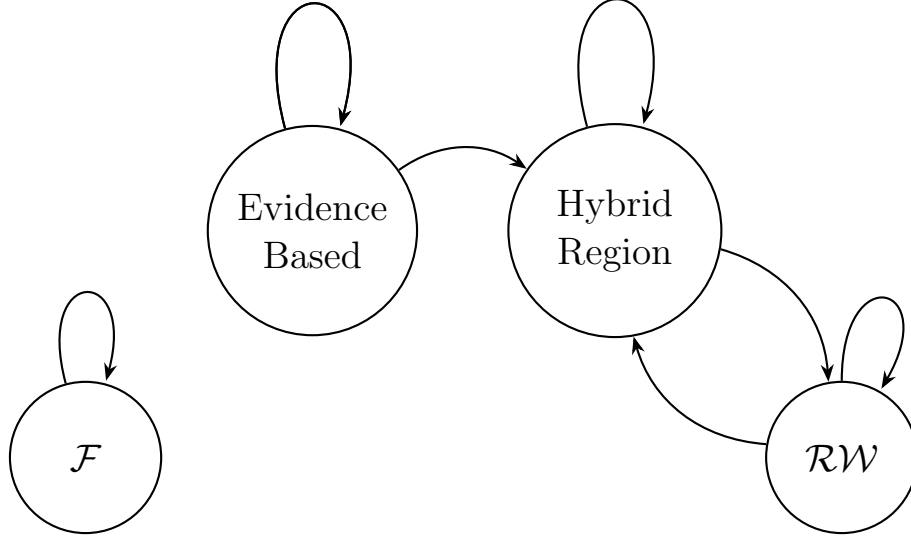
Figure 12: Equilibrium-path transitions between regions when both the hybrid and the evidence based regions are non empty.

5.3 Discussion

Trust Interpretation of the Model and Results. A natural interpretation of our model is to view the complement of verification as trust: the less frequently the principal verifies, the greater the trust in the agent. More precisely, when the agent is required to work and the principal chooses not to verify, we interpret this as an act of trust.

Our results answer the following questions: how often should the principal trust and delegate decisions to the agent, and how does trust evolve over time? How is the evolution of trust related to the agent's bias?

The results show that trust builds gradually. In the evidence-based trust region, the principal never fully trusts the agent's actions, even when he chooses the less-preferred action. Nevertheless, she places relatively more trust in that case. In this region, trust increases only after the agent provides evidence—that is, the probability of verification decreases in subsequent periods.

In the hybrid trust region, the principal fully trusts the agent when he takes L, and trust increases in the following period. By contrast, if the agent takes R and verification does not occur, trust decreases. In this region, trust is shaped by both hard information (evidence obtained through verification) and soft information (the agent's actions when not verified).

At the lowest level of trust—when verification is most frequent—the relationship is fragile and may terminate. At the highest level of trust—when verification is

30

rare—the agent benefits from a reduced workload.

Commitment. When the principal is assumed to have commitment power, the results remain largely the same. In this case, the principal's problem coincides with $\mathcal{P}$ except that the two constraints (Ve) and (Fi) are no longer required, since the principal's own incentives need not be considered. Theorems 1, 2, and 3 continue to hold, with the exception that the highest continuation utility is $\frac{w+b}{1-\delta}$, and the reduced-workload region may become piecewise linear (with a kink at $\frac{w}{1-\delta}$), in contrast to the linear form that arises in the no-commitment case.

Transfer. Suppose that, instead of the fixed transfer w, the principal can use a flexible and positive transfer $T \geq 0$ at the end of each period. Since transfers are a more efficient instrument, the principal never relies on reduced workload or firing on the equilibrium path. The upper boundary is linear and the principal's preferred equilibrium is stationary.

Denote T_V^a as the transfer when the agent takes action a and the principal verifies, and $T_{\neg V}^a$ as the transfer when the agent takes action a and the principal does not verify. As in the model without transfers ($u_V = u_V^L = u_V^R$), it is straightforward to show that $T_V := T_V^L = T_V^R$.

The decoupling of incentives continues to hold. When verification occurs, transfers provide incentives for working, i.e., $(\rho v^R + (1-\rho)v^L)T_V = c_A$. When verification does not occur, transfers provide incentives to match the action with the state, i.e., $(1 - v^R)T_{\neg V}^R = b + (1 - v^L)T_{\neg V}^L$. The value function is linear, it follows that $T_{\neg V}^R = 0$, while the transfer when the agent takes L and the principal does not verify compensates for the agent's bias: $(1 - v^L)T_{\neg V}^L = b$.

The principal's problem is therefore equivalent to minimizing her cost by selecting transfers and verification probabilities:

$$\rho(1 - v^L)T_{\neg V}^L + (\rho v^L + (1-\rho)v^R)T_V + c_P(\rho v^R + (1-\rho)v^R).$$

The first two terms are fixed by the incentive constraints. Hence the principal minimizes the last term subject to her own incentive constraints:

$$\text{verification:} \quad \delta\Pi - T_V - c_P \geq 0, \qquad \text{firing:} \quad \delta\Pi - T_{\neg V}^L \geq 0.$$

The verification constraint always binds; otherwise one could reduce v^R and increase T_V. Hence, the transfer after verification equals the principal's entire future payoff minus the current cost of verification. It follows that

$$\rho v^L + (1-\rho)v^R = \frac{c_A}{\delta\Pi - c_P}.$$

Therefore, Π satisfies

$$\Pi = \pi - b - c_A - c_P \, \frac{c_A}{\delta\Pi - c_P}.$$

Finally, two conditions must be satisfied. First, the expected probability of verification must not exceed one, $\rho v^R + (1 - \rho)v^L \le 1$, or, equivalently (using the verification incentive constraint), $\delta\Pi \ge c_A + c_P$. Second, the firing constraint requires $\delta\Pi(\delta\Pi - c_A - c_P) \ge (\delta\Pi - c_P)\rho b$. For sufficiently large π, both conditions are satisfied.

The agent's bias is toward the more likely state. Suppose the agent is biased toward the more likely state, i.e., $u(L) = w + b$ and $u(R) = w$. In this case, the preferences of the principal and the agent are congruent in terms of the action that should be taken when the state is unknown: both prefer L. Consequently, when the agent is expected to shirk (in the reduced-workload region), he chooses L. Aside from this difference, the structure of the principal's preferred equilibrium remains unchanged, and Theorems 1, 2, and 3 continue to apply.

The agent is allowed to mismatch the action with the sate when he works. We have focused so far on equilibria in which the agent matches the action to the state whenever he works. However, the principal may benefit from deliberately inducing a mismatch between the action and the state (using a prd). Allowing the agent to choose his preferred action R when the state is L relaxes his incentive constraint for truth-telling after working and learning the state. This instrument generates an additional linear region on the upper boundary, located between the verification and reduced-workload regions, with slope $\frac{-\pi}{b}$.

Similarly, requiring the agent to choose L when the state is R can serve as a punishment device. This creates a linear region between the verification and firing regions, with slope $\frac{\pi}{b}$.[37] Apart from these modifications, the structure of the principal's preferred equilibrium remains unchanged, and Theorems 1, 2, and 3 continue to apply.

---

[37]Note that depending on parameters these two regions might not exist.

6 Conclusion

We studied an infinite-horizon principal–agent relationship in which the agent exerts costly effort to learn a privately observed binary state before choosing an action. The principal wants action–state alignment, while the agent is biased toward one action. As a result, the principal faces an adverse selection problem followed by a moral hazard problem. The principal cannot observe payoffs directly but disciplines the agent through verification, workload adjustments, and termination. Continuation utilities summarize history. The upper boundary of the equilibrium payoff set is self-generating, allowing the problem to be formulated as a recursive optimization subject to the objectives of both the agent and the principal.

We showed that equilibrium behavior along the upper boundary of payoffs divides into three regions. Firing region (low continuation utilities), with firing probabilities declining as utility rises. Verification region (intermediate utilities), applied more frequently after the biased action. Reduced-workload region (high continuation utilities), with working probabilities declining as utility increases.

In the evidence-based region, continuation utilities evolve only through hard information from verification. The principal verifies after both actions with positive probability, including when the agent chooses his less-preferred action to ensure effort. Since both probabilities are strictly positive, the principal relies on differential verification rather than distorted continuation utilities. In contrast, in the hybrid region, continuation utilities evolve through both hard information (verification) and soft information (unverified actions). Here, the principal never verifies after L.

We showed that snowballing leads the principal to tolerate shirking even when the loss from shirking is high and the cost of verification is negligible. When the agent's bias is high, the threat of future firing complements verification and reduces verification costs across multiple periods. In contrast, for small biases, the evidence-based region is nonempty, and the principal must verify after both actions with positive probability. This higher level of verification substitutes for firing and prevents its snowballing effect.

References

Dhruva Bhaskar. Dynamic screening and the dual roles of monitoring. *Available at SSRN 2689503*, 2024.

Simon Board. Relational contracts and the value of loyalty. *American Economic Review*, 101(7):3349–3367, 2011.

Liang Dai, Yenan Wang, and Ming Yang. Dynamic contracting with flexible monitoring. *Available at SSRN 3496785*, 2024.

Arash Fahim, Simon Gervais, and R. Vijay Krishna. Equity prices and the dynamics of corporate governance. Working paper, 2021.

Kyna Fong. Evaluating skilled experts: Optimal scoring rules for surgeons. Working paper, 2008.

Yingni Guo and Johannes Hörner. Dynamic allocation without money. 2021.

Marina Halac. Relational contracts and the value of relationships. *American Economic Review*, 102(2):750–779, 2012.

Marina Halac and Andrea Prat. Managerial attention and worker performance. *American Economic Review*, 106(10):3104–3132, 2016.

Jonathan Levin. Relational incentive contracts. *American Economic Review*, 93 (3):835–857, 2003.

Jin Li, Niko Matouschek, and Michael Powell. Power dynamics in organizations. *American Economic Journal: Microeconomics*, 9(1):217–241, 2017.

Elliot Lipnowski and Joao Ramos. Repeated delegation. *Journal of Economic Theory*, 188:105040, 2020.

Qingmin Liu. Information acquisition and reputation dynamics. *The Review of Economic Studies*, 78(4):1400–1425, 2011.

Iván Marinovic, Andrzej Skrzypacz, and Felipe Varas. Dynamic certification and reputation for quality. *American Economic Journal: Microeconomics*, 10(2): 58–82, 2018.

Radoslawa Nikolowa. Motivate and select: Relational contracts with persistent types. *Journal of Economics & Management Strategy*, 26(3):624–635, 2017.

Dmitry Orlov. Frequent monitoring in dynamic contracts. *Journal of Economic Theory*, 206:105550, 2022.

Tomasz Piskorski and Mark M. Westerfield. Optimal dynamic contracts with moral hazard and costly monitoring. *Journal of Economic Theory*, 166:242–281, 2016.

Debraj Ray. The time structure of self-enforcing agreements. *Econometrica*, 70(2): 547–582, 2002.

Eilon Solan and Chang Zhao. Dynamic monitoring under resource constraints. *Games and Economic Behavior*, 129:476–491, 2021.

Felipe Varas, Iván Marinovic, and Andrzej Skrzypacz. Random inspections and periodic reviews: Optimal dynamic monitoring. *The Review of Economic Studies*, 87(6):2893–2937, 2020.

Yu Fu Wong. Dynamic monitoring design. *Available at SSRN 4466562*, 2023.

Yishu Zeng. *On Information Policy Design and Strategic Interactions*. PhD thesis, Stanford University, 2022.

7 Appendix

Let G denote the game presented in Section 3. Let $\tilde{\mathcal{E}}$ denote the set of all equilibria of G, and let $\mathcal{E} \subset \tilde{\mathcal{E}}$ be the subset of equilibria in which the agent matches the action to the state whenever the agent works. In Section 7.1 we provide proofs of the results in Section 5, which characterize the principal's preferred equilibrium in $\mathcal{E}$. Section 7.2 characterizes the principal's preferred equilibrium in $\tilde{\mathcal{E}}$.

7.1 Proofs of Results in Section 5

We begin by stating some claims; we then use these claims to establish the proofs of the theorems.Define the Lagrangian associated with problem $\mathcal{P}$. Let f denote the probability of firing, e the probability of working, and $\alpha_{\neg W}^{R}$ the probability of choosing R when the agent is not supposed to work. Using Lemma 3 (iii), the constraints (M) and (WM) can be replaced by (I) and (W). Therefore, the Lagrangian is given by

$$
\begin{aligned}
\mathcal{L} = \ & (1-f)\Bigg[e\Big(\pi - w + (\rho v^{R} + (1-\rho)v^{L})\delta\,\Pi(u_{V}^{R}) + \rho(1-v^{R})\delta\,\Pi(u_{\neg V}^{R}) \\
& \qquad + (1-\rho)(1-v^{L})\delta\,\Pi(u_{\neg V}^{L}) - \big(\rho v^{R} + (1-\rho)v^{L}\big)c_{p}\Big) \\
& \qquad + (1-e)\Big(\pi p + \alpha_{\neg W}^{R}(1-2\rho)\pi + \delta\,\Pi(u_{\neg W})\Big)\Bigg] \\
& - \lambda_{PK}\Bigg[u - (1-f)\Big(e\Big(w + (1-\rho)b + (\rho v^{R} + (1-\rho)v^{L})\delta u_{V} + \rho(1-v^{R})\delta u_{\neg V}^{R} \\
& \qquad + (1-\rho)(1-v^{L})\delta u_{\neg V}^{L}\Big) + (1-e)\big(w + \alpha_{\neg W}^{R}b + \delta u_{\neg W}\big)\Big)\Bigg] \\
& - \gamma\Big((\rho v^{R} + (1-\rho)v^{L})\delta u_{\neg V} - c_{A}\Big) \\
& - \eta\Big(w + \delta v^{L}u_{\neg V}^{L} + \delta(1-v^{L})u_{V} - \big(w + b + \delta(1-v^{R})u_{\neg V}^{R}\big)\Big). \qquad \text{(L)}
\end{aligned}
$$

Since the value function is concave, the firing constraint (Fi) is equivalent to assuming continuation utilities are non-negative and do not exceed $\overline{u}$. Moreover, the constraint (Ve) implies that $u_{V} \in [\underline{u}_{V}, \overline{u}_{V}]$ for some $\underline{u}_{V}, \overline{u}_{V} \in (0, \overline{u})$. Hence, the additional constraints are

$$
f, e, v^{R}, v^{L}, \alpha_{\neg W}^{R} \in [0,1], \quad u_{V} \in [\underline{u}_{V}, \overline{u}_{V}], \quad u_{\neg V}^{R}, u_{\neg V}^{L} \in [0, \overline{u}], \quad u_{\neg W} \in [0, \overline{u}].
$$

Observe that the value function $\Pi(\cdot)$ is concave and therefore locally Lipschitz continuous on the interior of its effective domain. By Rademacher's theorem, any locally Lipschitz function is differentiable almost everywhere with respect to Lebesgue measure. This property is employed in the subsequent analysis. The following claims characterize both the solution to problem $\mathcal{P}$ and the associated value function $\Pi(\cdot)$.

Finally, note that in the linear segments of $\Pi(\cdot)$, the equilibrium can be implemented through randomization over the extreme points of the segment. Accordingly, in the proofs that follow we always assume u is not located in the interior of a linear segment of $\Pi(\cdot)$.

Claim 1 (Firing). *There exists $u_f \in \left[w + b, \frac{w+b}{1-\delta} \right]$ such that $f > 0$ if and only if $u < u_f$. Moreover, $\Pi(u) = a_f u$ for some $a_f > 0$ and $f = 1 - \frac{u}{u_f}$ for $u \leq u_f$.*

Proof. If $f = 1$, then the value of $\mathcal{P}$ is zero, and by equation (PK) we obtain $u = 0$. Suppose $f < 1$. The first-order condition with respect to f implies

$$\frac{\Pi(u)}{1-f} + \lambda_{PK} \frac{u}{1-f} = 0,$$

which is equivalent to $\Pi(u) + \lambda_{PK} u = 0$ for $f \in (0, 1)$. Furthermore, if $\Pi(u) + \lambda_{PK} u > 0$, then $f = 0$. Define

$$u_f := \inf\{u \mid \Pi(u) - \Pi^+(u)u > 0\},$$

where $\Pi^+(u)$ denotes the right derivative. By concavity of $\Pi(\cdot)$, it follows that

$$\Pi(u_1) - \Pi^-(u_1)u_1 \;\geq\; \Pi(u_f) - \Pi^-(u_1)u_f \;\geq\; \Pi(u_f) - \Pi^+(u_f)u_f,$$

for all $u_1 > u_f$, where $\Pi^-(\cdot)$ denotes the left derivative. By the Subdifferential Envelope Theorem, we have $\Pi^-(u_1) \geq -\lambda_{PK}$. Thus,

$$\Pi(u_1) + \lambda_{PK}u_1 \;\geq\; \Pi(u_1) - \Pi^-(u_1)u_1 \;>\; 0,$$

which implies that $f = 0$ for all $u_1 > u_f$.

Moreover, optimality with respect to f requires

$$\Pi(u_2) - \Pi^+(u_2)u_2 \;\geq\; \Pi(u_2) + \lambda_{PK}u_2 \;\geq\; 0,$$

for all $u_2 \leq u_f$. By the definition of u_f, it follows that $\Pi(u_2) - \Pi^+(u_2)u_2 = 0$ for all $u_2 < u_f$. This condition implies that $\Pi(u)$ is differentiable for all $u \leq u_f$. Hence, for all $u_2 < u_f$ we have $\Pi(u_2) - \Pi'(u_2)u_2 = 0$, which yields $\Pi(u) = a_f u$ for some $a_f > 0$.

Since the value function is linear on the interval $[0, u_f]$, it can be implemented as a randomization between the two points 0 (with probability $1 - \frac{u}{u_f}$) and u_f (with probability $\frac{u}{u_f}$). Therefore, $f = 1 - \frac{u}{u_f}$. $\qquad\square$

Claim 2. *For all $u \in [0, \overline{u}]$, $\Pi(u) - u\Pi^+(u) \geq 0$.*

Proof. Suppose, for contradiction, that for some $u \in [0, \overline{u}]$ we have $\Pi(u) - u\Pi^+(u) < 0$. The first-order condition with respect to f implies that if $\Pi(u) + \lambda_{PK}u < 0$, then $f = 1$. This can occur only for $u = 0$; therefore, for all $u > 0$ we have $\Pi(u) + \lambda_{PK}u \geq 0$. The envelope theorem implies that $-\lambda_{PK} \in \partial\Pi(u)$, hence $-\lambda_{PK} \geq \Pi^+(u)$. Therefore,

$$\Pi(u) - \Pi^+(u)u \geq \Pi(u) + \lambda_{PK}u \geq 0.$$

$\qquad\square$

Claim 3 (Continuation utility in case of not working, i.e., $u_{\neg W}$)**.**

　　1 If $u_{\neg W} \leq u$, then either 1.1) $u_{\neg W} = u$ or 1.2) $\Pi^-(u_{\neg W}) = \Pi^+(u)$ and the value function is linear in $[u_{\neg W}, u]$.

　　2 If $u_{\neg W} \geq u$, then either 2.1) $u_{\neg W} = u$ or 2.2) $\Pi^+(u_{\neg W}) = \Pi^-(u)$ and the value function is linear in $[u, u_{\neg W}]$ or 2.3) $u_{\neg W} = \overline{u}$.

Proof. Without loss of generality, we can assume $f = 0$ and $e < 1$. Suppose $u_{\neg W} < \overline{u}$. The first-order condition with respect to $u_{\neg W}$ gives us $-\lambda_{PK} \in \partial\Pi(u_{\neg W})$. By the Subdifferential Envelope Theorem, $-\lambda_{PK} \in \partial\Pi(u)$.

　　Suppose $u \geq u_{\neg W}$. Then, by the concavity of $\Pi(\cdot)$, either 1) $u_{\neg W} = u$ or 2) $\Pi^-(u_{\neg W}) = \Pi^+(u)$ and the value function is linear in $[u_{\neg W}, u]$.

　　Suppose $u \leq u_{\neg W}$. Then, by the concavity of $\Pi(\cdot)$, either 1) $u_{\neg W} = u$ or 2) $\Pi^+(u_{\neg W}) = \Pi^-(u)$ and the value function is linear in $[u, u_{\neg W}]$. Now suppose $u_{\neg W} = \overline{u}$. The first-order condition with respect to $u_{\neg W}$ implies $\Pi'(\overline{u}) \leq -\lambda_{PK}$. $\qquad\square$

Claim 4 (Action in case of not working, i.e., $\alpha_{\neg W}^R$)**.** *There exists $\check{u} \in [u_e, \overline{u}]$ such that for all $u \leq \check{u}$, $\alpha_{\neg W}^R \in [0, 1)$, and for all $u > \check{u}$, $\alpha_{\neg W}^R = 1$. Moreover, $(1 - 2\rho)\pi + \lambda_{PK}b \geq 0$ if $\alpha_{\neg W}^R = 1$, $(1 - 2\rho)\pi + \lambda_{PK}b = 0$ if $\alpha_{\neg W}^R \in (0, 1)$, and $(1 - 2\rho)\pi + \lambda_{PK}b \leq 0$ if $\alpha_{\neg W}^R = 0$.*

Proof. Without loss of generality, we can assume $f = 0$ and $e < 1$. The first-order condition with respect to $\alpha_{\neg W}^R$ gives $(1 - 2\rho)\pi + \lambda_{PK}b \geq 0$ if $\alpha_{\neg W}^R = 1$, $(1 - 2\rho)\pi + \lambda_{PK}b = 0$ if $\alpha_{\neg W}^R \in (0, 1)$, and $(1 - 2\rho)\pi + \lambda_{PK}b \leq 0$ if $\alpha_{\neg W}^R = 0$.

　　By the Subdifferential Envelope Theorem, $-\lambda_{PK} \in \partial\Pi(u)$. Therefore, by concavity of $\Pi(\cdot)$, there exists $\check{u} \in \left[0, \frac{w+b}{1-\delta}\right]$ such that for all $u \leq \check{u}$, $\alpha_{\neg W}^R \in [0, 1)$, and for all $u > \check{u}$, $\alpha_{\neg W}^R = 1$. $\qquad\square$

Claim 5 (Shirking region). $\Pi(\cdot)$ *is linear on the interval* $[u_e, \overline{u}]$. *Moreover,* $\alpha_{\neg W}^R = 1$ *if and only if the slope of the value function is less than* $\frac{\pi(1-2\rho)}{1-\delta}$; *otherwise,* $\alpha_{\neg W}^R = 0$. *In addition, for* $u \in [u_e, \overline{u}]$ *suppose* $\Pi(u)$ *takes the form* $\Pi(u) = au + d$ *for some* $a, d \in \mathbb{R}$. *If* $a \geq \frac{\pi(1-2\rho)}{1-\delta}$, *then* $a\left(\frac{w}{1-\delta}\right) + d = \frac{\rho\pi - w}{1-\delta}$, *and if* $a \leq \frac{\pi(1-2\rho)}{1-\delta}$, *then* $a\left(\frac{w+b}{1-\delta}\right) + d = \frac{(1-\rho)\pi - w}{1-\delta}$.

Proof. Without loss of generality, we can assume $f = 0$ and $e \in (0, 1)$. Let

$$A = \mathbb{E}\Big[\pi(\theta, a) - c_P\,\mathbb{1}_V - w + \delta\,\Pi(u') \mid W\Big] - \mathbb{E}\Big[\pi(\theta, a) - c_P\,\mathbb{1}_V - w + \delta\,\Pi(u') \mid \neg W\Big],$$

$$B = \mathbb{E}\big[u(a) - c_A\,\mathbb{1}_W + \delta\,u' \mid W\big] - \mathbb{E}\big[u(a) - c_A\,\mathbb{1}_W + \delta\,u' \mid \neg W\big].$$

The first-order condition with respect to e implies $A + \lambda_{PK}B = 0$. Using the definition of the value function,

$$\Pi(u) - \mathbb{E}\Big[\pi(\theta, a) - c_P\,\mathbb{1}_V - w + \delta\,\Pi(u') \mid \neg W\Big] = eA.$$

In addition, by (PK),

$$e = \frac{u - \mathbb{E}\big[u(a) - c_A\,\mathbb{1}_W + \delta\,u' \mid \neg W\big]}{B}.$$

Therefore, by substituting e and using the first-order condition,

$$\Pi(u) - (\rho\pi + \alpha_{\neg W}^R(1 - 2\rho)\pi - w + \delta\Pi(u_{\neg W})) = -\lambda_{PK}(u - w - \alpha_{\neg W}^R b - \delta u_{\neg W}). \tag{1}$$

Suppose $\Pi(\cdot)$ is differentiable at u. Then, by the envelope theorem, $\Pi'(u) = -\lambda_{PK}$. Using the above equation, we consider two cases.

First, consider $u \leq \check{u}$. Using Claim 4, for all $u \leq \check{u}$ the above equation simplifies to

$$\Pi(u) - (-w + \rho\pi + \delta\Pi(u_{\neg W})) = \Pi'(u)\,(u - w - \delta u_{\neg W}).$$

By Claim 3, either $\Pi'(u) = \frac{\Pi(u_{\neg W}) - \Pi(u)}{u_{\neg W} - u}$, or $u_{\neg W} = u$, or $u_{\neg W} = \overline{u}$.

First, suppose $u_{\neg W} \neq \overline{u}$. Then

$$\Pi(u) + w - \rho\pi = \Pi'(u)\,(u - w) + \delta\Pi(u) - \delta u\Pi'(u).$$

This is a first-order ODE, and it is straightforward to show that $\Pi(\cdot)$ is affine. Suppose $\Pi(u)$ takes the form $\Pi(u) = a_1 u + d_1$ for $a_1, d_1 \in \mathbb{R}$, where $a_1\left(\frac{w}{1-\delta}\right) + d_1 = \frac{\rho\pi - w}{1-\delta}$.

Now suppose $u_{\neg W} = \overline{u}$. Then

$$\Pi(u) + w - \rho\pi = \Pi'(u)\,(u - w - \delta\overline{u}).$$

This is a first-order ODE, and again it is straightforward to show that $\Pi(\cdot)$ is affine. Suppose $\Pi(u)$ takes the form $\Pi(u) = a_2 u + d_2$ for $a_2, d_2 \in \mathbb{R}$, where

$$a_2 \tfrac{w}{1-\delta} + d_2 = \tfrac{\rho\pi - w}{1-\delta} - \tfrac{\delta(a_2 \overline{u} + d_2)}{1-\delta}.$$

Therefore, in both cases $\Pi(\cdot)$ is affine. Hence, for all $u \in [u_e, \check{u}]$, $\Pi(\cdot)$ is piecewise linear. However, it is straightforward to show it cannot have more than two linear parts. Formally, there exists $u_1 \in [u_e, \check{u}]$ such that $\Pi(u) = a_1 u + d_1$ for $u \in [u_e, u_1]$, and for all $u \in [u_1, \check{u}]$, $\Pi(u) = a_2 u + d_2$. Moreover, by Claim 4, $a_i \geq \tfrac{(1-2\rho)\pi}{b}$, and since $a_2 \overline{u} + d_2 \geq 0$, in both cases $a_i \tfrac{w}{1-\delta} + d_i \leq \tfrac{\rho\pi - w}{1-\delta}$ for $i \in \{1, 2\}$.

If $\check{u} = \overline{u}$, then $a_2 \overline{u} + d_2 = 0$. Since $\Pi(\cdot)$ is concave and both affine functions pass through the point $\left(\tfrac{w}{1-\delta}, \tfrac{\rho\pi - w}{1-\delta}\right)$, these two affine functions must coincide. Hence, for all $u \in [u_e, \overline{u}]$, $\Pi(u)$ takes the form $\Pi(u) = au + d$ for $a, d \in \mathbb{R}$, where $a\left(\tfrac{w}{1-\delta}\right) + d = \tfrac{\rho\pi - w}{1-\delta}$.

Second, consider $u \geq \check{u}$. Thus, $\alpha_{\neg W}^R = 1$. Equation 1 simplifies to

$$\Pi(u) - ((1-\rho)\pi - w + \delta\Pi(u_{\neg W})) = \Pi'(u)\,(u - w - b - \delta u_{\neg W}).$$

By Claim 3, either $\Pi'(u) = \tfrac{\Pi(u_{\neg W}) - \Pi(u)}{u_{\neg W} - u}$, or $u_{\neg W} = u$, or $u_{\neg W} = \overline{u}$.

First, suppose $u_{\neg W} \neq \overline{u}$. Then

$$\Pi(u) - \tfrac{(1-\rho)\pi - w}{1-\delta} = \Pi'(u)\left(u - \tfrac{w+b}{1-\delta}\right).$$

Thus, $\Pi(u)$ takes the form $\Pi(u) = a_3 u + d_3$ for $a_3, d_3 \in \mathbb{R}$, where $a_3\left(\tfrac{w+b}{1-\delta}\right) + d_3 = \tfrac{(1-\rho)\pi - w}{1-\delta}$.

Now suppose $u_{\neg W} = \overline{u}$. Then

$$\Pi(u) + w - (1-\rho)\pi = \Pi'(u)\,(u - w - b - \delta\overline{u}).$$

Thus, $\Pi(u)$ takes the form $\Pi(u) = a_4 u + d_4$ for $a_4, d_4 \in \mathbb{R}$, where

$$a_4 \tfrac{w+b}{1-\delta} + d_4 = \tfrac{(1-\rho)\pi - w}{1-\delta} - \tfrac{\delta(a_4 \overline{u} + d_4)}{1-\delta}.$$

Similar to the first case, there exists $u_3 \in [\check{u}, \overline{u}]$ such that $\Pi(u) = a_3 u + d_3$ for $u \in [\check{u}, u_3]$, and for all $u \in [u_3, \overline{u}]$, $\Pi(u) = a_4 u + d_4$. In addition, $a_4 \overline{u} + d_4 = 0$. Therefore, these two affine functions must coincide. Hence, for all $u \in [\check{u}, \overline{u}]$, $\Pi(u)$ takes the form $\Pi(u) = a_3 u + d_3$ for $a_3, d_3 \in \mathbb{R}$, where $a_3\left(\tfrac{w+b}{1-\delta}\right) + d_3 = \tfrac{(1-\rho)\pi - w}{1-\delta}$.

If $\check{u} = u_e$, then for all $u \in [u_e, \overline{u}]$, $\Pi(u)$ takes the form $\Pi(u) = a_3 u + d_3$ for $a_3, d_3 \in \mathbb{R}$, where $a_3\left(\frac{w+b}{1-\delta}\right) + d_3 = \frac{(1-\rho)\pi - w}{1-\delta}$. Moreover, by Claim 4, $a_3 \leq \frac{(1-2\rho)\pi}{b}$.

Finally, suppose $\check{u} \in (u_e, \overline{u})$. We show that this case cannot occur. Combining the first and the second cases, $\Pi(\cdot)$ would be a piecewise linear function with three linear parts. Formally, $\Pi(u) = a_1 u + b_1$ for $u \in [u_e, u_1]$, $\Pi(u) = a_2 u + b_2$ for $u \in [u_1, \check{u}]$, and $\Pi(u) = a_3 u + b_3$ for $u \in [\check{u}, \overline{u}]$. Using the facts that $a_1 \frac{w}{1-\delta} + d_1 = \frac{\rho\pi - w}{1-\delta}$, $a_2 \frac{w}{1-\delta} + d_2 \leq \frac{\rho\pi - w}{1-\delta}$, $a_1, a_2 \geq \frac{(1-2\rho)\pi}{1-\delta}$, $a_3\left(\frac{w+b}{1-\delta}\right) + d_3 = \frac{(1-\rho)\pi - w}{1-\delta}$, and $a_3 \leq \frac{(1-2\rho)\pi}{1-\delta}$, a simple algebra shows that $\Pi(\cdot)$ cannot be concave. This leads to a contradiction. $\quad\square$

Claim 6. *Let $u_1 := \min\{u \mid \Pi^+(u) = \Pi'(\overline{u})\}$. Without loss of generality, the probability of working at $\overline{u}$, denoted by $e \in (0, 1)$, satisfies*

$$\overline{u} = e u_1 + (1 - e)\left(w + \alpha^R_{\neg W} b + \delta \overline{u}\right),$$

and the continuation utility in case the agent is not supposed to work $(u_{\neg W})$ is equal to $\overline{u}$. Moreover, continuation utilities and verification probabilities when the agent is supposed to work are the same as u_1.

Proof. By Claim 5, $\Pi(\cdot)$ is linear for all $u \in [u_e, \overline{u}]$. Hence, $\Pi(\cdot)$ is also linear for all $u \in [u_1, \overline{u}]$. Suppose $\Pi(u) = au + d$ for some $a, d \in \mathbb{R}$ on this interval. Define

$$e = \frac{\overline{u} - \left(w + \alpha^R_{\neg W} b + \delta \overline{u}\right)}{u_1 - \left(w + \alpha^R_{\neg W} b + \delta \overline{u}\right)}.$$

Note that $\overline{u} - \left(w + \alpha^R_{\neg W} b + \delta \overline{u}\right) < 0$ since $\overline{u} < \frac{w + \alpha^R_{\neg W} b}{1-\delta}$. Therefore, $e \in (0, 1)$. Using $\Pi(\overline{u}) = a\overline{u} + d$, $\Pi(u_1) = au_1 + d$, and Claim 4, it follows that

$$\Pi(\overline{u}) = e\Pi(u_1) + (1 - e)\left(-w + \rho\pi + \alpha^R_{\neg W}(1 - 2\rho)\pi + \delta\Pi(\overline{u})\right).$$

$\square$

Claim 7. *Let $u \in [u_f, u_e]$ and suppose constraint (M–L) is slack. Then $u^L_{\neg V} = u^R_{\neg V} = u$.*

Proof. Since constraint (M–L) is not binding, we have $\eta = 0$. The first-order conditions of the Lagrangian with respect to $u^L_{\neg V}$ and $u^R_{\neg V}$ yield

$$-\lambda_{PK} \in \partial\Pi(u^L_{\neg V}), \qquad -\lambda_{PK} \in \partial\Pi(u^R_{\neg V}),$$

and by the envelope theorem, $-\lambda_{PK} \in \partial\Pi(u)$.

If $\Pi(\cdot)$ is strictly concave at u, it follows that $u^L_{\neg V} = u^R_{\neg V} = u$.

Now suppose instead that $\Pi(\cdot)$ is linear on a right neighborhood I^+ of u, so that $\Pi(x) = ax + d$ for $x \in I^+$. The envelope theorem then implies $-\lambda_{PK} \geq a$.

Consider the case $u_{\neg V}^L \geq u$. The derivative of the Lagrangian $\mathcal{L}$ with respect to $u_{\neg V}^L$ is $a + \lambda_{PK}$. If $-\lambda_{PK} > a$, then $u_{\neg V}^L = u$. Suppose instead that $-\lambda_{PK} = a$ and, for contradiction, assume $u_{\neg V}^L > u$. The value function is

$$\Pi(u) = \pi - w + \left(\rho v^R + (1-\rho)v^L\right)\delta\Pi(u_V^R) + \rho(1-v^R)\,\delta\Pi(u_{\neg V}^R)$$
$$+ (1-\rho)(1-v^L)\,\delta\Pi(u_{\neg V}^L) - \left(\rho v^R + (1-\rho)v^L\right)c_p.$$

Since (M–L) does not bind, the left derivative with respect to $u_{\neg V}^L$ implies $\Pi^-(u) = \Pi^-(u_{\neg V}^L)$. Moreover, fixing other variables and lowering $u_{\neg V}^L$ extends the linearity of the value function beyond u, so that u lies in the interior of a linear segment of $\Pi(\cdot)$. However, without loss of generality we assume that u is not in the interior of a linear segment of $\Pi(\cdot)$. This yields a contradiction.

If instead $u_{\neg V}^L < u$, then a symmetric argument using a left neighborhood of u leads to the same contradiction.

The same reasoning applies to $u_{\neg V}^R$. Hence $u_{\neg V}^L = u_{\neg V}^R = u$. $\qquad\square$

Claim 8. *Suppose that, for $u \in [u_f, u_e] \cap [\underline{u}_V, \overline{u}_V]$, (W) does not bind. Then $u_V = u$.*

Proof. The first-order condition of the Lagrangian with respect to u_V implies $-\lambda_{PK} \in \partial\Pi(u_V)$. By the envelope theorem, $-\lambda_{PK} \in \partial\Pi(u)$.

For a contradiction, suppose $u_V \neq u$. Then $\Pi(\cdot)$ must be linear on the interval $[\min\{u_V, u\}, \max\{u_V, u\}]$ with slope $-\lambda_{PK}$. Without loss of generality, assume $u_V > u$. The value function is

$$\Pi(u) = \pi - w + \left(\rho v^R + (1-\rho)v^L\right)\delta\Pi(u_V^R) + \rho(1-v^R)\,\delta\Pi(u_{\neg V}^R)$$
$$+ (1-\rho)(1-v^L)\,\delta\Pi(u_{\neg V}^L) - \left(\rho v^R + (1-\rho)v^L\right)c_p.$$

Since (W) is slack, the left-hand derivative with respect to u_V implies $\Pi^-(u) = \Pi^-(u_V)$. Moreover, holding other variables fixed and lowering u_V extends the linearity of the value function below u, so that u lies in the interior of a linear segment of $\Pi(\cdot)$. However, without loss of generality we assume that u is not in the interior of a linear segment of $\Pi(\cdot)$. This yields a contradiction.

The argument is analogous if $u_V < u$, again leading to a contradiction. $\qquad\square$

Claim 9. *For all $u \in [u_f, u_e]$,*

$$\Pi^+(u) = \rho\Pi^+(u_{\neg V}^L) + (1-\rho)\,\Pi^+(u_{\neg V}^R) \quad \text{and} \quad \Pi^-(u) = \rho\Pi^-(u_{\neg V}^L) + (1-\rho)\,\Pi^-(u_{\neg V}^R).$$

Proof. By the second part of Lemma 3,

$$w + \delta(1 - v^L)u_{\neg V}^L = w + b + \delta(1 - v^R)u_{\neg V}^R.$$

42

Substitute

$$u = w + (1 - v^L)\delta u^L_{\neg V} + (\rho v^L + (1 - \rho)v^R)\delta u_V - c_A, \quad u^R_{\neg V} = \tfrac{-b + \delta(1 - v^L)u^L_{\neg V}}{\delta(1 - v^R)}$$

into the expression

$$\Pi(u) = \pi - w + (\rho v^R + (1 - \rho)v^L)\,\delta\Pi(u_V) + \rho(1 - v^R)\,\delta\Pi(u^R_{\neg V})$$
$$+ (1 - \rho)(1 - v^L)\,\delta\Pi(u^L_{\neg V}) - (\rho v^R + (1 - \rho)v^L)c_p.$$

Taking the right and left derivatives with respect to $u^L_{\neg V}$ yields the stated result. $\square$

Claim 10. *Suppose $v^R > 0$ and $u_V \in (\underline{u}_V, \overline{u}_V)$. Then*

$$\Pi(u_V) - u_V\Pi^-(u_V) - \left[\Pi(u^R_{\neg V}) - u^R_{\neg V}\Pi^+(u^R_{\neg V})\right] \le \tfrac{c_P}{\delta},$$

$$\Pi(u_V) - u_V\Pi^+(u_V) - \left[\Pi(u^R_{\neg V}) - u^R_{\neg V}\Pi^-(u^R_{\neg V})\right] \ge \tfrac{c_P}{\delta}.$$

Proof. First suppose (W) binds. We have $e > 0$. Using (I) and the fact that (W) binds,

$$u^L_{\neg V} = \frac{\tfrac{u}{1-f} - (1 - e)(w + \alpha^R_{\neg W}b + \delta u_{\neg W}) - ew}{e\delta(1 - v^L)},$$

and

$$u^R_{\neg V} = \frac{\tfrac{u}{1-f} - (1 - e)(w + \alpha^R_{\neg W}b + \delta u_{\neg W}) - ew}{e\delta(1 - v^R)}.$$

Now substitute $u^L_{\neg V}$, $u^R_{\neg V}$, and $u_V = \tfrac{c_A}{\delta((1-\rho)v^R + \rho v^L)}$ into the objective

$$(1 - f)\Bigg[e\Big(\pi - w + ((1 - \rho)v^R + \rho v^L)\delta\,\Pi(u_V) + (1 - \rho)(1 - v^R)\delta\,\Pi(u^R_{\neg V})$$

$$+ \rho(1 - v^L)\delta\,\Pi(u^L_{\neg V}) - ((1 - \rho)v^R + \rho v^L)c_p\Big)$$

$$+ (1 - e)\Big(\pi p + \alpha^R_{\neg W}(1 - 2\rho)\pi + \delta\,\Pi(u_{\neg W})\Big)\Bigg].$$

The left and right derivatives with respect to v^R yield the result. Note that the argument is independent of whether ((M–L)) binds, or equivalently, of whether $v^L = 0$.

Now suppose ((M–L)) binds; hence $v^L = 0$ and (W) does not bind. Using (I),

$$u^L_{\neg V} = \frac{\tfrac{u}{1-f} - (1 - e)(w + \alpha^R_{\neg W}b + \delta u_{\neg W}) - e(w + (1 - \rho)v^R\delta u_V - c_A)}{e\delta},$$

43

and

$$u_{\neg V}^R = \frac{\frac{u}{1-f} - (1-e)(w + \alpha_{\neg W}^R b + \delta u_{\neg W}) - e(w + (1-\rho)v^R \delta u_V - c_A)}{e\delta(1 - v^R)}.$$

Now substitute $u_{\neg V}^L$ and $u_{\neg V}^R$ into the objective. Employing Claim 8 and Claim 9,

$$\Pi^+(u_V) = \rho\Pi^+(u_{\neg V}^L) + (1-\rho)\,\Pi^+(u_{\neg V}^R) \quad \text{and} \quad \Pi^-(u_V) = \rho\Pi^-(u_{\neg V}^L) + (1-\rho)\,\Pi^-(u_{\neg V}^R).$$

The left and right derivatives with respect to v^R yield the result. $\qquad\square$

Claim 11. *Suppose $u_f \leq \frac{w}{1-\delta}$. Then for every $u \in \left[u_f, \min\{\frac{w}{1-\delta}, u_e\}\right]$ such that $\frac{c_A}{\delta - 1 + \frac{w+(1-\rho)b}{u}} \in [\underline{u}_V, \overline{u}_V]$, the following statements hold:*

1 $u_{\neg V}^L = u_{\neg V}^R = u$.

2 $v^L = 1 - \frac{1}{\delta} + \frac{w}{\delta u}$ and $v^R = 1 - \frac{1}{\delta} + \frac{w+b}{\delta u}$.

3 $u_V = \frac{c_A}{\delta - 1 + \frac{w+(1-\rho)b}{u}} > u + \Delta$.

Proof. Consider the relaxed problem obtained by dropping constraint (M–L) from the Lagrangian $\mathcal{L}$. A mean-preserving contraction argument in this relaxed problem implies that, without loss of generality, we may impose $u_{\neg V}^L = u_{\neg V}^R$. Let $u_{\neg V} := u_{\neg V}^L = u_{\neg V}^R$ and define $v := \rho v^R + (1-\rho)v^L$.

First, note that constraint (W) binds. Otherwise, reduce v to some $v' < v$ so that (W) binds at v'. Next set the continuation utilities as follows:

- upon verification (which now occurs with probability v'), set the continuation utility to u_V;

- upon no verification (probability $1 - v'$), assign $u_{\neg V}$ with probability $1 - v$ and assign u_V with the remaining probability $v - v'$.

This modification leaves the distribution of continuation utilities—and hence the agent's expected utility—unchanged, while strictly increasing the principal's objective because the verification probability falls. Therefore (W) must bind at the optimum.

(1) Since (M–L) is slack and by Claim 7, it follows that $u_{\neg V} = u$.

(2) Using (PK) together with (I), we obtain

$$v^L = 1 - \frac{1}{\delta} + \frac{w}{\delta u} \qquad \text{and} \qquad v^R = 1 - \frac{1}{\delta} + \frac{w+b}{\delta u}.$$

(3) Because (W) binds, and using part (2), we have

$$u_V = \frac{c_A}{\delta - 1 + \frac{w+(1-\rho)b}{u}} \in [\underline{u}_V, \overline{u}_V].$$

We first show $u_V \geq u$. Suppose, to the contrary, that $u_V < u$. Dual feasibility (KKT) then implies $\gamma \leq 0$. The first-order condition with respect to u_V yields

$$-\lambda_{PK} + \gamma \in \partial\Pi(u_V),$$

and since $u_V \in [\underline{u}_V, \overline{u}_V]$, constraint (Ve) is slack.[38] By concavity of Π on $[\underline{u}_V, \overline{u}_V]$, together with $u_V < u_{\neg V} = u$, we conclude $\gamma \leq 0$ and thus $\gamma = 0$. Hence either $u_V = u_{\neg V} = u$ (contradicting $u_V < u$), or

$$\frac{\Pi(u_V) - \Pi(u_{\neg V})}{u_V - u_{\neg V}} = -\lambda_{PK}.$$

The first-order condition with respect to v in $\mathcal{L}$ is

$$\big(\Pi(u_V) - \Pi(u_{\neg V})\big) + \lambda_{PK}(u_V - u_{\neg V}) - \gamma u_V - \frac{c_P}{\delta} = 0.$$

Combining the last two displays yields $\gamma u_V - \frac{c_P}{\delta} = 0$. Since $\gamma = 0$, this implies $\frac{c_P}{\delta} = 0$, a contradiction. Therefore $u_V \geq u$.

We now show: there exists $\Delta > 0$ such that $u_V \geq u + \Delta$. Suppose not. Then there exists $\check{u}$ such that either $u_V(\check{u}) = \check{u}$, or there is a sequence $\{\tilde{u}_i\}$ with $\tilde{u}_i \downarrow \check{u}$ and $u_V(\tilde{u}_i) \to \check{u}$. In either case, using $u_{\neg V} = u$, one obtains

$$\check{u} = \frac{w + (1-\rho)b - c_A}{1 - \delta}.$$

By Claim 10 and $u_{\neg V} = u$, for all $u > \check{u}$,

$$-\big(\Pi(u) - u\Pi^-(u)\big) + \big(\Pi(u_V) - u_V\Pi^+(u_V)\big) \geq \frac{c_P}{\delta}.$$

Consider a decreasing sequence $\{u_i\}$ with $u_{i+1} = u_V(u_i)$ and $u_i \downarrow \check{u}$. Since

$$\Pi(u_i) - u_i\Pi^+(u_i) \leq \Pi(u_{i+1}) - u_{i+1}\Pi^-(u_{i+1}),$$

we obtain, for all i,

$$-\big(\Pi(u_i) - u_i\Pi^+(u_i)\big) + \big(\Pi(u_{i+2}) - u_{i+2}\Pi^+(u_{i+2})\big) \geq \frac{c_P}{\delta}.$$

---

[38] Because $u_V \in [\underline{u}_V, \overline{u}_V]$, (Ve) does not bind.

By Claim 2, $\Pi(u) - u\Pi^+(u) \geq 0$ for all $u \in [0, \overline{u}]$. Hence

$$0 \ \geq \ \Pi(u_{2i}) - u_{2i}\Pi^+(u_{2i}) \ \geq \ \Pi(u_1) - u_1\Pi^+(u_1) - \frac{nc_p}{\delta},$$

which yields a contradiction for n large enough. Therefore $u_V > u$, and thus $u_V \geq u + \Delta$ for some $\Delta > 0$.

Finally, since the solution to the relaxed problem satisfies (M–L) and (M–R), it is also a solution to the original problem $\mathcal{P}$. $\qquad\square$

Claim 12. *If (M–L) is slack, then (W) binds.*

Proof. If (M–L) is slack, then $v^L > 0$. Suppose, for contradiction, that (W) does not bind. Decrease v^L while keeping the distribution of continuation utilities the same. Let the reduced value be $\tilde{v}^L$. With probability $\tilde{v}^L$, the continuation utility is u_V^L. With probability $v^L - \tilde{v}^L$, the continuation utility is u_V^L. With probability $1 - v^L$, the continuation utility is $u_{\neg V}^L$. The distribution of continuation utilities remains unchanged. Since (M–R) is slack, this change does not affect (M–R). It does not affect (M–L), (Ve), or (W), and it keeps the agent's utility constant while increasing the principal's payoff. This yields a contradiction. $\qquad\square$

Claim 13. *The constraint $u_V \geq \underline{u}_V$ is always slack.*

Proof. Consider the relaxed problem obtained by ignoring the constraint $u_V \geq \underline{u}_V$. We show that, in this relaxed problem, $\underline{u}_V < u_f$. Since $u_V \geq u$ for all $u \in [u_f, \min\{u_e, \overline{u}_V\}]$, it follows that in the relaxed problem $u_V \geq \underline{u}_V$.

Suppose, for contradiction, that for some $u \geq u_f$ we have $u_V < \underline{u}_V$. Claim 10 implies that

$$\Pi(u_V) - u_V\Pi^+(u_V) - \left[\Pi(u_{\neg V}^R) - u_{\neg V}^R\Pi^-(u_{\neg V}^R)\right] \geq \frac{c_P}{\delta}.$$

By Claim 2, for all u we have $\Pi(u) - u\Pi^-(u) \geq 0$. Hence,

$$\Pi(u_V) - u_V\Pi^+(u_V) \geq \frac{c_P}{\delta}.$$

However, we know $\Pi(u_V) \leq \Pi(\underline{u}_V) = \frac{c_P}{\delta}$, and since $\underline{u}_V$ lies on the increasing part of $\Pi(\cdot)$, we have $\Pi^+(u_V) \geq 0$, yielding a contradiction. $\qquad\square$

Claim 14. *If (Ve) binds, then (M–L) binds.*

Proof. Suppose, for contradiction, that (M–L) does not bind, and therefore $v^L > 0$. By Claim 12, constraint (W) binds, and by Claim 7, we have $u_{\neg V}^R = u_{\neg V}^L = u$. By (I),

$$v^L = 1 - \tfrac{1}{\delta} + \tfrac{w}{\delta u} \quad \text{and} \quad v^R = 1 - \tfrac{1}{\delta} + \tfrac{w+b}{\delta u}.$$

Since (W) binds, u_V can be expressed as a function of u. Therefore

$$u_V = \frac{c_A}{\delta - 1 + \frac{w+(1-\rho)b}{u}} = \bar{u}_V.$$

Hence, except possibly at a particular knife-edge value of u, this yields a contradiction. Therefore (M–L) must bind. $\qquad\square$

Claim 15. *Suppose $v^R > 0$, and (W) does not bind. If $(1-\rho)u_V \geq u^R_{\neg V}$, then*

$$(u_V - \tfrac{u^R_{\neg V}}{1-\rho})\Pi^+(u) = \Pi(u_V) - \Pi(u^R_{\neg V}) - \tfrac{c_P}{\delta} - \tfrac{\rho}{1-\rho}u^R_{\neg V}\Pi^-(u^L_{\neg V}),$$

$$(u_V - \tfrac{u^R_{\neg V}}{1-\rho})\Pi^-(u) = \Pi(u_V) - \Pi(u^R_{\neg V}) - \tfrac{c_P}{\delta} - \tfrac{\rho}{1-\rho}u^R_{\neg V}\Pi^+(u^L_{\neg V}).$$

If $(1-\rho)u_V \leq u^R_{\neg V}$, then the direction reverses at u: in the preceding equations, replace $\Pi^+(u)$ with $\Pi^-(u)$ and replace $\Pi^-(u)$ with $\Pi^+(u)$.

Proof. We have $e > 0$. Using (I),

$$u^L_{\neg V} = \frac{b + \delta(1-v^R)u^R_{\neg V}}{\delta(1-v^L)}.$$

By the promise-keeping constraint,

$$u = (1-f)\Big[e\Big(w + (1-\rho)b + ((1-\rho)v^R + \rho v^L)\delta u_V + (1-\rho)(1-v^R)\delta u^R_{\neg V}$$

$$+ \rho(1-v^L)\delta u^L_{\neg V}\Big) - (1-e)\big(w + \alpha^R_{\neg W}b + \delta u_{\neg W}\big)\Big].$$

Now substitute u and $u^R_{\neg V}$ into the expression

$$\Pi(u) = (1-f)\Big[e\Big(\pi - w + ((1-\rho)v^R + \rho v^L)\delta\,\Pi(u_V) + (1-\rho)(1-v^R)\delta\,\Pi(u^R_{\neg V})$$

$$+ \rho(1-v^L)\delta\,\Pi(u^L_{\neg V}) - \big((1-\rho)v^R + \rho v^L\big)c_p\Big)$$

$$+ (1-e)\Big(\pi p + \alpha^R_{\neg W}(1-2\rho)\pi + \delta\,\Pi(u_{\neg W})\Big)\Big].$$

The left and right derivatives with respect to v^R yield the result. $\qquad\square$

Claim 16. *Suppose (M–L) binds and W does not bind. Then it cannot be the case that $u = u^L_{\neg V} = u^R_{\neg V} = u_V$.*

Proof. Suppose $u = u^L_{\neg V} = u^R_{\neg V} = u_V$. Claim 15 implies

$$-\frac{\rho u}{1-\rho}\Pi^+(u) = -\frac{c_P}{\delta} - \frac{\rho u}{1-\rho}\Pi^-(u),$$

$$-\frac{\rho u}{1-\rho}\Pi^-(u) = -\frac{c_P}{\delta} - \frac{\rho u}{1-\rho}\Pi^+(u).$$

There is no $\Pi^+(u)$ and $\Pi^-(u)$ such that solves the above equations. A contradiction.
$\square$

Claim 17. *Suppose (Ve) binds and (W) does not bind. Then $\Pi^-(u) = \Pi^-(\bar{u}_V)$. Hence $[\min\{u, \bar{u}_V\}, \max\{u, \bar{u}_V\}]$ is a linear segment of $\Pi(\cdot)$.*

Proof. By the promise-keeping constraint,

$$u = (1-f)\Big[e\Big(w + (1-\rho)b + ((1-\rho)v^R + \rho v^L)\delta u_V + (1-\rho)(1-v^R)\delta u^R_{\neg V}$$

$$+ \rho(1-v^L)\delta u^L_{\neg V}\Big) - (1-e)\big(w + \alpha^R_{\neg W}b + \delta u_{\neg W}\big)\Big].$$

Now substitute u into the expression

$$\Pi(u) = (1-f)\Big[e\Big(\pi - w + ((1-\rho)v^R + \rho v^L)\delta\,\Pi(u^R_V) + (1-\rho)(1-v^R)\delta\,\Pi(u^R_{\neg V})$$

$$+ \rho(1-v^L)\delta\,\Pi(u^L_{\neg V}) - \big((1-\rho)v^R + \rho v^L\big)c_p\Big)$$

$$+ (1-e)\Big(\pi p + \alpha^R_{\neg W}(1-2\rho)\pi + \delta\,\Pi(u_{\neg W})\Big)\Big].$$

The left derivative with respect to u_V yields the result. $\square$

Claim 18. *Suppose $(1-\rho)b \geq c_A$. Then $u_f \geq \tilde{u}$.*

Proof. Suppose $u \in [u_f, \tilde{u}]$. Then $u^L_{\neg V} = u^R_{\neg V} = u$ and $u_V > u$. By the promise-keeping constraint,

$$u = w + (1-\rho)b + ((1-\rho)v^R + \rho v^L)\delta u_V + (1-\rho)(1-v^R)\delta u^R_{\neg V} + \rho(1-v^L)\delta u^L_{\neg V} - c_A$$

$$> w + (1-\rho)b - c_A + \delta u.$$

Hence $u > \frac{w+(1-\rho)b-c_A}{1-\delta} \geq \frac{w}{1-\delta}$, a contradiction. $\square$

Claim 19. *Suppose $u_e \geq \frac{w}{1-\delta}$. Then for all $u \in [\max\{u_f, \frac{w}{1-\delta}\}, u_e]$,*

1) $v^R > v^L = 0$.

2) *There exists $\Delta > 0$ such that $u^R_{\neg V} + \Delta < u < u^L_{\neg V} - \Delta$.*

Proof. **Step 1: Show $v^L = 0$ (hence $v^R > v^L$).** Suppose, for contradiction, that $v^L > 0$. Then constraint (M–L) does not bind. By Claim 12, (W) binds. Moreover, since (M–L) does not bind, Claim 7 gives $u = u^L_{\neg V} = u^R_{\neg V}$. Using the promise-keeping constraint (PK), we obtain

$$v^L \;=\; 1 - \frac{1}{\delta} + \frac{w}{\delta u}.$$

Because $u > \frac{w}{1-\delta}$, this implies $v^L < 0$, a contradiction. Therefore $v^L = 0$ and (M–L) binds.

Step 2: Show $u^R_{\neg V} < u < u^L_{\neg V}$. We consider two cases according to the multiplier η.

Case A: $\eta = 0$. The derivatives of the Lagrangian $\mathcal{L}$ with respect to $u^L_{\neg V}$ and $u^R_{\neg V}$ imply

$$-\lambda_{PK} \in \partial\Pi(u^L_{\neg V}) \quad\text{and}\quad -\lambda_{PK} \in \partial\Pi(u^R_{\neg V}),$$

and the envelope theorem implies $-\lambda_{PK} \in \partial\Pi(u)$. Since the value function is concave, if $u \neq u^a_{\neg V}$ then $\Pi(\cdot)$ is linear on the interval $[\min\{u, u^a_{\neg V}\}, \max\{u, u^a_{\neg V}\}]$ for $a \in \{L, R\}$. Without loss of generality, we assume u is not in the interior of a linear segment of $\Pi(\cdot)$. Claim 9 then implies $u^L_{\neg V} = u = u^R_{\neg V}$.

If (W) binds, straightforward algebra shows $v^L > 0$ and $u \leq \frac{w}{1-\delta}$, a contradiction; hence (W) does not bind. If (Ve) does not bind, then by Claim 8 we conclude $u^L_{\neg V} = u = u^R_{\neg V} = u_V$, which contradicts Claim 16.

Now suppose (Ve) binds.

- If $(1-\rho)u_V \geq u$, then Claim 15 implies

$$(\overline{u}_V - \tfrac{u}{1-\rho})\Pi^+(u) = -\Pi(u) - \tfrac{\rho u}{1-\rho}\Pi^-(u), \qquad (\overline{u}_V - \tfrac{u}{1-\rho})\Pi^-(u) = -\Pi(u) - \tfrac{\rho u}{1-\rho}\Pi^+(u).$$

  Therefore $\overline{u}_V = \frac{u(1+\rho)}{1-\rho}$ and $\frac{\rho u}{1-\rho}(\Pi^+(u) + \Pi^-(u)) = -\Pi(u)$. Using Claim 17, since $u_V > u$, we have $\Pi^+(u) = \Pi^-(u)$. Hence

$$\Pi'(u) = \frac{-\Pi(u)}{\frac{2\rho u}{1-\rho}} = \frac{\Pi(u) - \frac{c_p}{\delta}}{u - \overline{u}_V} = \frac{\Pi(u) - \frac{c_p}{\delta}}{u - \frac{u(1+\rho)}{1-\rho}},$$

  which yields $\Pi(u) = \Pi(u) - \frac{c_P}{\delta}$, a contradiction.

- If $(1-\rho)u_V < u$, then Claim 15 implies

$$(\overline{u}_V - \tfrac{u}{1-\rho})\Pi^-(u) = -\Pi(u) - \tfrac{\rho u}{1-\rho}\Pi^-(u), \qquad (\overline{u}_V - \tfrac{u}{1-\rho})\Pi^+(u) = -\Pi(u) - \tfrac{\rho u}{1-\rho}\Pi^+(u).$$

49

Therefore $\Pi^-(u) = \Pi^+(u) = \frac{-\Pi(u)}{\overline{u}_V - u}$. By Claim 17, $[\min\{u, \overline{u}_V\}, \max\{u, \overline{u}_V\}]$ is a linear segment of $\Pi(\cdot)$. Hence $\Pi'(u) = \frac{\frac{c_P}{\delta} - \Pi(u)}{\overline{u}_V - u}$. This contradicts the earlier identity $\Pi'(u) = \frac{-\Pi(u)}{\overline{u}_V - u}$. Therefore we may conclude $\eta < 0$.

Case B: $\eta < 0$. The derivative of the Lagrangian $\mathcal{L}$ with respect to $u_{\neg V}^L$ implies $-\lambda_{PK} + \eta \in \partial\Pi(u_{\neg V}^L)$, and with respect to $u_{\neg V}^R$ it implies $-\lambda_{PK} - \eta \in \partial\Pi(u_{\neg V}^R)$. By concavity of Π, we have $u_{\neg V}^R \le u \le u_{\neg V}^L$. Therefore, if $\Pi(\cdot)$ is differentiable at u, it follows that $u_{\neg V}^R < u < u_{\neg V}^L$. Suppose instead that $\Pi(\cdot)$ is not differentiable at u. Consider three cases:

1. $u_{\neg V}^R = u = u_{\neg V}^L = u_V$. If (W) binds, straightforward algebra shows $v^L > 0$ and $u \le \frac{w}{1-\delta}$, a contradiction; hence (W) does not bind. If (Ve) does not bind, Claim 8 implies $u_{\neg V}^L = u = u_{\neg V}^R = u_V$, which contradicts Claim 16. If (Ve) binds, an argument analogous to the case $\eta = 0$ yields a contradiction.

2. $u_{\neg V}^L = u > u_{\neg V}^R$. Claim 9 yields a contradiction.

3. $u_{\neg V}^R = u > u_{\neg V}^L$. Similarly, Claim 9 yields a contradiction.

Thus $u_{\neg V}^R < u < u_{\neg V}^L$.

Step 3: Strict separation by a margin. We now show that there exists $\Delta > 0$ such that $u_{\neg V}^R + \Delta < u < u_{\neg V}^L - \Delta$. Suppose, toward a contradiction and without loss of generality, that there exists a sequence $u_i \to \check{u}$ with $u_{\neg V}^L(u_i) \to \check{u}$. Claim 9 implies $\Pi^-(u_{\neg V}^R(u_i)) \to \Pi^-(\check{u})$ and $\Pi^+(u_{\neg V}^R(u_i)) \to \Pi^+(\check{u})$. Since $\Pi(\cdot)$ is concave, the sequence $u_{\neg V}^R(u_i)$ also converges to $\check{u}$. Consider two cases:

- For a subsequence u_i', constraint (W) binds. Simple algebra then yields $\check{u} \le \frac{w}{1-\delta}$, a contradiction.

- There exists a subsequence u_i'' such that (W) does not bind. Using Claim 15 yields a contradiction.

This completes the proof of both (1) and (2). $\qquad\square$

Proof of Theorem 1:

Proof.

1 It follows from Claim 1.

2 It follows from Claims 11 and 19.

3 It follows from Claims 4 and 5.

$\square$

Proof of Theorem 2:

Proof. It follows from Claims 6, 11 and 19.

$\square$

Proof of Theorem 3:

Proof. It follows from Claims 11, 18 and 19.

$\square$

7.2 Characterization of the Principal's Preferred Equilibrium in $\tilde{\mathcal{E}}$

Lemma 1. *Suppose $f(\cdot) : \mathbb{R} \to \mathbb{R}$ is a continuous and concave function. Let $x, y \in \mathbb{R}$, $z \in \mathbb{R}^+$ and $a \in (0,1)$ such that $x > y$ and $x - \frac{z}{a} \geq y + \frac{z}{1-a}$ Then, the following inequality holds:*

$$af(x) + (1-a)f(y) \leq af(x - \frac{z}{a}) + (1-a)f(y + \frac{z}{1-a})$$

Furthermore, the inequality is strict if $f(\cdot)$ is not linear on $[y, x]$.

Proof. The inequality holds if and only if

$$\frac{f(x) - f(x - \frac{z}{a})}{x - (x - \frac{z}{a})} \leq \frac{f(y + \frac{z}{1-a}) - f(y)}{(y + \frac{z}{1-a}) - y}.$$

Let $\alpha \in (0,1)$ such that $\alpha x + (1-\alpha)y = x - \frac{z}{a}$ and $\beta \in (0,1)$ such that $\beta x + (1-\beta)y = y + \frac{z}{1-a}$. Therefore the inequality in the Lemma holds if and only of

$$\frac{f(x) - f(\alpha x + (1-\alpha)y)}{(1-\alpha)(x-y)} \leq \frac{f(\beta x + (1-\beta)y) - f(y)}{\beta(x-y)}.$$

Since $x > y$, equivalently

$$\beta f(x) + (1-\alpha)f(y) \leq \beta f(\alpha x + (1-\alpha)y) + (1-\alpha)f(\beta x + (1-\beta)y).$$

By concavity of $f(\cdot)$ we know $f(\alpha x + (1-\alpha)y) \geq \alpha f(x) + (1-\alpha)f(y)$ and $f(\beta x + (1-\beta)y) \geq \beta f(x) + (1-\beta)f(y)$. Combining both inequalities gives us the last inequality. In addition, if $f(\cdot)$ is not linear in interval $[y, x]$, then $f(\alpha x + (1-\alpha)y) > \alpha f(x) + (1-\alpha)f(y)$ and $f(\beta x + (1-\beta)y) > \beta f(x) + (1-\beta)f(y)$ which imply the inequality in the statement is a strict inequality. $\square$

51

Lemma 2. *It is without loss of generality to restrict attention to direct, truthful communications in $\tilde{\mathcal{E}}$ with message space $\hat{L}, \hat{R}$.*

Proof. Let G_M be the same game as G but with message space M. We show that for any equilibrium $e_M \in G$, there exists an outcome-equivalent equilibrium $e \in G$.

Fix a history h^{t-1} in e_M. Assume the agent is not fired and is supposed to work after h^{t-1}, i.e., $f^t_\xi = \neg F$ and $w^t_\xi = W$.[39] Assume he sends a message $m^t \in M$.

Define e' to be the same as e_M, except that after $\{h^{t-1}, \neg F, W\}$ the message space is restricted to $\{\hat{L}, \hat{R}\}$ and the agent reveals the state. We show that e' is an equilibrium. Consider two cases:

1. **Separating messages in e_M:** If the agent plays a separating strategy when sending messages m^t in e_M, then truthful direct communication in e' is clearly outcome-equivalent to e_M.

2. **Pooling messages in e_M:** If m^t is pooling in e_M, suppose instead that the agent reveals the state in e'. Continuation strategies in e' following revelation are defined to coincide with continuation strategies in e_M after sending m^t. The principal's incentive to verify or not verify in e' is the same as in e_M. After the action a is taken, the principal's payoff in e' is identical to that in e_M, since the action itself is sunk. Moreover, in the continuation game both players face the same incentives in e' and e_M. Because the agent has already taken the action, verifying and learning the state is payoff-irrelevant.

The argument above is not specific to the message space following h^{t-1}; it applies to all message spaces on the equilibrium path. □

7.2.1 Principal's Preferred Equilibrium of $\tilde{\mathcal{E}}$

Our objective in this section is to characterize the principal's preferred equilibrium within $\tilde{\mathcal{E}}$. The history of the game can be summarized by the agent's continuation utility. The upper boundary of the equilibrium payoff set is self-generating in the sense of Abreu, Pearce, and Stacchetti (1990). Consequently, the problem can be formulated as a recursive program subject to the incentive constraints of both the agent and the principal.

As noted earlier, the principal does not verify when the agent is recommended not work. Thus, on the equilibrium path, the agent's continuation utility at the end of each period can be classified into three cases:

1 u^{ma}_V, when the agent is supposed to work, he sends message $m \in \{\hat{L}, \hat{R}\}$ and he takes action $a \in \{L, R\}$ and the principal verifies.

---

[39]If the agent is not supposed to work ($w^t_\xi = \neg W$), messages are superfluous.

2 $u_{\neg V}^{ma}$, when the agent is supposed to work, he sends message $m \in \{\hat{L}, \hat{R}\}$ and he takes action $a \in \{L, R\}$ and the principal does not verify.

3 $u_{\neg W}^{a}$, when the agent is recommended not work and the he takes action $a \in \{L, R\}$.[40]

Let v^{ma} be the verification probability when the agent sends message $m \in \{\hat{L}, \hat{R}\}$ and he takes action $a \in \{L, R\}$. Let α_W^{ma} be the probability that the agent is recommended to take action action $a \in \{L, R\}$ when the he sends message $m \in \{\hat{L}, \hat{R}\}$ When he is supposed to work. Let $\alpha_{\neg W}^{Ra}$ be the probability that the agent is recommended to take action action $a \in \{L, R\}$ When he is supposed to shirk. Let $\tilde{\Pi}(u)$ be the principal's highest equilibrium payoff consistent with the agent's equilibrium payoff being u. Let $\tilde{\Pi}(u) = -\infty$ if no equilibrium exists that gives the agent u. Let $\tilde{\bar{u}}$ denote the maximum promise such that $\Pi(u) > -\infty$. Note that by Lemma 2 we can restrict attention to truthful direct mechanisms.

The principal's program is:

$$\tilde{\mathcal{P}}: \quad \Pi(u) = \sup \mathbb{E}\Big[\mathbb{1}_{\neg F}\big(\pi(\theta, a) - c_P \mathbb{1}_V - w + \delta \tilde{\Pi}(u')\big)\Big],$$

subject to the agent's incentive constraints: 1) Working and reporting truthfully the state

$$\mathbb{E}\left[u(a) + v^{\theta a}\delta u_V^{\theta a} + (1 - v^{\theta a})\delta u_{\neg V}^{\theta a}\right] - c_A \geq \max_{m=\{\hat{L}, \hat{R}\}} \{\mathbb{E}\left[u(a) + (1 - v^{ma})\delta u_{\neg V}^{ma} \mid m\right]\}. \tag{WT}$$

2) Conditional on working, reporting truthfully the state

$$\mathbb{E}\left[u(a) + v^{\theta a}\delta u_V^{\theta a} + (1 - v^{\theta a})\delta u_{\neg V}^{\theta a} \mid \theta\right] \geq \mathbb{E}[u(a) + (1 - v^{\theta' a})\delta u_{\neg V}^{\theta' a} \mid \theta'], \tag{T}$$

for all $\theta, \theta' \in \{L, R\}$. Specifically, let (T–L) denote the constraint corresponding to $\theta = L$ and $\theta' = R$, and let (T–R) denote the constraint corresponding to $\theta = R$ and $\theta' = L$. 3) The promise-keeping constraint:

$$u = \mathbb{E}\big[\mathbb{1}_{\neg F}\left(u(a) - c_A \mathbb{1}_W + \delta u'\right)\big]. \tag{PK}$$

Principal's incentive constraints: 1) Firing

$$\mathbb{E}\left[\pi(\theta, a) - c_P\mathbb{1}_V - w + \delta\tilde{\Pi}(u')\right] \geq 0, \tag{Fi}$$

---

[40]It is easy to see messages m is superfluous when the agent does not work.

2) Verifying

$$\delta \tilde{\Pi}(u_V^{\theta a}) - c_P \geq 0, \tag{Ve}$$

for all $a, \theta \in \{L, R\}$. The supremum is taken over the continuation utility u', the probability of firing, the probability of working, the probability of verification and the probability of actions. The continuation utility u' can take the three aforementioned forms, $u_V^{ma}, u_{\neg V}^{ma}, u_{\neg W}^a \in [0, \tilde{\bar{u}}]$.[41],[42]

Lemma 3. *In the principal's preferred equilibrium, without loss of generality:*

1. *The probabilities of verification and the continuation utilities are independent of the action recommendations. Formally, $u_V^{ma} = u_V^{ma'}$, $u_{\neg V}^{ma} = u_{\neg V}^{ma'}$, $u_{\neg W}^a = u_{\neg W}^{a'}$ and $v^{ma} = v^{ma'}$ for $m \in \{\hat{L}, \hat{R}\}$ and $a, a' \in \{L, R\}$.*

2. *If the agent deviates and shirks whenever he is supposed to work, he will obtain the same utility from reporting $\hat{L}$ or $\hat{R}$. Formally,*

$$\mathbb{E}\left[u(a) + (1 - v^{\hat{R}a})\delta u_{\neg V}^{\hat{R}a} \mid \hat{R}\right] = \mathbb{E}\left[u(a) + (1 - v^{\hat{L}a})\delta u_{\neg V}^{\hat{L}a} \mid \hat{L}\right].$$

3. *The continuation utility of the agent after verification is independent of the agent message. Formally $u_V^m = u_V^{m'}$, for $m, m' \in \{\hat{L}, \hat{R}\}$.*

Proof. 1 By Lemma 1, a mean-preserving contraction applied to each pair $(u_V^{ma}, u_V^{ma'})$, $(u_{\neg V}^{ma}, u_{\neg V}^{ma'})$, and $(u_{\neg W}^a, u_{\neg W}^{a'})$ weakly increases the principal's payoff while leaving all incentive constraints unchanged.

Furthermore, since these continuation utilities are independent of messages, a mean-preserving contraction between v^{ma} and $v^{ma'}$ does not alter the principal's payoff, the agent's utility, or any incentive constraint.

2 Using part (1) of the lemma, we adopt the convention of dropping the superscript "a" from continuation utilities and verification probabilities. We proceed the proof by contradiction. Without loss of generality, suppose

$$\mathbb{E}\left[u(a) + (1 - v^{\hat{R}})\delta u_{\neg V}^{\hat{R}} \,\middle|\, \hat{R}\right] < \mathbb{E}\left[u(a) + (1 - v^{\hat{L}})\delta u_{\neg V}^{\hat{L}} \,\middle|\, \hat{L}\right].$$

Then $v^{\hat{R}} > 0$; otherwise, this would contradict constraint (T–R). Now decrease $v^{\hat{R}}$ slightly while keeping the distribution of continuation utilities the same.

---

[41] Expectations are taken with respect to the random outcomes, namely the state, firing status, verification status, working status, and the agent's action.

[42] The agent has another incentive constraint. The agent must follow prd's action recommendation. For now we ignore this constraint and later we show with our assumptions this constraint never binds on the equilibrium path.

Let the reduced value be $\tilde{v}^{\hat{R}}$. With probability $\tilde{v}^{\hat{R}}$, the continuation utility is $u_V^{\hat{R}}$. With probability $v^{\hat{R}} - \tilde{v}^{\hat{R}}$, the continuation utility is $u_V^{\hat{R}}$. With probability $1 - v^{\hat{R}}$, the continuation utility is $u_{\neg V}^{\hat{R}}$. Therefore, the distribution of continuation utilities remains unchanged. Since (T–L) is slack, this change does not affect that constraint. It does not affect other constraints, because the distribution of continuation utilities remains unchanged. It keeps the agent's utility constant while increasing the principal's payoff. This yields a contradiction.

3 Using parts (1) and (2) of the lemma, constraints WT and T are equivalent to

$$\mathbb{E}\left[u(a) + (1 - v^{\hat{R}a})\delta u_{\neg V}^{\hat{R}a} \;\middle|\; \hat{R}\right] = \mathbb{E}\left[u(a) + (1 - v^{\hat{L}a})\delta u_{\neg V}^{\hat{L}a} \;\middle|\; \hat{L}\right],$$

and

$$\rho v^{\hat{L}}\delta u_V^{\hat{L}} + (1 - \rho)v^{\hat{R}}\delta u_V^{\hat{R}} \;\geq\; c_A.$$

By Lemma 1, a mean-preserving contraction between u_V^R and u_V^L leaves all incentive constraints unchanged while weakly increasing the principal's payoff. Hence $u_V^R = u_V^L$.

$\square$